

Reinforcement Learning Lecture Notes

A Student

October 28, 2025

These are unofficial lecture notes for the course Reinforcement Learning [M-INFO-105623] at the Karlsruhe Institute of Technology (KIT). I began writing them to deepen my own understanding of the material. The structure of the notes reflects my personal learning process and what made the most sense to me. At times, I used language models to improve phrasing, since English isn't my first language. If anything sounds off or unclear, feel free to let me know or, if you have access to the source files, you're welcome to make edits yourself.

The content is based on a variety of sources, including lectures, blog posts, and research papers, which I've tried to reference wherever possible. If you're interested in exploring more about reinforcement learning, here are some great additional resources and courses in addition to the course at KIT:

- Sergey Levine's course: CS285: Deep Reinforcement Learning 2020 [[RAI20](#), [HLF⁺19](#)]
- Katerina Fragkiadaki's lectures: CMU 10403: Deep Reinforcement Learning and Control (Spring 2019) at Carnegie Mellon University, as well as her newer courses [[Fra19](#)]
- Emma Brunskill's course: CS234: Reinforcement Learning (Spring 2024) at Stanford University [[Onl24](#)]
- The lecture series from Google DeepMind in collaboration with UCL [[Dee21](#)]
- The Deep RL Bootcamp lecture series from UC Berkeley [[Pri18](#)]

Of course, you don't need to go through all of these courses, sometimes one instructor's explanation of a topic just clicks better than another's. For a solid foundational textbook, I highly recommend Reinforcement Learning: An Introduction by Sutton and Barto [[SB18](#)]. Please note that these notes are not a substitute for the actual course; they're meant to be used as a supplement to it.

Contents

1	What is Reinforcement Learning	6
2	Markov Decision Processes	6
3	Terminology	7
3.1	Policy	7
3.2	Episodes, Roll-out, Trajectories	7
3.3	Reward and Return	7
3.4	Value Functions	7
3.4.1	Value Function	7
3.4.2	Action Value Function	7
3.4.3	Advantage Function	8
3.5	Optimisation objective	8
3.6	Bellman Equations	8
3.6.1	Bellman Optimality Equations	9
4	Optimal decision making via dynamic programming	10
4.1	Self-Test Questions	11
5	K-Armed-Bandits	12
5.1	Greedy action selection	12
5.2	Exploration and Entropy	12
5.3	Optimistic Value Initialization	13
5.4	Upper-Confidence-Bounds (UCB)	13
6	Temporal Difference (TD)- Methods	16
6.1	Monte Carlo policy evaluation	16
6.2	Temporal Difference Learning	17
6.3	Generalized Policy Improvement	17
6.3.1	TD-Learning with Q-Values	18
6.3.2	Multi-Step Returns	19
6.4	Function Approximation for policy evaluation	20
6.4.1	Max Action Calculation	21
6.4.2	Improving DQN	21
6.4.3	Double Deep Q Learning	22
6.4.4	Dueling DQN	23
6.5	Self-Test Questions	23
6.6	Resources	23
7	Policy Gradients	24
7.1	Use Temporal Structure	25
7.2	REINFORCE	25
7.3	Baseline	26
7.4	Data Reuse	26
7.5	Policy Gradients with Q-Values	27
7.6	Advantage Estimation	27
7.6.1	Generalized Advantage Estimation	28
7.7	Self-Test Questions	28
7.8	Resources	29
8	Actor Critic	30

9	Off Policy Actor Critic	31
9.1	First Attempt	31
9.2	Updating the critic	32
9.3	Deterministic updates	32
9.3.1	Deep Deterministic Policy Gradient (DDPG)	32
9.3.2	Twin Delayed Deep Deterministic (TD3)	32
9.4	Variational Actor Updates (Soft Actor Critic (SAC))	34
9.4.1	Max-Entropy Reinforcement Learning	34
9.4.2	Actor and Critic	34
9.4.3	Reparametrization Trick	35
9.5	CrossQ	37
9.6	BRO	37
9.7	Self-Test Questions	37
9.8	Resources	38
10	Control as Inference	39
10.1	Variational Inference	39
10.1.1	Minimizing Reverse KL	39
10.1.2	Minimizing Forward KL	39
10.2	Connection between Variational Inference and Control	40
10.3	Resources	42
11	Trust Regions for Policy Gradients	43
11.1	Trust regions	43
11.2	Lagrangian Multipliers	43
11.2.1	Lagrangian Multipliers cookbook	44
11.3	Trust Region Policy Optimization (TRPO)	44
11.4	Proximal Policy Optimization (PPO)	45
11.5	Differentiable Trust Region Layers	47
11.6	Maximum Posteriori Optimization (MPO)	47
11.7	Self-Test Questions	47
11.8	Resources	48
12	Evolution Strategies	49
12.1	Gaussian Evolution Strategies	49
12.1.1	Canonical Evolution Strategy (CES)	50
12.1.2	Cross-Entropy Method (CEM)	50
12.2	Trust-Region Methods for Stochastic Search	51
12.3	Self-Test Questions	51
12.4	Resources	52
13	Bayesian Optimization	53
13.1	Surrogate Model (Gaussian Process Regression)	53
13.2	Acquisition Function	53
13.3	Optimization Loop	53
13.4	Self-Test Questions	54
14	Planning	55
14.1	Open-/Closed-Loop Planning	55
14.2	Closed-Loop Planning in Discrete Domains - Monte-Carlo Tree Search	56
14.2.1	Alpha-Go and its successors	56
14.3	Resources	57

15 Model Predictive Control / Closed-Loop Planning in Continuous Domains – Optimal Control	58
15.1 Linear Quadratic Regulator	58
15.2 Linear–quadratic–Gaussian Control	58
15.3 Approximating Non-Linear Systems (iLQR)	59
15.4 Resources	59
16 Model-Based Reinforcement Learning (MBRL)	60
16.1 Distributional shift	60
16.2 Model Exploitation	60
16.3 Uncertainty	61
16.3.1 Epistemic uncertainty	62
16.3.2 Aleatoric uncertainty	62
16.4 Replanning	62
16.5 Probabilistic Ensembles with Trajectory Sampling (PETS)	62
16.5.1 Core Components	63
16.6 Model-Based Policy Optimization (MBPO)	63
16.7 Model-Based-RL under Partial Observability	63
16.8 Self-Test Questions	64
16.9 Resources	65
17 Imitation Learning	66
17.1 Behavioural Cloning	66
17.1.1 Distributional Shift	66
17.2 Inverse Reinforcement Learning	67
17.2.1 Feature Matching	67
17.2.2 Max-Entropy Inverse Reinforcement Learning	68
17.2.3 Adversarial Methods	69
17.3 Challenges of human data	70
17.3.1 Trajectory-ranked Reward EXtrapolation (T-REX)	71
17.3.2 Goal-Conditioned Imitation Learning	71
17.4 Resources	72
18 Offline Reinforcement Learning	73
18.1 Challenges in Offline-RL	74
18.1.1 Intuitive Explanation	75
18.1.2 Mathematical Explanation	75
18.2 Approaches to Offline-RL	76
18.2.1 Explicit Policy Constraints	77
18.2.2 Implicit Policy Constraints	77
18.2.3 Implicit Q-Learning	77
18.2.4 Conservative Q-Learning	78
18.3 Resources	78
19 Transformers in Reinforcement Learning	79
19.1 Embedding	79
19.2 Positional Encoding	80
19.2.1 Binary Representation and Positional Encoding	81
19.3 Attention Mechanism	82
19.3.1 Query, Key, and Value Vectors	82
19.3.2 Attention Scores	82
19.3.3 Softmax	82
19.3.4 Updating the Embedding	82
19.3.5 Multi-Head Attention	83
19.4 Feedforward Neural Network and Add & Norm	84
19.5 Modifications to the Vanilla Transformer	84
19.5.1 Transformer-XL	84

19.5.2	Stabilizing Transformers for Reinforcement Learning	85
19.6	Transformers in RL	86
19.7	Pros and Cons of Transformers in RL	86
19.8	Resources	86
20	Generative Models	87
20.1	Autoencoders	87
20.2	Variational Autoencoders (VAEs)	87
20.3	Diffusion	88
20.3.1	Forward Diffusion Process	89
20.3.2	Reverse Diffusion Process	90
20.3.3	Training	90
20.4	Score-Based Generative Models	90
20.5	Flow Matching	91
20.6	Connecting Diffusion Models and Reinforcement Learning	91
20.7	Resources	91
21	Distributional Reinforcement Learning	92
21.1	Resources	92

1 What is Reinforcement Learning

Reinforcement learning (RL), a core machine learning paradigm alongside supervised and unsupervised learning, integrates machine learning and optimal control to enable an intelligent agent to maximize a reward signal in an environment. It formalises the idea that rewarding or punishing an agent for its behaviour makes it more likely to repeat or refrain from that behaviour in the future.




Figure 1: Agent-Environment interaction loop

The main characters are the agent and the environment. At every step, the agent receives a reward and either the full state or just an observation of the environment. A state is a complete description of the world, like the configuration of a chess game. An observation, though, is a partial description. Imagine an agent in a video game where the state is defined by underlying game mechanics. Since most of the time you don't actually have access to these things you might only infer information about the world through the current frame (this is the observation).

The reward indicates how good or bad the agent's last action was, providing feedback about what happened and how the environment might have changed (sometimes the environment changes on its own). Based on this, the agent selects its next action, aiming to accumulate as much reward as possible over time, known as the return.

2 Markov Decision Processes

The Markov Decision Process (MDP) provides the standard mathematical framework for modelling the interaction between an agent and its environment. It serves as a formal model for sequential decision-making under uncertainty and is defined as a 5-tuple $(S, A, P_{sa}, \gamma, R)$, where:

- S is a set of states
- A is a set of actions
- P_{sa} are the state transition probabilities. It describes the probability of getting into a state s_{t+1} when being in a state s_t and performing action a_t , i.e. $P_{sa}(s_{t+1}|s_t, a_t)$
- $\gamma \in [0, 1)$ is the discount factor
- R is the reward function. It can depend on the actions and the states, but sometimes also only on the states

A defining feature of an MDP is the Markov property, which says that the outcome of an action depends only on the current state and not on the prior history of states or actions. Mathematically, this is expressed as

$$P_{sa}(s_{t+1}|s_t, a_t, \dots, s_0, a_0) = P_{sa}(s_{t+1}|s_t, a_t).$$

3 Terminology

To understand the key concepts and algorithms of reinforcement learning, we need to introduce some terminology.

3.1 Policy

A policy π determines what action an agent should take when it is in a particular state. A deterministic policy is defined as $\pi(s) = a$. A stochastic policy is defined as $\pi(a|s)$. It tells us what the probability is of performing the action a while in state s . The probability function here is not the same as the state transition probabilities P_{sa} !

3.2 Episodes, Roll-out, Trajectories

All three represent a sequence of (s, a, r) tuples, where an episode is generally from an initial state to a terminal state, a roll-out is from an initial or intermediate state to a terminal state, and a trajectory is between any two states. But there is no clear distinction in the RL-World

3.3 Reward and Return

The goal of the agent is to maximise some notion of cumulative reward over episodes, given a sequence of state action reward tuples (s_t, a_t, r_t) . The total **discounted return** from time step t , called G_t , is defined as

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+1+k}.$$

The restriction of the discount factor γ to the interval $[0, 1)$ is now more evident. Allowing $\gamma = 1$ in the context of an infinite episode would result in an infinite return G_t , which offers little insight into the agent's behaviour. For finite episodes, setting $\gamma = 1$ is permissible. The value of γ also significantly influences the agent's decision-making. When γ approaches 0, the agent behaves "myopic", prioritizing the maximization of immediate rewards. As γ approaches 1, the agent becomes "farsighted", placing greater emphasis on long-term future rewards.

3.4 Value Functions

In order to maximise the return, the agent needs to know which actions to take in each state. Therefore, it is often useful to know "how good" a particular action performed in a particular state is.

3.4.1 Value Function

The state value function $V^\pi(s)$ of an MDP is the expected discounted return starting from state s and following a policy π

$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s]. \quad (1)$$

3.4.2 Action Value Function


The action-value function represents the expected return when starting in state s , taking action a , and thereafter following policy π

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t | s_t = s, a_t = a]. \quad (2)$$

The value function V and the action-value function Q can also be defined in terms of each other:

$$V^\pi(s) = Q^\pi(s, \pi(s)) = \mathbb{E}_{a \sim \pi}[Q^\pi(s, a)]. \quad (3)$$

Figure 2: Backup diagram illustrating the different value functions.



3.4.3 Advantage Function

In RL, we often care about the relative advantage of an action over some other action rather than the absolute value. The advantage function $A^\pi(s, a)$ quantifies how much better taking action a in state s is compared to randomly selecting an action according to $\pi(\cdot|s)$ assuming the policy π is followed thereafter. Mathematically, the advantage function is defined by

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s).$$

A thing to note is that $\mathbb{E}_{a \sim \pi}[A^\pi(s, a)] = 0$ since

$$\mathbb{E}_{a \sim \pi}[A^\pi(s, a)] = \mathbb{E}_{a \sim \pi}[Q^\pi(s, a) - V^\pi(s)] = \mathbb{E}_{a \sim \pi}[Q^\pi(s, a)] - \mathbb{E}_{a \sim \pi}[V^\pi(s)] \stackrel{(3)}{=} V^\pi(s) - V^\pi(s) = 0^1.$$

3.5 Optimisation objective

As previously noted, the agent's objective is to find a policy that maximizes the return. To achieve this, the agent seeks to optimize either the value function or the state-action value function, which then serves as a foundation for deriving the policy. The optimal policy is defined as follows

$$V^*(s) = \max_{\pi} \mathbb{E}_{\pi}[G_t | s_t = s]$$

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\pi}[G_t | s_t = s, a_t = a].$$

Given the relationship between the value function and the state-action value function, it follows that $V^*(s) = \max_a Q^*(s, a)$. With Q^* determined, the optimal policy can be straightforwardly expressed as $\pi^*(s) = \arg \max_a Q^*(s, a)$. In the subsequent sections, we will explore algorithms designed to compute the optimal policy across various environmental settings.²

3.6 Bellman Equations

The Bellman equations are a fundamental set of equations in reinforcement learning and dynamic programming that express a relationship between the value of a state (or state-action pair) and the values of its successor states.

The Bellman equation for the state value function $V^\pi(s)$ can be derived as follows

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_{\pi}[G_t | s_t = s] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s\right] \\ &= \mathbb{E}_{\pi}[r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} \dots) | s_t = s] \\ &= \mathbb{E}_{\pi}[r_{t+1} + \gamma G_{t+1} | s_t = s] \\ &= \mathbb{E}_{\pi}[r_{t+1} + \gamma V(s_{t+1}) | s_t = s]. \end{aligned} \tag{4}$$

Similarly, the action-value function $Q^\pi(s, a)$ can be expressed recursively

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E}_{\pi}[G_t | s_t = s, a_t = a] \\ &= \mathbb{E}_{\pi}[r_{t+1} + \gamma V(s_{t+1}) | s_t = s, a_t = a] \quad (\text{Eq. (4)}) \\ &= \mathbb{E}_{\pi}[r_{t+1} + \gamma \mathbb{E}_{a \sim \pi}[Q^\pi(s_{t+1}, a_{t+1})] | s_t = s, a_t = a]. \quad (\text{Eq. (3)}) \end{aligned} \tag{5}$$

The V-Bellman (4) and Q-Bellman (5) equations can be decomposed further by expanding the expectations

$$V^\pi(s) = \sum_a \pi(a | s) \left(R(s, a) + \gamma \sum_{s'} p(s' | s, a) V^\pi(s') \right), \tag{6}$$

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s'} p(s' | s, a) \sum_{a'} \pi(a' | s') Q^\pi(s', a').^3 \tag{7}$$

¹The reason why $\mathbb{E}_{a \sim \pi}[V^\pi(s)] = V^\pi(s)$ is that $\mathbb{E}[\mathbb{E}[X]] = \mathbb{E}[X]$, and since $V^\pi(s)$ is itself an expectation, this holds [1h].

²In the literature, the term "Bellman backup" is sometimes used. It simply refers to the sum of the immediate reward and the value of the next state. $Q^\pi(s, a) = \mathbb{E}_{\pi}[\underbrace{r_{t+1} + \gamma Q^\pi(s_{t+1}, a_{t+1})}_{\text{Bellman backup}} | s_t = s, a_t = a]$

3.6.1 Bellman Optimality Equations

the optimal value function can also be expressed in terms of the Bellman equation

$$\begin{aligned}
V^*(s) &= \max_a \mathbb{E}_{s' \sim P(\cdot | s, a)} [r(s, a) + \gamma V^*(s')] \\
&= \max_a \left(R(s, a) + \gamma \sum_{s'} p(s' | s, a) V^*(s') \right), \\
Q^*(s, a) &= \mathbb{E}_{s' \sim P(\cdot | s, a)} \left[r(s, a) + \gamma \max_{a'} Q^*(s', a') \right] \\
&= R(s, a) + \gamma \sum_{s'} p(s' | s, a) \max_{a'} Q^*(s', a').
\end{aligned}$$

The idea of Bellman's optimality principle is that an optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with respect to the state resulting from the first decision.

³Since $\sum_{a'} \pi(a' | s') Q^\pi(s', a') = \mathbb{E}_{a' \sim \pi} [Q^\pi(s', a')] = V^\pi(s') \rightarrow Q^\pi(s, a) = R(s, a) + \gamma \sum_{s'} p(s' | s, a) V(s')$

4 Optimal decision making via dynamic programming

We begin with the simplest scenario, where the agent has full knowledge of the environment, including the state transition probabilities and the reward function. A straightforward approach to computing the value function for a given policy π is the Policy-Evaluation algorithm. This method iteratively employs the Bellman equation (6) to update and refine the value function.

Algorithm 1 Policy Evaluation

```

 $V_0^\pi(s) \leftarrow 0$ 
for  $k = 1$  until convergence do
  for all  $s$  in  $S$  do
     $V_k^\pi(s) \leftarrow \sum_a \pi(a|s) (R(s, a) + \gamma \sum_{s'} p(s'|s, a) V_{k-1}^\pi(s'))$ 
  end for
end for

```

Now we know how to evaluate a policy, but we are interested in finding the optimal policy. One naive way is to try every possible policy to find the optimal one. But already the number of deterministic policies is $|A|^{|S|}$ (because we have in each state $|A|$ possible actions we could take), which is intractable. So we use a more efficient algorithm called Policy-Iteration (2). The idea is to start with a random policy and improve it iteratively.

Algorithm 2 Policy Iteration

```

 $\pi_0(s) \leftarrow$  randomly for all states  $s$ 
while  $\|\pi_i - \pi_{i-1}\|_1 > 0$  do
   $V^{\pi_i} \leftarrow$  Policy Evaluation of  $\pi_i$ 
   $\pi_{i+1} \leftarrow$  Policy Improvement
   $i \leftarrow i + 1$ 
end while

```

Algorithm 3 Policy Improvement

```

for  $s$  in  $S$  and  $a$  in  $A$  do
   $Q^{\pi_i}(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} p(s'|s, a) V^{\pi_i}(s')$ 
end for
for  $s$  in  $S$  do
   $\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$ 
end for

```

It can be shown that this procedure leads to the optimal policy due to its monotonic improvement of the policy in every iteration, meaning that $V^{\pi_i}(s) \leq V^{\pi_{i+1}}(s)$.

Another way to find the optimal policy in this setting is to use the Value-Iteration algorithm (4). It can also be demonstrated that value iteration converges to an optimal solution.

Algorithm 4 Value Iteration

```

 $V_0^\pi(s) \leftarrow 0 (\forall s)$ 
while not converged (for ex.  $\|V_{k+1} - V_k\| > \epsilon$ ) do
  for  $s$  in  $S$  do
     $V_{k+1}(s) = \max_a [R(s, a) + \gamma \sum_{s'} p(s'|s, a) V_k(s')]$ 
     $\pi_{k+1}(s) = \arg \max_a [R(s, a) + \gamma \sum_{s'} p(s'|s, a) V_k(s')]$ 
  end for
end while

```

The problem with the Policy- and the Value-Iteration algorithm is that they require full knowledge of the environment, i.e. dynamics model and reward function. This is often not available, so we will now look at methods that allow us to find a optimal policy even if we do not have a model of the environment.

4.1 Self-Test Questions

1. **The definition of a MDP** → 2

2. **Why do we need discounting ?**

If we would no discount then our return (discounted future rewards) would be infinite which results in every policy being an optimal policy and making it hard to come up with algorithms to find the optimal policy in these cases. Another thing is that we can control the behaviour of the agent by choosing the the discount factor. If $\gamma \rightarrow 0$ (myopic) then focused on maximizing immediate reward. If $\gamma \rightarrow 1$ (farsighted) focused on future rewards.

3. **The definition of the optimal and the policy based V- and Q-Functions**

$$V^\pi(s) = \mathbb{E}[G_t | s_t = s], \quad V^*(s) = \max_{\pi} V^\pi(s) \\ Q^\pi(s, a) = \mathbb{E}[G_t | s_t = s, a_t = a], \quad Q^*(s, a) = \max_{\pi} Q^\pi(s, a)$$

4. **What is the bellman principle of optimality ?**

An optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with regard to the state resulting from the first decision.

$$V^*(s) = \max_a \left(R(s, a) + \gamma \sum_{s'} p(s' | s, a) V^*(s') \right) \\ Q^*(s, a) = R(s, a) + \gamma \sum_{s'} p(s' | s, a) \max_{a'} Q^*(s', a')$$

5. **How the value iteration algorithm works** → 4

6. **How the policy iteration algorithm works** → 2

7. **If in a policy iteration step the policy doesn't change, can it ever change again? (not from the lecture)**

No (if π^* is unique $\forall s$). Because if $\pi_i = \pi_{i+1}$ then $Q^{\pi_i} = Q^{\pi_{i+1}}$ and this results in the policy never changing again.

8. **Is there a maximum number of iterations of policy iteration? (not from the lecture)**

$|A|^{|S|}$ since that is the maximum number of policies, and as the policy improvement step is monotonically improving, each policy can only appear in one round of policy iteration unless it is an optimal policy.

9. **What are the conceptual differences between value- and policy- iteration?**

In policy iteration we first do multiple passes that update the value function and then update the policy with these newly computed values. In value iteration when updating the value function we simultaneously/implicitly also update the policy. Value iteration can be seen as an extreme case of policy iteration where we stop the evaluation after one update.

10. **What are the limitations of these approaches ?**

We must know the dynamics and reward model and it only works for discrete settings or Linear Systems, Quadratic Reward, Gaussian Noise (LQR)

5 K-Armed-Bandits

Before we look at algorithms for the case where we do not have a model of the world, we will talk about the k-armed-bandit problem, which helps us to understand decision making under uncertainty. A k-armed bandit is a tuple $(A, p(r|a))$, where A is the known set of k actions we can take (arm) and $p(r|a)$ is the unknown probability distribution over rewards. At each time step t we take an action a_t and the environment generates a reward r_t . The goal is to maximise the expected reward. To achieve this goal we have to decide which arm to pull, so we need to know the expected reward per action $\hat{q}$ (action value function)

$$\hat{q} = \mathbb{E}[r_t|a_t = a] = \int p(r|a)r \, dr$$

Since we do not know $p(r|a)$, we approximate the integral by Monte Carlo simulation.

$$\hat{q} = \mathbb{E}[r_t|a_t = a] \approx \frac{\sum_{i=1}^{t-1} \mathbb{1}(a_i = a)r}{\sum_{i=1}^{t-1} \mathbb{1}(a_i = a)} = q_t(a)$$

Instead of storing all the rewards to calculate the sum, one can also reformulate it using an incremental update rule.

$$\begin{aligned} q_{n+1} &= \frac{1}{n} \sum_{i=1}^n r_i = \frac{1}{n} \left(r_n + \sum_{i=1}^{n-1} r_i \right) \\ &= \frac{1}{n} \left(r_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} r_i \right) \\ &= \frac{1}{n} (r_n + (n-1)q_n) \\ &= q_n + \frac{1}{n}(r_n - q_n) \end{aligned} \tag{8}$$

This formulation is called the sampled average method, whereas the more general update rule looks like this.

$$q_{n+1} = q_n + \alpha_n(r_n - q_n)$$

We will now consider different strategies for choosing actions to maximise the expected reward.

5.1 Greedy action selection

In greedy action selection, our next action a_t is defined as

$$a_t = \arg \max_{a \in A} q_t(a)$$

The problem with this approach is that we could get stuck in a suboptimal policy. Consider the following example where we have three actions a_1, a_2, a_3 to choose from with unknown true rewards $q(a_1) = 0.95, q(a_2) = 0.9, q(a_3) = 0.1$. We know have as the initial q-values the values $q(a_1) = q(a_3) = 0, q(a_2) = 1$. From there on we do greedy action selection, which results in only picking the action a_2 even though the actual best action to take is a_1 , but for that we would have to drift from our strategy. This is the core problem in RL. We need to make a trade-off between Exploration: Improve knowledge for long term benefit and Exploitation: Using knowledge for short-term benefit. One way to achieve this is called ϵ -greedy exploration, where we choose a value $\epsilon \in [0, 1]$ representing the probability of exploring/exploiting.

$$a = \begin{cases} \arg \max_a q_t(a) & \text{with probability } 1 - \epsilon \\ \text{Uniform}(a_1, \dots, a_k) & \text{with probability } \epsilon \end{cases}$$

5.2 Exploration and Entropy

Although ϵ greedy action selection is an effective means of balancing exploration and exploitation in reinforcement learning, one drawback is that when it explores, it chooses equally among all actions.

This means that it is just as likely to choose the worst appearing action as it does the next best action. In tasks where the worst actions are very bad, this can be unsatisfactory. The obvious solution is to vary the action probabilities as a graded function of the estimated value (exploitation). But at the same time, we are uncertain about the rewards and therefore need to keep uncertainty in the policy, for which entropy H is a good measure (exploration). Entropy is a concept from information theory that measures the randomness or uncertainty of a probability distribution, which, in our case, is the stochastic policy $\pi(a|s)$. A policy with High Entropy means the action probabilities are nearly equal for all actions. This signifies high uncertainty on the agent's part, as its action choices are very random. This state promotes greater exploration because the agent is not strongly committed to any single action.

Conversely, a policy with Low Entropy means the agent assigns a high probability to one or a few actions (approaching a deterministic policy). This signifies high confidence in those actions, as the agent's choices are predictable. This state promotes greater exploitation as the agent focuses on the action it believes to be optimal. The new the objective which includes this idea of entropy looks like this

$$\begin{aligned}\pi_t^* &= \arg \max_{\pi} \sum_a \pi(a) q_t(a) + \tau^{-1} H(\pi) = \arg \max_{\pi} \sum_a \pi(a) q_t(a) + \tau^{-1} \left(- \sum_a \pi(a) \log \pi(a) \right) \\ &= \arg \max_{\pi} \sum_a \pi(a) (q_t(a) - \tau^{-1} \log \pi(a))\end{aligned}$$

Calculation the derivative and setting it equal to zero gives us the policy which maximises the objective

$$\pi(a) = \frac{\exp(\tau q_t(a))}{\sum_{a'} \exp(\tau q_t(a'))}$$

5.3 Optimistic Value Initialization

The idea is simply to initialise the q-values with optimistic values, this has the effect of adding a positive bias to less explored actions, thus forcing exploration.

5.4 Upper-Confidence-Bounds (UCB)

The UCB formula is defined as

$$a_t = \arg \max_a \underbrace{q_t(a)}_{\text{exploitation}} + c \underbrace{\sqrt{\frac{\log t}{N_t(a)}}}_{\text{exploration}}$$

c is a constant that allows us to set the exploration/exploitation trade-off. $\sqrt{1/N_t(a)}$ essentially says that the more times you pull an arm, there is less unknown about that arm. So the fewer times action a is chosen, the higher the exploration bonus. $\sqrt{\log t}$ ensures that you don't stop exploring too early. In the following, we will examine how different strategies perform across various parameter settings. To do this, we utilize a 10-armed bandit experiment, averaging the results over 2000 runs with sample average action-value estimates.




Figure 3: ϵ -Greedy




Figure 4: Softmax




Figure 5: UCB




Figure 6: Comparison of different strategies

6 Temporal Difference (TD)- Methods

We now shift to the setting where we lack a model of the environment from which to derive an optimal policy. Instead, we are provided with sample trajectories, denoted as $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots)$, and must learn from these directly.

6.1 Monte Carlo policy evaluation

Recall that the value function under a policy π is defined as $V^\pi(s) = \mathbb{E}_\pi [G_t | s_t = s]$, where G_t represents the discounted return starting from state s at time t . The core idea of Monte Carlo methods is to estimate an expectation, $\mathbb{E}_{p(x)} [f(x)] = \int p(x)f(x) dx$, by sampling from the distribution and averaging the results:

$$\mathbb{E}_{p(x)} [f(x)] = \int p(s)f(x)dx \approx \frac{1}{N} \sum_{x_i \sim p(x)} f(x_i)$$

This translates to sampling random episodes or trajectories under the policy π and averaging the returns observed for each state to approximate $V^\pi(s)$. There are two primary approaches to compute this estimate: First-Visit Monte Carlo and Every-Visit Monte Carlo. These methods differ in how they handle state visits within an episode and update the value estimates. The First-Visit approach updates the estimate only the first time a state s is encountered at a timestep t in an episode. In contrast, the Every-Visit approach updates the estimate for every occurrence of a state across all timesteps in the episode. At their core, they work as follows:

Algorithm 5 Monte Carlo Policy Evaluation

```

 $N(s) \leftarrow 0, G(s) \leftarrow 0 \forall s \in S$ 
for do
    sample episode  $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}, a_{i,T_i}, r_{i,T_i}$ 
    define  $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-t} r_{i,T_i}$  as the return from time step  $t$  onwards in the  $i$ th episode
    for  $t \in T_i$  do
        First-Visit Monte Carlo
        if this is the first time  $t$  that state  $s$  is visited in episode  $i$ 
             $N(s) \leftarrow N(s) + 1$ 
             $V^\pi(s) \leftarrow V^\pi(s) + \alpha(G_{i,t} - V^\pi(s))$  {for  $\alpha = \frac{1}{N(s)}$  equal to running average see (8)}
        Every-Visit Monte Carlo
        state  $s$  is being visited at time step  $t$  in episode  $i$ 
             $N(s) \leftarrow N(s) + 1$ 
             $V^\pi(s) \leftarrow V^\pi(s) + \alpha(G_{i,t} - V^\pi(s))$  {for  $\alpha = \frac{1}{N(s)}$  equal to running average see (8)}
    end for
end for

```

Convergence can be guaranteed depending on the choice of α , provided the following conditions are satisfied:

- $\sum_{n=1}^{\infty} \alpha_n(s, a) = \infty \forall (s, a)$
- $\sum_{n=1}^{\infty} \alpha_n^2(s, a) < \infty \forall (s, a)$

The sample average method, using $\alpha_n = \frac{1}{n}$, fulfills these requirements. Additionally, it's worth noting that the First-Visit Monte Carlo method is an unbiased estimator, whereas the Every-Visit Monte Carlo method is both unbiased and consistent (meaning $\forall \epsilon > 0 : \lim_{n \rightarrow \infty} P(|V_n^\pi - V| > \epsilon) = 0$).

The problem with the Monte Carlo estimator is that it is a high variance estimator, because if the policy or the environment is highly stochastic, the total return G_t from a state can vary a lot even

if you start from the same state. So, you're averaging a bunch of noisy, random returns = variance. Reducing the variance can require a lot of data. Therefore, in cases where data is expensive to obtain MC may be impractical. Another disadvantage is that it only works if we have an episodic setting, which makes it very sample inefficient because we have to wait for the whole episode to end before we can compute the returns. But a nice feature is that it does not require the model to be Markov.

6.2 Temporal Difference Learning

Temporal Difference (TD) Learning combines the Monte Carlo method with dynamic programming which is used in the Value- and Policy-Iteration algorithms. This model-free technique is applicable to both episodic and non-episodic environments. It updates the value function V immediately after each transition tuple (s, a, r, s') , with an update rule resembling that of incremental Monte Carlo methods:

$$V^\pi(s) = V^\pi(s) + \alpha \underbrace{[r + \gamma V^\pi(s')] - V^\pi(s)}_{\text{TD target}}$$

This version is known as TD(0), as it uses a one-step lookahead/ one-step return for updating the value estimates.

Algorithm 6 Temporal Difference (TD(0)) Learning Algorithm

```

 $V_0^\pi(s) \leftarrow 0, \forall s \in S$ 
Loop for each episode:
  init.  $s$ 
  Loop
    sample tuple  $(s, a, r, s')$  from  $\pi$ 
     $V^\pi(s) = V^\pi(s) + \alpha(r + \gamma V^\pi(s')) - V^\pi(s)$ 
     $s \leftarrow s'$ 
  until  $s$  is terminal

```

A key limitation of this approach is its inability to directly support policy improvement (see algo. 3). To compute the Q-values, we rely on the dynamics model, but such a model is unavailable in a model-free setting:

$$Q^{\pi_i}(s, a) \leftarrow R(s, a) + \gamma \sum_{s'} p(s'|s, a) V^{\pi_i}(s')$$

To address this, we shift to learning the Q-function directly instead of the value function, enabling policy improvement without dependence on the dynamics model.




Figure 7: Visualization of Monte Carlo methods (left), Temporal Difference (TD) learning (middle), and Dynamic Programming (right) in a backup diagram. The red highlights indicate how each method computes value functions: Monte Carlo uses the entire episode, TD learning performs a one-step lookahead using bootstrapping, and Dynamic Programming computes exact expectations but only when the transition probabilities are known.

6.3 Generalized Policy Improvement

The first algorithm we'll look at is Monte Carlo Online Control, which builds on Monte Carlo Policy Evaluation by incorporating updates to the Q-values.

Algorithm 7 Monte Carlo Online Control

```
 $N(s, a) \leftarrow 0, Q(s, a) \leftarrow 0 \forall (s, a) \in S \times A$ 
for (over episodes) do
  sample episode  $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}, a_{i,T_i}, r_{i,T_i}$  according to  $\pi$ 
  define  $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-t} r_{i,T_i}$ 
  for  $t \in T_i$  do
    if First-Visit // Every-Visit then
       $N(s, a) = N(s, a) + 1$ 
       $Q(s, a) = Q(s, a) + \frac{1}{N(s, a)}(G_{i,t} - Q(s, a))$ 
       $\pi(s) \leftarrow \arg \max_{a'} Q(s, a')$ 
    end if
  end for
end for
```

The Monte Carlo Online Control algorithm ⁴ converges to the optimal state-action value function Q^* if the two conditions are fulfilled:

- Every state-action pair (s, a) is visited infinitely often, expressed as $\lim_{i \rightarrow \infty} N_i(s, a) \rightarrow \infty$.
- The behavior policy (the policy guiding actions in the environment) converges to a greedy policy with probability 1, i.e., $\lim_{i \rightarrow \infty} \pi(a|s) \rightarrow \arg \max_a Q(s, a)$.

These conditions are sometimes referred to as the "Greedy in the Limit with Infinite Exploration" (GLIE) property.

6.3.1 TD-Learning with Q-Values

Next, we will examine two variants of Temporal Difference (TD) Learning that update Q-values rather than value functions: SARSA and Q-Learning. While these methods are nearly identical, they only differ in their Q-value update rules.

We begin with SARSA, the on-policy variant. In SARSA, the next action used to compute the Q-value is determined by the current policy.

Algorithm 8 SARSA (on-policy TD control)

```
 $Q(s, a) \leftarrow 0, \forall (s, a) \in S \times A$ 
Loop for each episode:
  init. s
  choose a from s using policy derived from Q(e.g.,  $\epsilon$ -greedy)
  Loop for each step of episode
    take action a, observe  $r, s'$ 
    choose  $a'$  from  $s'$  using policy derived from Q(e.g.,  $\epsilon$ -greedy)
     $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s', a \leftarrow a'$ 
  until s is terminal
```

Q-Learning closely resembles SARSA, but operates off-policy due to the difference in its Q-value update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

The max operator is what makes Q-Learning off-policy. Rather than sticking to the current policy to select the next action, it greedily chooses the action with the highest Q-value, diverging from the policy being followed. Both SARSA and Q-Learning can converge to the optimal solutions provided they satisfy the GLIE condition (see 6.3) and the step sizes α_t meet the following two criteria:

⁴Monte Carlo Online Control isn't sourced from the KIT lecture; it think i got it from Stanford's CS234 course, so feel free to skip it.

- $\sum_{n=1}^{\infty} \alpha_n(s, a) = \infty \forall (s, a)$
- $\sum_{n=1}^{\infty} (\alpha_n(s, a))^2 < \infty \forall (s, a)$

6.3.2 Multi-Step Returns

We have thus far considered two extremes for policy evaluation. On one hand, Monte Carlo methods use the full, unbiased episodic return (G_t), but they suffer from high variance because the estimate depends on the entire sequence of stochastic transitions and rewards until the episode's end. On the other hand, one-step Temporal-Difference learning (TD(0)) uses a low-variance estimate by bootstrapping after a single step ($R_{t+1} + \gamma V(S_{t+1})$), but this introduces significant bias from its reliance on the current (and likely imperfect) value estimate.

To find a middle ground and control this bias-variance trade-off, we can use *n-step returns*. The core idea is to look n steps into the future, accumulating actual rewards, and only then bootstrap using the value estimate of the resulting state. The **n -step return** at time t , denoted $G_t^{(n)}$, is formally defined as:

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

where $V(S_{t+n})$ is the current estimate of the value of state S_{t+n} . If the episode terminates before $t+n$, the return is simply the full Monte Carlo return G_t , and $V(S_{\text{terminal}}) \equiv 0$. The corresponding n -step TD update for the state-value function $V(S_t)$ is then:

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t^{(n)} - V(S_t)]$$

This formulation elegantly unifies MC and TD(0). Note that for $n = 1$, we recover the TD(0) update, as $G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$. As $n \rightarrow \infty$ (or n exceeds the episode length), $G_t^{(n)}$ becomes the standard Monte Carlo return G_t . The parameter n is now a hyperparameter. Rather than selecting a single "best" n , we can combine all possible n -step returns. This leads to the **λ -return**, G_t^λ , which creates a compound return by averaging all n -step returns, weighting each by $(1-\lambda)\lambda^{n-1}$. This is the foundation of the TD(λ) algorithm.

Eligibility traces However, a practical problem persists, which we first encountered with Monte Carlo methods. To compute the n -step return $G_t^{(n)}$ (or the full λ -return), we must wait until time $t+n$ (or the end of the episode for G_t^λ) to observe the required rewards and the state S_{t+n} . Only then can we compute the update for $V(S_t)$. This delay is undesirable, as it prevents the algorithm from being fully online. This "forward-view" perspective motivates the need for an equivalent "backward-view" mechanism, which is provided by eligibility traces.

To solve the problem of update delays, we introduce *eligibility traces*. This "backward-view" mechanism provides a computationally efficient and fully online method that is mathematically equivalent to the "forward-view" λ -return. The intuition is as follows: an eligibility trace is a short-term memory that tracks which states have been visited recently. When a "surprise" occurs (i.e., the TD error is high, meaning our prediction was wrong), we assign credit or blame for this surprise not just to the immediately preceding state, but to all recently visited states, proportional to their eligibility.

Mechanically, we maintain an eligibility trace, $E_t(s)$, for every state s . At the beginning of an episode, all traces are zero. At each time step t , the trace for the visited state S_t is incremented, and all other traces decay:

$$E_t(s) = \begin{cases} \gamma\lambda E_{t-1}(s) + 1 & \text{if } s = S_t \\ \gamma\lambda E_{t-1}(s) & \text{if } s \neq S_t \end{cases}$$

This is known as an **accumulating trace**. The parameter $\lambda \in [0, 1]$ controls the rate of decay:

- If $\lambda = 0$, $E_t(s)$ is 1 only for the current state $s = S_t$ and 0 otherwise. This reduces to the standard one-step TD(0) update.
- If $\lambda = 1$, the trace does not decay (within an episode), which relates the algorithm to Monte Carlo methods.

With eligibility traces, the update process becomes simple and online. At time $t + 1$, we observe the reward R_{t+1} and the next state S_{t+1} . We then compute the one-step TD error:

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

This single error signal δ_t is then used to update *all* states in the environment, weighted by their current eligibility trace:

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s), \quad \text{for all } s \in \mathcal{S}$$

This backward-view update, performed at every step, has been proven to be equivalent to the forward-view TD(λ) algorithm. It achieves the same effect as averaging all n -step returns, but does so efficiently, without waiting for future time steps.

But there is still a problem with the algorithms shown here. They only work when we can store all the values in tables. In fact, every algorithm so far only works if we can store the values in tables. For systems with many states and actions, this is a computational burden. And also the conditions for GLIE are somewhat impossible to fulfil if we have too many states and actions. Instead, we want a more compact representation that generalises across states or states and actions.

6.4 Function Approximation for policy evaluation

Rather than relying on tabular representations, we now use parametrized Q or V functions. These can range from simple linear models to complex neural networks. The algorithms introduced earlier remain mostly unchanged, except that we now also train a function approximator. Since the true Q-, V- or Reward-Function values are unavailable during training, we rely on approximated targets as labels. The most well-known function approximation algorithm is the Deep Q-Learning (DQN) algorithm, defined as follows:

Algorithm 9 Vanilla Deep-Q-Learning (DQN)

Input α
 init $\omega, t = 0$
 init. s
Loop for each episode:
 choose a from s using policy derived from $Q_\omega(s, a)$
 observe r, s'
 $y_i = \begin{cases} r & \text{episode terminated at step } i+1 \\ r + \gamma \max_{a'} Q_\omega(s', a') & \text{otherwise} \end{cases}$
 Do gradient descent step on $(y_i - Q_\omega(s, a))^2$
 $\omega \leftarrow \omega + \alpha(y_i - Q_\omega(s, a)) \nabla_\omega Q_\omega(s, a)$
 $t \leftarrow t + 1$
 $s \leftarrow s'$
end loop

One might argue that the gradient calculation of the loss contains an error, as y_i also depends on w , yet its derivative seems to be missing. This is intentional. The update rule incorporating the full gradient is known as the "Residual Gradient Algorithm". However, this approach minimizes the mean squared temporal difference error (MSTD), defined as:

$$\text{MSTD} = \mathbb{E}_\pi \left[\left(r_t + \gamma \max_{a'} Q_\omega(s', a') - Q_\omega(s, a) \right)^2 \right]$$

Minimizing the MSTD, however, does not guarantee Bellman optimality. Furthermore, the bias inherent in the residual gradient algorithm often results in suboptimal performance and slow convergence, making it less used in practice.

An alternative approach to achieve Bellman optimality involves minimizing the Mean Squared Bellman Error (MSBE), defined as:

$$\text{MSBE} = \mathbb{E}_{\pi} \left[\left(r_t + \gamma \mathbb{E}_{p(\cdot|s,a)} [\max_{a'} Q_{\omega}(s', a')] - Q_{\omega}(s, a) \right)^2 \right]$$

However, computing the inner expectation in this formulation requires knowledge of the transition model, which poses a significant challenge in the model free setting. One could simply use a single sample to estimate the expectation, which would effectively reduce the MSBE to the MSTD. However, there are also methods that employ two samples to address this, known as the double sampling approach. But these will not be discussed further.

Most of the time people will use the MSTD but don't take the gradient of the target values.

6.4.1 Max Action Calculation

To calculate the targets, we need to identify the action a' with the highest Q-value for a given state. In discrete settings, the typical approach is to evaluate all possible actions and select the best one. However, this becomes highly inefficient and impractical for high-dimensional or continuous action spaces. There are several strategies to address this challenge, which we will explore in later chapters. Some of these include:

- **Gradient optimization on the Q-value network:** One approach is to optimize the Q-value network directly to find the best action. However, this introduces an additional inner loop alongside the existing loop, which can significantly slow down the algorithm.
- **Random shooting:** This method involves sampling a set of actions $(a_1, \dots, a_n)$ (e.g., uniformly), and selecting the best action from this set to maximize the Q-value when calculating the target. While this approach is valid, it does not guarantee optimal results.
- **Iterative stochastic optimization:** This category includes methods such as the Cross-Entropy Method (CEM), Covariance Matrix Adaptation Evolution Strategy (CMA-ES), and others. These techniques operate similarly to random shooting but improve upon it by iteratively refining the action selection process to converge toward the optimal action. (See Chapter 12 for more details.)
- **Learning an approximate maximizer with a neural network:** In this approach, a neural network μ_{β} is trained to approximate the action that maximizes the Q-value, i.e., $\mu_{\beta} \approx \arg \max_a Q_{\psi}(s, a)$.

6.4.2 Improving DQN

The vanilla DQN algorithm, as we have seen in algorithm 9, is known to diverge very easily. There are two main problems. The first is that sequential states are highly correlated because some of the time steps are dependent on each other. The Q function would be updated many times for similar states while forgetting the Q values for other states (called catastrophic forgetting). In the DQN paper, they used a biologically inspired mechanism called experience replay (replay buffer). This approach stores past experiences and randomly samples mini-batches from the buffer during training, removing correlations in the observation sequence.

The second challenge arises because the targets (y_i) are computed using the same model which is being trained. With each update, the target values shift alongside the Q-function, making the training process unstable. The solution is to use a separate set of older weights to calculate the target values.

Algorithm 10 Deep-Q-Learning (DQN) with Replay-Buffer and Target Networks [MKS⁺13]

Input C, α, τ
 $D \leftarrow \{\}$, init $\omega, \omega' = \omega, t = 0$
 init. s
Loop for each episode:
 choose a from s using policy derived from $Q_\omega(s, a)$
 observe r, s'
 store transition (s, a, r, s') in replay buffer D
 sample random minibatch of tuples (s_t, a_t, r_t, s_{t+1}) from D
 for i in minibatch **do**
 $y_i = \begin{cases} r_i & s_{i+1} \text{ is terminal} \\ r_i + \gamma \max_{a'} Q_{\omega'}(s_{i+1}, a') & \text{otherwise} \end{cases}$
 Do gradient descent step on $(y_i - Q_\omega(s_i, a_i))^2$ for parameters
 $w : \Delta w = \alpha (y_i - Q_\omega(s_i, a_i)) \nabla_w Q_\omega(s_i, a_i)$
 endfor
 if $\text{mod}(t, C) == 0$ **then**
 $\omega' \leftarrow \tau \omega' + (1 - \tau) \omega$
 endif
end loop

6.4.3 Double Deep Q Learning

The modified DQN algorithm works very well, but there is still a problem with this approach. The issue is that the Q-values are overestimated, as depicted in Figure 8. The problem lies in the max/argmax operator used for calculating the target values.

$$y_i = r_i + \gamma \max_{a'} Q_{\omega'}(s_{i+1}, a') = r_i + \gamma Q_{\omega'}(s_{i+1}, \arg \max_{a'} Q_{\omega'}(s_{i+1}, a'))$$

The max operator is biased under noise, leading to overestimation. For random variables, this implies:

$$\mathbb{E}[\max(X_1, X_2, \dots)] \geq \max(\mathbb{E}[X_1], \mathbb{E}[X_2], \dots)$$

To understand this intuitively, imagine trying to determine the heaviest person in a group where everyone has the same true weight (e.g., 70kg), but the scale introduces random noise of $\pm 1\text{kg}$. If you take the maximum of these noisy measurements, it will almost always be greater than 70kg. Thus, the expectation of the maximum exceeds the maximum of the expectations, resulting in an upward bias. Since the Q-values we use are estimates and inherently noisy, this max-based target leads to systematic overestimation during training. The solution is to use two different networks for action selection and value computation, i.e. the current network ω and the target network ω' .

$$y_i = r + \gamma Q_\beta(s', \arg \max_{a'} Q_\alpha(s', a'))$$




Figure 8: Value estimates by DQN (orange) and Double DQN (blue) on some Atari games, from [vHGS15]

6.4.4 Dueling DQN

Recall that the advantage function is defined as $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$. The key idea behind the *Dueling DQN* architecture [WSH⁺16] is to separately estimate two components: the **state-value function** $V(s)$ and the **advantage function** $A(s, a)$. These two outputs are then combined to compute the action-value function $Q(s, a)$. This decoupling allows the network to independently learn the value of a given state and the relative advantages of actions within that state. An important detail is that when computing the advantages, we ensure they have zero mean across all actions in order to be compliant with the definition of the advantage function:

$$A(s, a) \leftarrow A(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a')$$

Overall, the Dueling DQN architecture leads to more stable and faster convergence, particularly in environments where many actions have similar Q-values.

6.5 Self-Test Questions

1. **What we mean by Monte Carlo estimates** $\rightarrow$ 6.1
2. **Why TD learning can be seen as a combination of MC and Dynamic Programming**
DP because we compute the value function iteratively starting from 0 for all states and MC because we sample an action in every iteration and use that action to compute the target y_t for time step t . This target is then incorporated by moving average
3. **What is the TD error ?** $V^\pi(s_t) = V^\pi(s_t) + \alpha \underbrace{([r_t + \gamma V^\pi(s_{t+1})] - V^\pi(s_t))}_{\text{TD error}}$
4. **How Q-Learning works** $\rightarrow$ 6.3.1
5. **What Value function approximation is**
We don't store the values in tables any more. Instead, we use parametrised function approximations of the V/Q-Function. Here we calculate the gradients of specific loss functions in order to update the parameters.
6. **Why Q-Learning is not actual gradient ascent?**
There is no clearly defined objective since the targets are always changing since we also approximate them ourself.
7. **How to fix Q-learning for deep neural networks?**
Use a replay buffer where we randomly select transitions to solve the problem that successive states are highly correlated. And use a separate target network to compute the targets to solve the problem that the targets would keep changing because they were computed by the same network that was supposed to learn the policy.
8. **Why do we get the over-estimation effect and how to fix it?**
Under noise the max operator is biased to be larger than it should be. And since our estimate of Q is a noisy estimate this leads to the overestimation effect. We fix it by using the double DQN method. There we use two separate networks. The first one selects the (what it thinks) best possible action and the second one calculates the target by using the suggested action.

6.6 Resources

Much of the content here on Deep Q-Learning is based on Sergey Levine's CS 285: Lecture 8 [LSM⁺20, RAI20].

7 Policy Gradients

In this section we change the approach to finding the optimal policy. The previous approaches were designed to approximate the value or action value function, which was then used to generate a policy. But there are some difficulties with this approach. In order to derive a policy from the value function, we would need the dynamics model. By learning the q-value we are able to derive a policy, but this also becomes more challenging to solve efficiently for high dimensional/continuous action spaces. With policy gradients we parametrise the policy π_θ directly, where the goal is to find the optimal parameter θ in order to maximise the value function.

The goal is to maximise the value function by choosing the right parameter θ for our policy.

$$V(s_0; \theta) = \mathbb{E}_{\pi_\theta} \left[\sum_{k=0}^T R(s_k, a_k) | s_0 \right]$$

We can now rewrite this in terms of trajectories $\tau = (s_0, a_0, \dots, s_T)$

$$V(s_0; \theta) = \int_{\tau} P(\tau; \theta) \underbrace{R(\tau)}_G \quad (9)$$

To find the optimal policy, we need to compute the gradient with respect to θ . The challenge is that the resulting gradient of (9) does not resemble an expectation that we can easily Monte Carlo sample. There are two methods that can be applied here to enable Monte Carlo integration: the likelihood-ratio gradient and the reparameterization trick. For now we will use the likelihood-ratio gradient, and later discuss scenarios where the reparameterization trick would be more appropriate (more on this can be found in [Vie19]). Taking the gradient with respect to θ :

$$\begin{aligned} \nabla_{\theta} V(s_0; \theta) &= \nabla_{\theta} \int_{\tau} P(\tau; \theta) R(\tau) \\ &= \int_{\tau} R(\tau) \nabla_{\theta} P(\tau; \theta) \\ &= \int_{\tau} R(\tau) \frac{P(\tau; \theta)}{P(\tau; \theta)} \nabla_{\theta} P(\tau; \theta) \quad \left(\text{log trick: } \nabla_{\theta} \log P(\tau; \theta) = \frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)} \right) \\ &= \int_{\tau} R(\tau) P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) \end{aligned}$$

The probability of a trajectory is defined as $P(\tau; \theta) = p_0(s_0) \prod_{t=0}^{T-1} p(s_{t+1}|s_t, a_t) \pi_{\theta}(a_t|s_t)$, which gives us the following equation

$$\begin{aligned} \nabla_{\theta} V(s_0; \theta) &= \int_{\tau} R(\tau) P(\tau; \theta) \nabla_{\theta} \log \left(p_0(s_0) \prod_{t=0}^{T-1} p(s_{t+1}|s_t, a_t) \pi_{\theta}(a_t|s_t) \right) \\ &= \int_{\tau} R(\tau) P(\tau; \theta) \left(\nabla_{\theta} \log p_0(s_0) + \sum_{t=0}^{T-1} \nabla_{\theta} \log p(s_{t+1}|s_t, a_t) + \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right) \\ &= \int_{\tau} R(\tau) P(\tau; \theta) \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[R(\tau) \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right] \quad (10) \end{aligned}$$

This is an expectation, which means that we can estimate it with a sample mean. If we collect a set of trajectories $\mathcal{D} = \{\tau_i\}_{i=1, \dots, N}$ where each trajectory is obtained by letting the agent act in the environment using the policy π_{θ} , the policy gradient can be estimated with

$$\frac{1}{|\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) R(\tau) \quad (11)$$

This has the nice property that we do not need a dynamic model. A side note, the term $\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)$ is called the score function and $\frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)}$ is called the likelihood ratio.

Depending on whether the action space is discrete or continuous, different policy gradient strategies are employed, each with a corresponding form of the score function.

- Softmax policy (for discrete):

$$\pi_{\theta}(a|s) = \frac{e^{\psi(s,a)^T \theta}}{\sum_a e^{\psi(s,a)^T \theta}} \rightarrow \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) = \psi(s, a) - \psi(s, a) \pi_{\theta}$$

- Gaussian (for continuous):

$$\pi_{\theta}(a|s) = N(a|f_{\theta}(s), \sigma^2) \rightarrow \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) = \frac{(a - f_{\theta}(s)) \nabla_{\theta} f(s)}{\sigma^2}$$

As formulated thus far this method is unbiased but very noisy (high variance) since we are estimating the expectation via our policy (see 11). The second issue is that since the expectation is defined over our parametrised policy, meaning the policy gradient is an on-policy method, we cannot reuse previously sampled data. This makes the approach sample inefficient, as the policy is updated at each step. However, there are potential solutions that we are going to look at in the following.

7.1 Use Temporal Structure

If we examine the equation 10, we see that taking a step with this gradient pushes up the log-probabilities of each action in proportion to $R(\tau)$, the sum of all rewards ever received. But this makes not so much sense, because the agent only should reinforce actions based on their consequences. Rewards received before an action should not determine how good that action was, only rewards received after. We can fix this by deriving the gradient estimator for a single reward.

$$\nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[r_t] = \mathbb{E}_{\pi_{\theta}} \left[r_t \sum_{k=0}^t \nabla_{\theta} \log \pi_{\theta}(a_k|s_k) \right]$$

Summing this formula gives

$$\begin{aligned} \nabla_{\theta} V(s_0; \theta) &= \nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[R] = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^T r_t \sum_{k=0}^t \nabla_{\theta} \log \pi_{\theta}(a_k|s_k) \right] \\ &= \mathbb{E}_{\pi_{\theta}} \left[\sum_{k=0}^T \nabla_{\theta} \log \pi_{\theta}(a_k|s_k) \sum_{t=k}^T r_t \right] \end{aligned} \quad (12)$$

This rearrangement shows that the log-probability of each action is scaled by the return following that action, rather than the total return. This aligns better with the principle that actions should be credited according to their future outcomes.

7.2 REINFORCE

A common algorithm that uses policy gradients is REINFORCE. (Note: Depending on the lecture, REINFORCE may sometimes be introduced without considering temporal structure, as seen in KIT. Please refer to your course materials for clarification. This definition aligns with that in Sutton and Barto [SB18].)

Algorithm 11 REINFORCE : Monte-Carlo Policy-gradient Control (episodic)

Input π_{θ}, α

Loop for each episode $\sim \pi_{\theta}$:

 Loop for each step of the episode $t = 0, 1, \dots, T - 1$

$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} r_k$

$\theta \leftarrow \theta + \alpha \gamma^t G \nabla_{\theta} \log \pi_{\theta}(a_t|s_t)$

7.3 Baseline

This is a further improvement on the vanilla policy gradient. Before discussing the improvement, we will introduce a lemma.

Lemma 1 (Grad-Log-Prob).

$$\mathbb{E}_{x \sim P_\theta} [\nabla_\theta \log P_\theta(x)] = 0 \quad (13)$$

The proof goes as follows

$$\begin{aligned} \overbrace{\int_x P_\theta(x) = 1}^{\text{property of a PDF}} &\rightarrow 0 = \nabla_\theta \int_x P_\theta(x) \\ &= \int_x \nabla_\theta P_\theta(x) \\ &= \int_x P_\theta(x) \nabla_\theta \log P_\theta(x) \\ &= \mathbb{E}_{x \sim P_\theta} [\nabla_\theta \log P_\theta(x)] \end{aligned}$$

A consequence of this is that we can subtract any function b which does not depend on the actions a while not adding any bias but are able to reduce the variance (last part is not shown here).

$$\begin{aligned} \mathbb{E}_{a \sim \pi_\theta} \left[\sum_{k=0}^T \nabla_\theta \log \pi_\theta(a_k | s_k) \left(\sum_{t=k}^T r_t - b(s_k) \right) \right] &= - \mathbb{E}_{a \sim \pi_\theta} \left[\sum_{k=0}^T \nabla_\theta \log \pi_\theta(a_k | s_k) b(s_k) \right] + \underbrace{\mathbb{E}_{a \sim \pi_\theta} \left[\sum_{k=0}^T \nabla_\theta \log \pi_\theta(a_k | s_k) \sum_{t=k}^T r_t \right]}_x \\ &= - \sum_{k=0}^T \mathbb{E}_{a \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_k | s_k) b(s_k)] + x \\ &= - \sum_{k=0}^T b(s_k) \underbrace{\mathbb{E}_{a \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_k | s_k)]}_0 + x \quad (\text{Lemma 1}) \\ &= x \\ &= \nabla_\theta V(s_0; \theta) \end{aligned}$$

The most common choice for the baseline is to use the value function $V_\omega^\pi(s)$, important here is that the value function does not depend on π_θ .

7.4 Data Reuse

As mentioned earlier, one limitation of the policy gradient approach is that it is inherently on-policy—the expectation is defined with respect to the current policy. This means that once the policy is updated, previously collected trajectories become outdated (off-policy) and can no longer be used for further training. As a result, new trajectories must be generated after every gradient step, which can be highly inefficient.

A common way to address this issue is through importance sampling. This technique allows us to reuse trajectories collected under an earlier policy by reweighting them appropriately. Suppose we want to evaluate the following expectation:

$$\mu = \mathbb{E}_{p(x)}[f(x)] = \int p(x) f(x) \, dx$$

one can then rewrite μ in terms of another PDF $q(x)$ as follows

$$\mu = \int p(x) f(x) \, dx = \int q(x) \frac{p(x)}{q(x)} f(x) \, dx = \mathbb{E}_{q(x)} \left[\frac{p(x)}{q(x)} f(x) \right]$$

But this can suffer from high variance if p and q are quite different. In expectation it is exact, but we would need to take a lot of samples to reduce the variance. Nevertheless, we can use this to compute our gradient while still using older trajectories:

$$\begin{aligned}
\nabla_{\theta} V(s_0; \theta) &= \nabla_{\theta} \mathbb{E}_{\pi_{\text{old}}} \left[\sum_{k=0}^T \frac{\pi_{\theta}(a_k | s_k)}{\pi_{\text{old}}(a_k | s_k)} \left(\sum_{t=k}^T r_t - b(s_t) \right) \right] \\
&= \mathbb{E}_{\pi_{\text{old}}} \left[\sum_{k=0}^T \nabla_{\theta} \frac{\pi_{\theta}(a_k | s_k)}{\pi_{\text{old}}(a_k | s_k)} \left(\sum_{t=k}^T r_t - b(s_t) \right) \right] \\
&\vdots \quad (\text{log trick}) \\
&= \mathbb{E}_{\pi_{\text{old}}} \left[\sum_{k=0}^T \frac{\pi_{\theta}(a_k | s_k)}{\pi_{\text{old}}(a_k | s_k)} \nabla_{\theta} \log \pi_{\theta}(a_k | s_k) \left(\sum_{t=k}^T r_t - b(s_t) \right) \right]
\end{aligned}$$

One remark is that, while we needed to use the likelihood ratio gradient for equation (9), since it would not have represented an expectation that we could sample using Monte Carlo. In the case of this expectation, this is actually not necessary. Even when we take the gradient of the expectation (first equation) with respect to θ , it remains an expectation (because the expectation is defined over π_{old}). As a result, many libraries do not implement it using the likelihood ratio form (last equation).

7.5 Policy Gradients with Q-Values

So far, we have used a single-sample estimate for calculating the return, as shown in equation (12). While the gradient estimator is unbiased, it suffers from high variance since we rely on a single trajectory to estimate the return for a given state-action pair. A better approach would be to use the expected sum of returns for the state-action pair, which is precisely the definition of the Q-value function. Therefore, we can also express the policy gradient via the Q-value which results in a lower-variance estimates.

Theorem 1 (Policy gradient theorem). *For any differentiable policy $\pi_{\theta}(s, a)$, let p_0 be the starting distribution over the states in which we begin an episode. Then, the policy gradient of $\mathbf{J}(\theta) = \mathbb{E}_{\pi_{\theta}}[G_0 | s_0 \sim p_0]$ is*

$$\nabla_{\theta} \mathbf{J}(\theta) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^T \nabla \log \pi_{\theta}(a_t | s_t) \gamma^t Q^{\pi_{\theta}}(s_t, a_t) \mid s_0 \sim p_0 \right]$$

The proof will not be shown here but can be looked up [here](#)

7.6 Advantage Estimation

We have seen that we can express the policy gradients through the Q-value. Additionally, we know that we can subtract any baseline, as long as it does not depend on the actions, such as the value function V^{π} . Naturally, this leads to the conclusion that we can also express the policy gradient using the advantage function.

$$\mathbb{E}_{a \sim \pi_{\theta}} \left[\sum_{k=0}^T \nabla_{\theta} \log \pi_{\theta}(a_k | s_k) (Q^{\pi}(s_k, a_k) - V^{\pi}(s_k)) \right] = \mathbb{E}_{a \sim \pi_{\theta}} \left[\sum_{k=0}^T \nabla_{\theta} \log \pi_{\theta}(a_k | s_k) A^{\pi}(s_k, a_k) \right]$$

The question now is, which function should we fit/learn in order to calculate A^{π} ? The simplest option would be to learn/fit the value function V^{π} , since it only depends on the states as inputs, unlike the Q-function or advantage function, which also depend on actions. From equation (??), we know that the Q-value can be expressed through the value function. However, since calculating the full expectation would be cumbersome, we approximate it instead:

$$\begin{aligned}
Q^{\pi}(s_k, a_k) &= r(s_k, a_k) + \mathbb{E}_{s_{k+1} \sim p(s_{k+1} | s_k, a_k)} [V^{\pi}(s_{k+1})] \approx r(s_k, a_k) + V^{\pi}(s_{k+1}) \\
A^{\pi}(s_k, a_k) &\approx r(s_k, a_k) + V^{\pi}(s_{k+1}) - V^{\pi}(s_k)
\end{aligned} \tag{14}$$

7.6.1 Generalized Advantage Estimation

With equation (14), we achieve lower variance but introduce higher bias, as our function approximation of the value function could be incorrect. In contrast, equation (12) has no bias but suffers from high variance. As we saw with n-step returns in Section ??, we can balance variance and bias by rolling out the estimates, resulting in:

$$Q_n^\pi(s_k, a_k) = \sum_{t=k}^{k+n} \gamma^{t-k} r(s_t, a_t) + \gamma^n V^\pi(s_{k+n})$$

Generalized Advantage Estimation (GAE) takes this a step further by considering all possible n-steps and calculating an exponentially-weighted average of n-step returns with decay parameter λ :

$$Q^\pi(s_k, a_k; \lambda) = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} Q_n^\pi(s_k, a_k) \quad \text{with } (1 - \lambda) \sum_{n=1}^{\infty} \lambda^n = 1$$

The idea is to assign exponentially smaller weights to future time steps, as they have more variance. This method, which involves function fitting combined with policy gradients, is known as actor-critic methods, and will be discussed in detail the following sections.

7.7 Self-Test Questions

1. **What are the advantages/disadvantages of policy search vs value-based methods**

Policy gradients are easier to use and tune, more compatible with rich architectures, more versatile, almost no bias $\rightarrow$ finds good solutions. But it needs much more samples. Value-based Methods are much more sample-efficient (allows off-policy) but No conv. guarantees, often hard to tune, hard to use cont. actions space and Aprox. errors in Q-function can bias the quality of the resulting policy.

2. **What is the main idea of policy gradient algorithms?**

With policy gradients we parametrise the policy π_θ directly, where the goal is to find the optimal parameter θ in order to maximise the value function.

3. **What kind of policies are used in discrete action and continuous action domains?**

Softmax policy for discrete and Gaussian policy for continuous.

4. **How can we use the log-ratio trick to compute the policy gradient?** $\rightarrow$ 7

5. **Why can we compute gradients even if the reward or the dynamics are not differentiable?**

After using the log-ratio trick we only have to compute the gradient $\nabla_\theta \log P(\tau; \theta)$ where $P(\tau; \theta) = p_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$ and since only our policy is dependent on θ everything else has zero gradient hence we do not need to have a dynamics model. And the reward is not dependent on θ .

6. **Explain the intuition of the REINFORCE update equation**

We want to maximise the value function parametrised by θ . To find the maximum, we take the derivative with respect to θ and take a step in the direction. As this gradient requires an expectation to be calculated, we estimate it using Monte Carlo simulation.

7. **Why do we need a baseline in policy gradient algorithms?**

Estimating the gradient via samples leads to high variance. Subtraction the baseline from the reward does not add any bias but can be chosen such that we reduce the variance.

8. **Why is optimizing the advantage beneficial to optimizing the Q-function?**

Advantage is given by $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$ It tells us how much action a is better than the expected value of the policy π in state s . The advantage function helps the agent focus on which actions lead to the most improvement over the average. This relative comparison helps in complex environments where absolute values can be hard to estimate.

9. **How can we exploit temporal structure for policy gradients?**

We can ignore rewards before time step t for the action a_t since they do not determine how good that action was

7.8 Resources

Much of the content here is based on Sergey Levine's CS 285: Lecture 5 [[LSM⁺20](#), [RAI20](#)] and also on the blog post "Part 3: Intro to Policy Optimization" from OpenAI Spinning Up [[Ach18](#)].

8 Actor Critic

So far, we've explored two methods for deriving optimal policies. The first approach involves value-based methods, such as dynamic programming and Q-learning, where we focus on learning value functions (V-, Q- or Advantage-Function) and then use them to derive the policy. The second approach is through policy gradient methods, where we directly learn the policy without the need for an explicit value function. Actor-Critic methods combine both of these approaches. The critic learns a value function, while the actor learns the policy. Instead of directly using the critic or the actor to derive the policy, the critic's value estimates are used to guide the actor, providing direction for updating the policy gradient.




Figure 9: Illustrating the relationship between Actor-Critic, policy-based and value-based methods

The most general actor-critic algorithm would look something like this:

Algorithm 12 One-Step Actor Critic

```

Init.  $s, \theta, w$ 
for  $t = 0, 1, \dots, T$  do
    sample  $a_t \sim \pi_\theta(s_t)$ 
    sample  $s_{t+1} \sim P(s_{t+1}|s_t, a_t)$  and  $r_{t+1} \sim R(s_t, a_t)$ 
    Update  $V_w^\pi$ 
     $A^\pi(s_t, a_t) = r_{t+1} + \gamma V_w^\pi(s_{t+1}) - V_w^\pi(s_t)$ 
     $\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi(a_t|s_t) A^\pi(s_t, a_t)$ 
     $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$ 
end for

```

There are multiple variants of the actor-critics method which will be looking at in the following.

9 Off Policy Actor Critic

The first actor-critic algorithm we examined, presented in Algorithm 12, is a fully on-policy method. This approach has significant drawbacks, as we discussed in the policy gradient chapter. Training samples are collected based on the current policy—the very policy we aim to optimize. As a result, after each gradient update, we can no longer use the previously collected samples. If the method were off-policy, we could achieve better sample efficiency and enhanced exploration by following a variety of different policies. This concept will be further explored in the following section.

9.1 First Attempt

The first idea that comes to mind for making the algorithm off-policy is to introduce a replay buffer, similar to what we did in the DQN section, resulting in an approach like this:

Algorithm 13

```

Init.  $s, \theta, w, D = \{\}$ 
for  $t = 0, 1, \dots, T$  do
    sample  $a_t \sim \pi_\theta(s_t)$ 
    sample  $s_{t+1} \sim P(s_{t+1}|s_t, a_t)$  and  $r_{t+1} \sim R(s_t, a_t)$ 
    store  $(s_t, a_t, s_{t+1}, r_{t+1})$  in  $D$ 
    sample a batch  $((s_i, a_i, s_{i+1}, r_{i+1}))$  from buffer  $D$ 
    execute Actor Critic Algorithm 12 on batch
end for

```

However, the algorithm has some significant flaws and cannot be used as it is currently written. There are two main issues.

Value Function The first concerns the target values

$$y_i = r_{i+1} + \gamma V_w^\pi(s_{i+1})$$

These targets are incorrect because the batch samples come from older actors, yet we use them to compute target values for training the current critic. As a result, the current actor is being guided by targets that reflect outdated policies, rather than its own behaviour.

To address this issue, we will learn the Q-value function instead of the V-function. Recall that the Q-function represents the expected return starting from a given state, taking a specific action, and then following the policy thereafter. Using the relationship between the value function and the Q-function given in Equation (3), we can rewrite the target as:

$$y_i = r_{i+1} + \gamma V_w^\pi(s_{i+1}) = r_{i+1} + \gamma Q_w^\pi(s_{i+1}, a') \quad a' \sim \pi_\theta(\cdot|s_i)$$

This allows us to reuse s_{t+1} from the replay buffer, while sampling the corresponding action a' from the current policy. In doing so, our target values remain aligned with the actor's present behaviour. If we had insisted on using the value function in an off-policy setting, we would have needed to invoke the simulator to generate the actions the actor would take and observe the resulting next states—an expensive and impractical process. In contrast, with this Q-function approach, we only need to sample a' from the current policy, which is straightforward since the policy is represented by a neural network. Thus, the procedure remains both accurate and efficient.

Policy Gradient The second problem is that the policy gradient is defined as an expectation over the current policy, as seen in Equation 12. This implies that we cannot directly use the samples from the batch. We've previously discussed how to mitigate this using importance sampling, but another option is to simply sample again from the current actor, resulting in:

$$\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi(a_i^\pi | s_i) A^\pi(s_i, a_i^\pi) \quad a_i^\pi \sim \pi_\theta(\cdot | s_i) (\text{current actor})$$

While some parts of what we have looked at here may seem a bit incoherent/random, I believe this explanation helps to understand state of the art off-policy actor-critic methods. These methods are quite similar to what we’ve discussed here, so by the time we look at them next, it should feel more familiar when reading the code for these algorithms. That’s the goal, at least! As a reminder, policy gradients don’t have to be defined solely through the advantage function; they can also be defined directly using the Q-function removing the need to also learn a V-function.

9.2 Updating the critic

There are multiple ways to update the critic (see Section 6 for details) but in most cases, the update resembles the approach used in DQN.

In contrast, there are several different strategies for updating the actor, which we will explore next.

9.3 Deterministic updates

In methods described above, the policy is always modelled as a stochastic policy $\pi(a|s)$. Deterministic policy gradient (DPG) instead models the policy as a deterministic decision $a = \pi(s)$. With this policy the objective of the actor is given by

$$\max_{\theta} \mathbb{E}_{s \sim \mathcal{D}} [Q_{\phi}(s, \pi_{\theta}(s))]$$

The gradient of this with respect to θ is given by

$$\begin{aligned} \nabla_{\theta} \mathbb{E}_{s \sim \mathcal{D}} [Q_{\phi}(s, \pi_{\theta}(s))] &= \mathbb{E}_{s \sim \mathcal{D}} [\nabla_{\theta} Q_{\phi}(s, \pi_{\theta}(s))] \\ &= \mathbb{E}_{s \sim \mathcal{D}} [\nabla_a Q_{\phi}(s, a) \nabla_{\theta} \pi_{\theta}(s)|_{a=\pi_{\theta}(s)}] \end{aligned}$$

At first glance, this may not seem like the policy gradient we derived in the policy gradient chapter. However, the deterministic policy gradient is actually a special case of the stochastic policy gradient. The proof can be found in [SLH⁺14](Section 3.3).

9.3.1 Deep Deterministic Policy Gradient (DDPG)

DDPG [LHP⁺19] is essentially a combination of the deterministic policy gradient method and DQN, with DQN being used to train the critic. It makes use of the replay buffer and employs the technique of using an additional target network for more stable training. The calculation of the max in the target values is addressed by using a policy network to compute an action that approximately maximizes the Q-function. Since the policy is deterministic, exploration is limited. To encourage better exploration, noise is added to the deterministic policy. (see algorithm 14)

But unfortunately DDPG is quite unstable and requires carefully chosen hyper-parameters. Some solutions to this are given with the SAC and TD3 algorithms which we will be looking at in the following.

9.3.2 Twin Delayed Deep Deterministic (TD3)

One of the problems with DDPG is that it suffers from overestimation in the learned Q-Function. This then leads to the policy breaking, because it exploits the errors in the Q-function. Twin Delayed DDPG (TD3) [FvHM18] is an algorithm (see 15) that addresses this issue by introducing three changes to the original DDPG algorithm.

1. **Double-Q Learning:** learns two Q-functions instead of one (hence “twin”), and uses the smaller of the two Q-values to form the targets in the Bellman error loss functions. Meaning we have critics (Q_{ϕ_1}, Q_{ϕ_2})
2. **Delayed Policy Updates:** update the policy (and target networks) less frequently than the Q-function.
3. **Target Policy Smoothing:** add noise to the target action, to make it harder for the policy to exploit Q-function errors by smoothing out Q along changes in action.

Algorithm 14 DDPG algorithm from [LHP⁺19]

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .
Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$
Initialize replay buffer R
for episode = 1, M **do**
 Initialize a random process $\mathcal{N}$ for action exploration
 Receive initial observation state s_1
 for t = 1, T **do**
 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise
 Execute action a_t and observe reward r_t and observe new state s_{t+1}
 Store transition (s_t, a_t, r_t, s_{t+1}) in R
 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R
 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$
 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$
 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\begin{aligned}\theta^{Q'} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} &\leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}\end{aligned}$$

end for
end for

Algorithm 15 TD3 algorithm from [FvHM18]

Initialize critic networks $Q_{\theta_1}, Q_{\theta_2}$, and actor network π_ϕ with random parameters θ_1, θ_2, ϕ
Initialize target networks $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$
Initialize replay buffer $\mathcal{B}$
for $t = 1$ **to** T **do**
 Select action with exploration noise $a \sim \pi_\phi(s) + \epsilon$,
 $\epsilon \sim \mathcal{N}(0, \sigma)$ and observe reward r and new state s'
 Store transition tuple (s, a, r, s') in $\mathcal{B}$

 Sample mini-batch of N transitions (s, a, r, s') from $\mathcal{B}$
 $\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon$, $\epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$ {Target policy smoothing}
 $y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$ {Clipped Double Q-Learning}
 Update critics $\theta_i \leftarrow \arg \min_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$
 if $t \bmod d$ {Delayed update}
 then
 Update ϕ by the deterministic policy gradient:
 $\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$
 Update target networks:
 $\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$
 $\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$
 end if
end for

The comments and underlines added in the following algorithm are not part of the original paper. They have been included solely to highlight the areas where the new changes have been made, inspired by [Wen18c].

9.4 Variational Actor Updates (Soft Actor Critic (SAC))

A problem with the deterministic policy approaches is that they might not explore enough leading to sub optimal policies. SAC [HZAL18] tries to balance exploration and exploitation by using max-entropy reinforcement learning. It is an algorithm that optimizes a stochastic policy in an off-policy way, forming a bridge between stochastic policy optimization and DDPG-style approaches.

9.4.1 Max-Entropy Reinforcement Learning

To remind ourself, the entropy of a random variable X roughly speaking says how random X is and is defined as

$$H(X) = \mathbb{E}_{x \sim X} [-\log p(x)]$$

In the context of entropy-regularised reinforcement learning, the agent is provided with not only the usual reward, but also an entropy term, which changes the problem of finding the optimal policy to:

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha H(\pi|s_t)) \right]$$

Here α is the **trade-off coefficient** and controls if the focus should be on maximizing reward or keeping entropy, i.e. a high α lets the agent focus more on exploration while a small α leads to more exploitation. In the course of this, we can also define the corresponding Value/Action-Value-Function

$$\begin{aligned} V_{\text{soft}}^{\pi}(s) &= \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (r_t + \alpha H(\pi|s_t)) \mid s_0 = s \right] \\ Q_{\text{soft}}^{\pi}(s, a) &= \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t r_t + \alpha \sum_{t=1}^{\infty} \gamma^t H(\pi|s_t) \mid s_0 = s, a_0 = a \right] \\ &= \mathbb{E}_{\substack{s' \sim P \\ a' \sim \pi}} [R(s, a) + \gamma (Q_{\text{soft}}^{\pi}(s', a') + \alpha H(\pi(\cdot, s')))] \end{aligned}$$

We can then define the functions in terms of each other:

$$V_{\text{soft}}^{\pi}(s) = \mathbb{E}_{a \sim \pi} [Q_{\text{soft}}^{\pi}(s, a)] + \alpha H(\pi(\cdot|s)) \quad (15) \quad Q_{\text{soft}}^{\pi}(s, a) = \mathbb{E}_{s' \sim P} [R(s, a) + \gamma V_{\text{soft}}^{\pi}(s')] \quad (16)$$

This V- and Q-Functions are also called the soft V- and soft Q-Function since the entropy regularization prevents that the policy becomes a "hard" maximum operator and always incorporate some exploration.

9.4.2 Actor and Critic

SAC sets up the MSBE loss for the critic similar to TD3 by using the clipped double-Q trick with additional target networks, and taking the minimum Q-value between the two Q approximators. Putting it all together, the loss functions for the Q-networks in SAC are:

$$\begin{aligned} L(\phi_i, \mathcal{D}) &= \mathbb{E}_{(s, a, r, s') \sim \mathcal{D}} \left[\left(Q_{\phi_i}(s, a) - y(r, s') \right)^2 \right], \text{ with} \\ y(r, s') &= r + \gamma \left(\min_{j=1,2} Q_{\phi_{\text{target},j}}(s', \tilde{a}') - \alpha \log \pi_{\theta}(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_{\theta}(\cdot|s'). \end{aligned}$$

In order to learn the policy the actor tries to maximize (15) which results into

$$V^{\pi}(s) = \mathbb{E}_{a \sim \pi_{\theta}} [\min_{j=1,2} Q_{\phi_j}(s, a) - \alpha \log \pi_{\theta}(a|s)]$$

In order to take the derivative with respect to θ , one could use the likelihood ratio gradient, as we have seen in the policy gradient chapter. However, this approach does not leverage the gradient information from the critic. Although the action fed into the network is sampled from the policy π_{θ} , the sampling process itself is non-differentiable, preventing the direct application of the chain rule to propagate the gradients. To solve this we are going to use the reparametrization trick.

9.4.3 Reparametrization Trick

Consider an objective of the form:

$$\arg \min_{\theta} \mathbb{E}_{p_{\theta}}[f(x)] = \int p_{\theta}(x) f(x) dx$$

In the context of a computational graph, the challenge arises from the fact that $x \sim p_{\theta}$ is a random/stochastic node, through which backpropagation cannot be performed.

The idea is to introduce a random variable $\xi \sim q(\xi)$ and a deterministic, differentiable function $g_{\theta}(\xi)$ such that $x = g_{\theta}(\xi)$, while still having $x \sim p_{\theta}(x)$. This allows us to rewrite the original expectation as:

$$\mathbb{E}_{x \sim p_{\theta}}[f(x)] = \mathbb{E}_{\xi \sim q}[f(g_{\theta}(\xi))]$$

The advantage of this formulation is that when we take the derivative, we are left with an expectation that can now be Monte Carlo sampled. Specifically, we have:

$$\begin{aligned} \nabla_{\theta} \mathbb{E}_{\xi \sim q}[f(g_{\theta}(\xi))] &= \nabla_{\theta} \int q(\xi) f(g_{\theta}(\xi)) d\xi \\ &= \int q(\xi) \nabla_{\theta} f(g_{\theta}(\xi)) d\xi \\ &= \int q(\xi) \frac{df}{dg_{\theta}(\xi)}(g_{\theta}(\xi)) \frac{dg_{\theta}}{d\theta}(\xi) d\xi \\ &= \mathbb{E}_{\xi \sim q} \left[\frac{df}{dg_{\theta}(\xi)}(g_{\theta}(\xi)) \frac{dg_{\theta}}{d\theta}(\xi) \right] \end{aligned}$$

In SAC for continuous action spaces, the policy is Gaussian distributed, i.e., $\pi_{\theta}(a|s) \sim \mathcal{N}(\mu, \Sigma)$. Here, the choice of $q(\xi)$ and $g_{\theta}(\xi)$ is given by:

$$g_{\theta}(\xi) = \mu + A\xi \quad \text{s.t.} \quad A^T A = \Sigma$$

After applying the reparameterization trick, we are left with the following objective for the actor

$$\max_{\theta} \mathbb{E}_{\substack{s \sim \mathcal{D} \\ \xi \sim \mathcal{N}}} \left[\min_{j=1,2} Q_{\phi_j}(s, \tilde{a}_{\theta}(s, \xi)) - \alpha \log \pi_{\theta}(\tilde{a}_{\theta}(s, \xi)|s) \right]$$

Unlike in TD3, which uses Q_{ϕ_1} (just the first Q approximator), SAC uses $\min_{j=1,2} Q_{\phi_j}$ (the minimum of the two Q approximators).




Figure 10: Computational graphs for the objective without the reparameterization trick (left) and with it (right). By isolating the stochasticity into a separate node, we enable gradient computation with respect to the model parameters.

Ultimately, we are left with the following SAC algorithm:

Algorithm 16 Soft Actor-Critic from [HZAL18] (original)

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.
for each iteration **do**
 for each environment step **do**
 $a_t \sim \pi_\phi(a_t|s_t)$
 $s_{t+1}p \sim p(s_{t+1}|s_t, a_t)$
 $\mathcal{D} \leftarrow \mathcal{D} \cup \{(s_t, a_t, r(s_t, a_t), s_{t+1})\}$
 end for
 for each gradient step **do**
 $\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$
 $\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i)$ for $i \in \{1, 2\}$
 $\phi \leftarrow \phi - \lambda_\phi \hat{\nabla}_\phi J_\phi(\phi)$
 $\bar{\psi} \leftarrow \tau\psi + (1 - \tau)\bar{\psi}$
 end for
end for

Algorithm 17 Soft Actor-Critic (OpenAI’s Spinning Up version [Ach18])

Input: initial policy parameters θ , Q-function parameters ϕ_1, ϕ_2 , empty replay buffer $\mathcal{D}$
Set target parameters equal to main parameters $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
repeat
 Observe state s and select action $a \sim \pi_\theta(\cdot|s)$
 Execute a in the environment
 Observe next state s' , reward r , and done signal d to indicate whether s' is terminal
 Store (s, a, r, s', d) in replay buffer $\mathcal{D}$
 If s' is terminal, reset environment state.
 if it’s time to update **then**
 for j in range(however many updates) **do**
 Randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from $\mathcal{D}$
 Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left(\min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$

where $\tilde{a}_\theta(s)$ is a sample from $\pi_\theta(\cdot|s)$ which is differentiable wrt θ via the reparametrization trick.

Update target networks with

$$\phi_{\text{targ},i} \leftarrow \rho\phi_{\text{targ},i} + (1 - \rho)\phi_i \quad \text{for } i = 1, 2$$

end for
 end if
until convergence

9.5 CrossQ

CrossQ [BPB⁺24] attempts to introduce batch normalization into Q-learning. But what exactly is batch normalization?

Batch normalization is a technique used in deep learning to improve the training of neural networks by normalizing the inputs to each layer. It works by adjusting and scaling the activations (outputs of neurons). Specifically, it standardizes these activations by subtracting the batch mean and dividing by the batch standard deviation, followed by scaling and shifting using learnable parameters. Batch normalisation has many benefits such as:

- **Faster Training:** It helps the network converge faster by reducing internal covariate shift (i.e., the changing distribution of inputs to each layer during training), making the process more stable.
- **Improved Performance:** It often leads to better generalization, reducing overfitting, and improving model performance on unseen data.
- **Allows Higher Learning Rates:** Batch normalization enables the use of higher learning rates without causing instability in the training process.

However, several people have found that standard batch normalization doesn't work well with Q-learning and often makes the learning process more unstable. The issue lies in how the Temporal Difference (TD) error is computed in reinforcement learning. So in general the loss is defined as

$$\begin{aligned} L_\psi &= \left(\max_{a'} Q_{\psi'}(s_{t+1}, a') + r_t - Q_\psi(s_t, a_t) \right)^2 \\ &= (Q_{\psi'}(s_{t+1}, \mu_\psi(s)) + r_t - Q_\psi(s_t, a_t))^2 \end{aligned}$$

where (s_t, a_t, r_t, s_{t+1}) are samples from the replay buffer. While s_t and s_{t+1} come from the same distribution, a_t is from the replay buffer, and a' is the action predicted by the neural network (which corresponds to the argmax action). This means that a_t and a' are distributed differently. Since batch normalization requires samples to be independent and identically distributed (iid), this discrepancy in distributions introduces the problem. The solution to this issue is to concatenate the different state and action batches into a single batch, ensuring that all inputs come from the same mixture distribution.

$$\begin{aligned} Q_\psi \left(\begin{bmatrix} s_t \\ s_{t+1} \end{bmatrix}, \begin{bmatrix} a_t \\ a' \end{bmatrix} \right) &= \begin{bmatrix} q_t \\ q_{t+1} \end{bmatrix} \\ L_\psi &= (\text{StopGradient}(q_{t+1}) + r_t - q_t)^2 \end{aligned}$$

This procedure makes training faster and more efficient overall, as it eliminates the need for a target network.

9.6 BRO

The BRO (Bigger, Regularized, Optimistic) paper [NOJ⁺24b] shows that scaling model capacity—alongside strong regularization and optimistic exploration—can significantly boost sample efficiency in reinforcement learning. Unlike traditional approaches that rely mainly on algorithmic tweaks, BRO leverages large critics and task-specific enhancements to achieve state-of-the-art performances.

9.7 Self-Test Questions

1. What is the role of the critic in actor-critic methods?

The critic calculates the TD-Error and tries to learn the value/action-value function

2. Which methods do you know for computing the critic?

Single sample estimates, N-step returns and Generalized advantage estimation

3. Why are off-policy methods more efficient than on-policy methods

The can use previous collected data to do updates and do not always have to generate new data/discard old data. It also promotes exploring if the samples are chosen randomly from the buffer.

4. **How can we achieve off-policy RL with continuous actions**

Computing the maximum over actions in the target is a challenge in continuous action spaces. Some solutions are

- finding the max by doing gradient descent on the Q-network
- random shooting: randomly select some actions and take the best performing out of them for calculating the target.
- iterative stochastic optimization: same as random shooting but we iteratively update the distribution from which we draw samples to get a better action at the end
- learning a neural network to predict the max action

5. **What are the different options to optimize the actor?**

- stochastic: normal policy gradient where the reward is now changed to be the advantage function
- deterministic: learn a deterministic policy that approximates the max operator
- variational: regularize RL objective with entropy

6. **Why do we have a bias in actor critic algorithms and how to fix it?**

The problem is that we use approximators to calculate the targets and since they are bias actor-critics are also bias. That approximator will usually be randomly initialized so it will not give a true estimation of the return, it will be biased towards some random value that was initialized with.

7. **What type of gradient is DDPG using and why can we not apply the same gradient for the sampled return?**

DDPG uses $\nabla_{\theta} \mathbb{E}_{s \sim \mathcal{D}} [Q_{\phi}(s, \pi_{\theta}(s))] = \mathbb{E}_{s \sim \mathcal{D}} [\nabla_a Q_{\phi}(s, a) \nabla_{\theta} \pi_{\theta}(s) |_{a=\pi_{\theta}(s)}]$

8. **What is the objective of max-ent reinforcement learning and why is that useful?**

In entropy-regularized reinforcement learning, the agent gets a bonus reward at each time step proportional to the entropy of the policy at that time step. The new objective is to maximise this reward. The entropy maximization leads to policies that can explore more.

9. **How to we obtain the policy update in SAC?** We take the derivative of the Value-Function (Q-Value + Entropy) but with the help of the reparametrization trick in order to make the gradient flow through the critic.

10. **What is the reparametrization trick and when should it be preferred to the likelihood policy gradients?** It is a trick, in which a sample from $\pi_{\theta}(\cdot|s)$ is drawn by computing a deterministic function of state, policy parameters, and independent noise. The reparametrization trick makes it possible to use the knowledge of the gradient of the Q-function approximation to calculate the gradient of the objective.

9.8 Resources

A helpful introduction to actor-critic methods, including the off-policy variants, is provided in Sergey Levine's CS 285: Lecture 6 [LSM⁺20, RAI20]. Additionally, the blog posts from OpenAI's Spinning Up series on DDPG, TD3, and SAC are excellent resources [Ach18]. For CrossQ, the video [LA24] was particularly informative. For BRO, the official website [NOJ⁺24a] offers well-written explanations for those seeking more detailed information.

10 Control as Inference

This is a brief exploration into the question of what constitutes optimal behaviour. Traditionally, we've assumed that the behaviour we want to imitate is perfectly optimal. However, this assumption is not realistic since humans rarely act in strictly optimal ways.

In this section, we explore Control as Inference, a paradigm that reimagines reinforcement learning not as a strict reward-maximisation problem, but as an inference problem within a probabilistic framework. This approach allows us to model and tolerate sub-optimal behaviour. We begin with an introduction to variational inference which forms the foundation for this reformulation

10.1 Variational Inference

A central challenge in modern statistics, especially Bayesian statistics, is approximating complex, intractable probability distributions, such as posterior densities like

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int_z p(x, z') dz'}$$

where z represents the latent variables and x represents the observed data. The marginalization over z to calculate $p(x)$ in the denominator is typically intractable, because, for example, the search space of z is too large. Traditional methods like Markov Chain Monte Carlo offer accurate sampling but are often computationally expensive, particularly for large datasets or complex models.

Variational Inference (VI) provides a faster alternative by reframing inference as an optimization problem. Instead of sampling, VI approximates the true posterior by selecting the closest distribution $q_\theta(z)$ (from a simpler, predefined family) using a dissimilarity measure such as Kullback-Leibler (KL) divergence in order to find the best parameters θ .

In the following, we will explore how the most common discrepancy measure, the Kullback-Leibler divergence, can be used for inference. Additionally, in the upcoming discussion, $p^*(x)$ will represent the true probability distribution, while $q(x)$ will denote the approximation.

10.1.1 Minimizing Reverse KL

Minimizing the reverse KL divergence is often referred to as (I)nformation-projection and is defined as:

$$\text{KL}(q(x)||p^*(x)) = \int_x q(x) \log \frac{q(x)}{p^*(x)} dx$$

This objective is also known for its zero-forcing behaviour. This term arises from how the reverse KL divergence behaves in regions where the true distribution $p^*(x)$ has high probability mass, but the approximate distribution $q(x)$ assigns little or no mass. Consider two cases:

- If $p^*(x) > 0$ but $q(x) \approx 0$, then the ratio $\frac{q(x)}{p^*(x)}$ inside the logarithm becomes very small. However, since $q(x) \log(\cdot)$ is multiplied by a value close to zero, the overall contribution to the KL divergence is negligible.
- If $q(x) > 0$ and $p^*(x) = 0$, then the ratio becomes infinite, and $\log(\infty) = \infty$. This causes the KL divergence to become infinitely large, leading the optimization to avoid placing mass in such regions.

Therefore, minimizing the reverse KL forces $q(x)$ to avoid assigning probability mass where $p^*(x) = 0$. This behavior leads to mode-seeking: reverse KL tends to concentrate $q(x)$ on one mode of $p^*(x)$, ignoring regions of low probability. And as long as we are able to sample from $q(x)$, we can approximate the KL divergence using Monte Carlo integration.

10.1.2 Minimizing Forward KL

Minimizing the forward KL divergence is often referred to as (M)oment-projection and is defined as:

$$\text{KL}(p^*(x)||q(x)) = \int_x p^*(x) \log \frac{p^*(x)}{q(x)} dx$$

This objective exhibits probability forcing: if the true distribution $p^*(x)$ is high in a region, then the approximation $q(x)$ must also be high there to avoid a large penalty. So it got the following properties:

- It tends to average over multiple modes of $p^*(x)$, even when the approximation family cannot represent such multimodality well.
- It requires samples from the true distribution $p^*(x)$, making it generally unsuitable for variational inference

$$\int_x p^*(x) \log \frac{p^*(x)}{q(x)} dx = \mathbb{E}_{x \sim p^*(x)} \left[\frac{p^*(x)}{q(x)} \right]$$

- Minimizing the forward KL is equivalent to maximum likelihood estimation (MLE).

$$\int_x p^*(x) \log \frac{p^*(x)}{q(x)} dx = \text{const} - \int_x p^*(x) \log q(x) dx$$





Figure 11: Illustrating forward vs. reverse KL. The blue contours represent the true distribution p^* , and the red contours are from the approximation q . (a) Minimizing forward KL. (b–c) Minimizing reverse KL, where q fits onto different modes each time.

10.2 Connection between Variational Inference and Control

As mentioned earlier, our current understanding does not allow us to account for suboptimal behaviour in our algorithms. Consider, for example, a task where we are required to move a dot on a screen to a goal using a joystick. The optimal behaviour in this scenario would be to take the straight path, but small deviations may occur due to factors such as lack of concentration or external distractions. Despite these deviations, we still achieve the same outcome as if we had followed the optimal path. Intuitively, we recognize that small errors do not necessarily prevent the task from being successfully completed. Therefore, it is clear that some mistakes are more consequential than others when completing a task while still achieving the overall goal.

To incorporate suboptimality into our model, we begin by adding the concept of optimality to the general structure of a MDP. We introduce an optimality variable o_t , where $o_t = 1$ indicates that time step t is optimal, and $o_t = 0$ indicates that it is not. The corresponding graphical model is illustrated as follows



This model resembles a Hidden Markov Model. We now define some probabilities, which may initially appear arbitrary. However, they will lead to an elegant framework. For now, we take them as

given:

$$\begin{aligned}
p(\tau) &= p(s_1, a_1, \dots, s_T) = p(s_1) \prod_{t=1}^{T-1} p(s_{t+1} \mid s_t, a_t) \underbrace{\pi(a_t \mid s_t)}_{1 \text{ (uniform)}} \\
p(o_t \mid s_t, a_t) &\propto \exp(r(s_t, a_t)) \\
p(o_{1:T} \mid \tau) &= \exp\left(\sum_{t=1}^T r(s_t, a_t)\right)
\end{aligned}$$

Our goal is to infer the most likely trajectory given that the agent acts optimally. This means we want to compute the posterior:

$$p(\tau \mid o_{1:T} = 1) = \frac{p(\tau) p(o_{1:T} = 1 \mid \tau)}{p(o_{1:T} = 1)} \quad (17)$$

Once we have this posterior, we can use marginalization and conditioning to compute the optimal policy:

$$p(a_t \mid s_t, o_{1:T} = 1)$$

The distribution (17) can be challenging to compute directly, so we turn to variational inference. Specifically, we minimize the reverse KL divergence between an approximate distribution $q_\theta(\tau)$ and the true posterior.

We define our variational distribution $q_\theta(\tau)$ as follows (again, this form might seem arbitrary, but it is motivated by deeper theory—see resources for details):

$$q_\theta(\tau) = p(s_1) \prod_t p(s_{t+1} \mid s_t, a_t) q_\theta(a_t \mid s_t)$$

Minimizing the reverse KL divergence gives us the following objective:

$$\begin{aligned}
\text{KL}(q_\theta(\tau) \parallel p(\tau \mid o_\tau = 1)) &= \int q_\theta(\tau) \log \frac{p(\tau) p(o_\tau = 1 \mid \tau)}{q_\theta(\tau) p(o_\tau = 1)} d\tau \\
&= \int q_\theta(\tau) \log p(o_\tau \mid \tau) d\tau - \text{KL}(q_\theta(\tau) \parallel p(\tau)) + \text{const} \\
&= \int q_\theta(\tau) \log p(o_\tau \mid \tau) d\tau - \int q_\theta(\tau) \log \frac{q_\theta(\tau)}{p(\tau)} d\tau \\
&= \int q_\theta(\tau) \left(\log p(o_\tau \mid \tau) - \log \frac{q_\theta(\tau)}{p(\tau)} \right) d\tau \\
&= \int q_\theta(\tau) \left(\log p(o_\tau \mid \tau) - \log \frac{p(s_1) \prod_t p(s_{t+1} \mid s_t, a_t) q_\theta(a_t \mid s_t)}{p(s_1) \prod_t p(s_{t+1} \mid s_t, a_t)} \right) d\tau \\
&= \int q_\theta(\tau) \left(\log p(o_\tau \mid \tau) - \log \prod_t q_\theta(a_t \mid s_t) \right) d\tau \\
&= \int q_\theta(\tau) \left(\log \left(\exp \left(\sum_{t=0}^T r(s_t, a_t) \right) \right) - \log \prod_t q_\theta(a_t \mid s_t) \right) d\tau \\
&= \int q_\theta(\tau) \left(\sum_t r(s_t, a_t) - \log q_\theta(a_t \mid s_t) \right) d\tau \\
&= \mathbb{E}_{(s_{1:T}, a_{1:T}) \sim q_\theta} \left[\sum_t r(s_t, a_t) - \log q_\theta(a_t \mid s_t) \right] \\
&= \sum_t \mathbb{E}_{(s_t, a_t) \sim q_\theta} [r(s_t, a_t) + H(q_\theta(a_t \mid s_t))]
\end{aligned}$$

This objective is identical to the one encountered in maximum entropy reinforcement learning, specifically when we discussed Soft Actor-Critic (SAC) in Section 9.4. In that sense, SAC can be interpreted as performing the expected information projection (I-projection).

It's worth noting that, in this setting, exact inference over the Hidden Markov Model (HMM) could also be performed using the forward-backward algorithm. However, we will not cover that here. For those interested, the resources provided include further materials that explore this topic in more depth.

10.3 Resources

A great introduction to variational inference can be found in [BKM17]. Another excellent resource is "CS 285: Lecture 18 – Variational Inference" by Sergey Levine, which provides an intuitive and accessible overview. In "Lecture 19 – Control as Inference", he derives many of the equations referenced above, which we treated as given, and shows how to perform **exact** inference using the forward-backward algorithm to compute the posterior [LSM⁺20, RAI20]. Much of the material in these lectures is based on his tutorial Reinforcement Learning and Control as Probabilistic Inference [Lev18].

11 Trust Regions for Policy Gradients

One problem that arises with policy gradients methods is that the distance in parameter space does not translate into policy space. As a result, small changes in the parameters can lead to a large changes in policy behaviour. This could even go so far as to make our policy so bad that we cannot recover from it. A popular method is to limit the difference between subsequent policies to increase stability.

11.1 Trust regions

To ensure that the difference between subsequent policies remains limited, we can reformulate the general objective of maximizing with respect to both the previous and current policies as we did in section 7.4. This formulation includes an additional constraint to prevent large differences between consecutive policies which is achieved by limiting a dissimilarity measure $D(\cdot, \cdot)$ (i.e. Kullback-Leibler divergence) between the distributions:

$$\theta_{\text{new}} = \arg \max_{\theta} \mathbb{E}_{(s,a) \sim \pi_{\text{old}}} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\text{old}}(a|s)} A^{\pi_{\text{old}}}(s, a) \right] \quad \text{s.t. } D(\pi_{\theta}(a|s), \pi_{\text{old}}(a|s)) \leq \epsilon, \forall s \in S$$

We now have a constrained optimization problem that requires specialized methods for solving, which we will look at next.

11.2 Lagrangian Multipliers

We start by introducing the following optimisation problem:

$$\underbrace{\min_x f(x)}_{\text{objective}} \quad \text{s.t. } \underbrace{h_i(x) \geq b_i, \text{ for } i = 1 \dots K}_{\text{constraints}}$$

We then define the Lagrangian function/ Lagrangian, which combines the objective and constraint into a single function:

$$L(x, \lambda) = f(x) - \sum_{i=1}^K \lambda_i (h_i(x) - b_i) \quad \text{s.t. } \lambda_i \geq 0, \text{ for } i = 1 \dots K$$

We call the λ_i the Lagrange multipliers. There are now two ways how we can solve this problem in order to solve our original optimization problem :

Primal optimization problem

$$\min_x \max_{\lambda} L(x, \lambda)$$

Dual optimization problem

$$\lambda^* = \arg \max_{\lambda} g(\lambda), \quad g(\lambda) = \min_x L(x, \lambda)$$

The question now is: which problem should we choose to solve, and do we actually obtain the optimal solutions through this approach?

The solutions of the primal problem also minimize the unconstrained objective. However, solving the primal problem can often be challenging, especially since it might not possess the necessary convexity properties, making it hard to solve directly. But what about the dual function?

Since $g(\lambda)$ is a pointwise minimum of affine functions (note that $\mathcal{L}(x, \lambda)$ is affine in λ , even though $f(x)$ might not be linear or affine but we treat it as a constant with respect to λ), the dual function is concave (see the proof [hc]). Since $g(\lambda)$ is concave and the constraints are linear (meaning they are convex), maximizing $g(\lambda)$ over λ is a convex optimization problem. Furthermore, because convex functions have the property that every local minimum is also a global minimum, it is easier to solve using standard optimization techniques than the primal problem.

But what can we say about the solutions from both methods? Any feasible solution to the primal problem is at least as large as any feasible solution to the dual problem. Therefore, the solution to the primal provides an upper bound to the solution of the dual, and the solution of the dual provides a lower bound to the solution of the primal. This fact is known as **weak duality**:

$$g(\lambda) = \min_x \mathcal{L}(x, \lambda) \leq \min_x f(x) = p^* \quad \Rightarrow \quad d^* = \max_{\lambda} g(\lambda) \leq p^*$$

In general the optimal values of the primal and dual problems, p^* and d^* do not need be equal. The difference between them is called the **duality gap**.

11.2.1 Lagrangian Multipliers cookbook

1. Write down Lagrangian $L(x, \lambda) = f(x) - \sum_{i=1}^K \lambda_i (h_i(x) - b_i)$
2. Obtain optimal solution for primal parameters (primal solution) $\frac{dL(x, \lambda)}{dx} = 0 \rightarrow x^* = u(\lambda)$
3. put x^* back into the Lagrangian $g(\lambda) = L(u(\lambda), \lambda)$
4. obtain optimal solution for $g(\lambda) \rightarrow \lambda^* = \arg \max_{\lambda} g(\lambda)$, s.t. $\lambda_i \geq 0 \forall i$
5. $x^* = u(\lambda^*)$

An important aspect of using the Lagrangian is how to construct it, depending on the type of optimization problem and constraints. Consider the following simple example:

$$\min_x x^2 \text{ s.t. } x \geq 1 \longrightarrow L(x, \lambda) = x^2 - \lambda(x - 1)$$

The Lagrangian must be constructed in such a way that violating the constraint will not lead to an optimal solution. Since Lagrange multipliers are non-negative (i.e., $\lambda \geq 0$), violating the constraint $x \geq 1$ would result in adding a positive term, even though the goal is to minimize the objective function. This ensures that when $x < 1$ the Lagrangian penalizes this violation, and vice versa for a maximization problem. To gain a deeper understanding, let's now consider a more complex example.

$$\begin{aligned} & \arg \max_{\omega} \int p_{\omega}(\theta) g(\theta) d\theta \quad \text{s.t.} \quad \text{KL}(p_{\omega}(\theta) || p_{\text{old}}(\theta)) \leq \epsilon, \quad \text{H}(p_{\omega}(\theta)) \geq \beta, \quad \int p_{\omega}(\theta) d\theta = 1. \\ \rightarrow L(p_w, \eta, \kappa) &= \int p_w(\theta) g(\theta) d\theta + \eta (\epsilon - \text{KL}(p_w(\theta) || p_{\text{old}}(\theta))) + \kappa (\text{H}(p_w(\theta)) - \beta) + \lambda \left(\int p_w(\theta) d\theta - 1 \right) \end{aligned}$$

Notice how we now use + instead of - because this is a maximization problem. And try verifying for yourself, for each of the constraints, what the outcome (in terms of sign) would be if they were violated. Does this help or penalize in finding the optimal solution?

11.3 Trust Region Policy Optimization (TRPO)

One of the first methods to introduce this idea of constraining the update step was TRPO [SLM⁺17]. The objective is defined as follows

$$\max_{\pi} J(\theta) = \mathbb{E}_{(s,a) \sim \pi_{\text{old}}} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\text{old}}(a|s)} A^{\pi_{\text{old}}}(s, a) \right] \quad \text{Constraint: s.t.: } \mathbb{E}_{\pi_{\text{old}}} [\text{KL}(\pi_{\theta} || \pi_{\text{old}})] \leq \epsilon$$

In order to derive the update step the TRPO Paper uses something called the Natural-Gradient. In standard gradient-based optimization, we use the gradient of the objective function with respect to the parameters to update the model. However, this method treats all directions in parameter space as if they are equally important, which may not be the case since parameters in different regions may have different scales/curvatures leading to inefficient learning/bad model performance.

The natural gradient takes into account the geometry of the parameter space. It adjusts the gradient by scaling it with respect to the "natural" metric, which is defined by the Fisher information matrix (a matrix that measures the sensitivity of the model's output with respect to its parameters).

Natural gradients use a Taylor expansion approximation of the trust region problem

$$\mathcal{L}(\theta_k, \theta) \approx \nabla_{\theta} J(\theta)^T (\theta - \theta_k)$$

$$\bar{D}_{KL}(\theta || \theta_k) \approx \frac{1}{2} (\theta - \theta_k)^T F (\theta - \theta_k) \quad F = \underbrace{\frac{d \text{KL}(\pi_{\theta} || \pi_{\theta_{\text{old}}})}{d\theta d\theta}}_{\text{Fisher information matrix}} = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) \nabla_{\theta} \log \pi_{\theta}(a|s)^T]$$

This approximate problem can be analytically solved by the methods of Lagrangian duality, yielding the solution:

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\epsilon}{\nabla_{\theta} J(\theta)^T F^{-1} \nabla_{\theta} J(\theta)}} F^{-1} \nabla_{\theta} J(\theta).$$

There is one small issue with this approach when implemented naively: it concerns storing the inverse of the Fisher information matrix. Since the matrix F is $|\theta| \times |\theta|$ dimensional, it would require excessive memory for a neural network with millions of parameters. TRPO sidesteps the issue by using the conjugate gradient algorithm to solve $Fx = \nabla_{\theta} J(\theta)$ for $x = F^{-1} \nabla_{\theta} J(\theta)$, requiring only a function which can compute the matrix-vector product Fx instead of computing and storing the whole matrix F directly.

Algorithm 18 TRPO (OpenAI’s Spinning Up version [Ach18])

Input: initial policy parameters θ_0 , initial value function parameters ϕ_0

Hyperparameters: KL-divergence limit δ , backtracking coefficient α , maximum number of backtracking steps K

for $k = 0, 1, 2, \dots$ **do**

Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_{\text{old}} = \pi(\theta_k)$ in the environment.

Compute rewards-to-go $\hat{R}_t$.

Compute advantage estimates, $\hat{A}_t$ (using any method of advantage estimation) based on the current value function V_{ϕ_k} .

Estimate policy gradient as

$$\hat{g}_k = \frac{1}{|\mathcal{D}_k|} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \nabla_{\theta} \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\text{old}}(a_t|s_t)} \hat{A}_t$$

Use the conjugate gradient algorithm to compute

$$\hat{x}_k \approx \hat{H}_k^{-1} \hat{g}_k,$$

where $\hat{H}_k$ is the Hessian of the sample average KL-divergence.

Update the policy by backtracking line search with

$$\theta_{k+1} = \theta_k + \alpha^j \sqrt{\frac{2\delta}{\hat{x}_k^T \hat{H}_k \hat{x}_k}} \hat{x}_k,$$

where $j \in \{0, 1, 2, \dots, K\}$ is the smallest value which improves the sample loss and satisfies the sample KL-divergence constraint.

Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

end for

11.4 Proximal Policy Optimization (PPO)

PPO [SWD⁺17] is another method that incorporates the idea of trust regions. There are two main variants: PPO-Penalty and PPO-Clip. We will focus here only on PPO-Clip.

PPO-Clip does not rely on an explicit dissimilarity measure between policies. Instead, it controls the difference between successive policies using a clipping mechanism and a min function, which are built directly into the loss function.

$$\theta_{k+1} = \arg \max_{\theta} \mathbb{E}_{s, a \sim \pi_{\theta_k}} [L(s, a, \theta_{\text{old}}, \theta)]$$

$$L(s, a, \theta_{\text{old}}, \theta) = \min \left(\frac{\pi_{\theta}(a|s)}{\pi_{\theta_{\text{old}}}(a|s)} A^{\pi_{\theta_{\text{old}}}}(s, a), \text{clip} \left(\frac{\pi_{\theta}(a|s)}{\pi_{\theta_{\text{old}}}(a|s)}, 1 - \epsilon, 1 + \epsilon \right) A^{\pi_{\theta_{\text{old}}}}(s, a) \right),$$

As stated in the original PPO-Paper [SWD⁺17] : "...The second term (clipped ratio), modifies the objective by clipping the probability ratio, which removes the incentive for moving it outside of the

interval $[1 - \epsilon, 1 + \epsilon]$. Finally, we take the minimum of the clipped and unclipped objective, so the final objective is a lower bound (i.e. a pessimistic bound) on the unclipped objective. With this scheme, we only ignore the change in probability ratio when it would make the objective improve, and we include it when it makes the objective worse “. The intuition for the clipping can be best seen in figure 12.




Figure 12: Plots showing one term (i.e a single timestep) of the surrogate function L as a function of the probability ratio r , for positive advantages left and negative advantages right, from [SWD⁺17]

When the advantage is positive meaning the action we took lead to a better value then we would want the policy to do more of this actions but the clipping in this case lets us not go beyond the $1 + \epsilon$ ratio. When the advantage is negative meaning the action was less than favourable then there has to be at minimum a punishment of $1 - \epsilon$ but there is no lower bound to it.

Algorithm 19 PPO-Clip (OpenAI’s Spinning Up version [Ach18])

Input: initial policy parameters θ_0 , initial value function parameters ϕ_0

for $k = 0, 1, 2, \dots$ **do**

Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_{\text{old}} = \pi(\theta_k)$ in the environment.

Compute rewards-to-go $\hat{R}_t$.

Compute advantage estimates, $\hat{A}_t$ (using any method of advantage estimation) based on the current value function V_{ϕ_k} .

Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left(\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} A^{\pi_{\theta_{\text{old}}}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_{\text{old}}}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

end for

PPO shows very impressive results, but there are many additional hacks needed to make it really perform well, to list some: Value function clipping, Reward scaling, orthogonal initialisation, Adam learning rate annealing, reward clipping, observation normalisation, observation clipping, hyperbolic tan activations, Global gradient clipping

11.5 Differentiable Trust Region Layers

Recent work has shown that convex optimization problems can be embedded as differentiable layers within neural networks [AAB⁺19], enabling end-to-end training through optimization- based components. Building on this idea, differentiable trust region layers enforce trust region constraints directly within the policy network [OBV⁺21]. This transforms what would otherwise be a complex constrained optimization problem into a standard differentiable layer that integrates seamlessly into deep reinforcement learning pipelines.

11.6 Maximum Posteriori Optimization (MPO)

MPO is the idea of Trust-Regions for Actor Critic Methods. MPO uses 2 steps to optimize the policy (inspired by Expectation Maximization)

- E-Step: Solve Trust Region problem for non-parametric distribution
- M-Step: Weighted maximum likelihood using weights from E-step

11.7 Self-Test Questions

1. Why trust regions are important for policy gradients?

Vanilla policy gradients suffer from the fact that the distance in parameter space does not translate into policy space which can cause Overshooting (The update misses the reward peak and lands in a sub- optimal policy region) or Undershooting (Taking needlessly small steps in the gradient direction causes slow convergence). Trust regions introduce a constraint to the objective such that subsequent policies do not have a high difference.

2. Which divergence is typically used for the trust region and what are its properties?

Typically the Kullback-Leiber-Divergence gets used which has the following properties

- always positive $KL(p||q) \geq 0$
- $KL(p||q) = 0 \iff p = q$
- Non-symmetric $KL(p||q) \neq KL(q||p)$

3. How does Lagrangian optimization work and why is it often easier to solve the dual than solving the primal objective.

For method see 11.2. Solving the dual is often easier because maximizing $g(\lambda)$ is typically a convex optimization problem.

4. How to implement trust-regions with discrete actions using Lagrangian optimization?

We use lagrangian optimization to solve the optimization objective

$$\arg \max_{\pi} \sum_a \pi(a)r(a), \quad \text{s.t. } KL(\pi(a)||q(a)) \leq \epsilon, \sum_a \pi(a) = 1$$

5. How natural gradients are connected to trust regions?

Natural gradients approximate the Trust Region problem by using Taylor Expansion Approximation.

6. How the Fisher Information Matrix is used in the natural gradient algorithms?

The Fisher Information Matrix is used to approximate the KL divergence

$$\bar{D}_{KL}(\theta||\theta_k) \approx \frac{1}{2}(\theta - \theta_k)^T F(\theta - \theta_k)$$

7. Why are natural gradients difficult to use for bigger networks?

The Fisher Information Matrix is defined as (in the context of TRPO)

$$\mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) \nabla_{\theta} \log \pi_{\theta}(a|s)^T]$$

The number of elements of that gradient is $|\theta|$ (in NN every weight) so the Fisher Informations Matrix is of dimension $|\theta| \times |\theta|$. For a DNN with millions of parameters this is very computational burden.

8. **What is the main idea of PPO (clipped version) and what are the benefits?**

The idea here is also to limit the gradient step according to the distance of the policies. But it does this without the KL constraint instead it clips the gradient if the ratio of the current and old policy is too big. This is much simpler to implement than TRPO and can be used with optimizers like ADAM.

11.8 Resources

A short and easy-to-understand introduction to Lagrangian duality can be found in [Kno10]. There's also a great blog post about natural gradients by [Wil19]. For learning about TRPO and PPO, the OpenAI Spinning Up series explains the ideas very clearly [Ach18].

12 Evolution Strategies

Traditional deep learning methods often rely on stochastic gradient descent (SGD) to optimize a given objective. However, this approach assumes that gradients of the objective function are available — an assumption that doesn’t always hold, especially when the function is non-differentiable or its analytic form is unknown.

Evolution Strategies (ES) offer an alternative: they are gradient-free, black-box optimization algorithms that can optimize objective functions without requiring explicit gradient information. This makes them particularly useful in settings where gradients are difficult or impossible to compute. Below is the general form of ES for optimizing $p_\theta(x)$:

1. Generate a population of samples $D = \{(x, f(x_i))\}$ where $x_i \sim p_\theta(x)$
2. Evaluate the “fitness” of samples in D
3. Select the best subset of individuals and use them to update θ , generally based on fitness or rank. Go to 1

In reinforcement learning (RL), episode-based algorithms often use black-box optimization techniques like ES to maximize the expected return of an entire trajectory — rather than focusing on step-by-step actions. This approach is particularly effective when dealing with sparse or non-Markovian reward structures, where traditional, step-based methods struggle. Such an episode-based Task is for instance the Hopper task, where the goal is to jump as high as possible. Rewarding based on the current height (Markovian reward) would be less effective since height can fluctuate quickly. A better approach is to reward the agent based on the highest point reached during the episode, providing clearer guidance towards the desired behavior.

So the overall objective is to find a distribution $\pi_\theta(x)$ over variables $x \in \mathbb{R}^n$ that maximizes the expected value of the episodic return $R(\tau)$:

$$\theta = \arg \max_{\theta} \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$$

In this setting, Evolution Strategies (ES) can be used as a black-box optimizer to iteratively improve the policy parameters. The process typically follows these steps:

1. Explore: sample parameters $\theta_i \sim p_k(\theta)$
2. Evaluate: assess quality of parameters by generating trajectory $\tau_i \sim p_{\theta_i}(\tau)$ and define its quality through $g(\theta_i) = \mathbb{E}_{\tau \sim p_{\theta_i}(\tau)} [R(\tau)]$
3. Update: compute new search distribution p_{k+1}

12.1 Gaussian Evolution Strategies

In this section, we consider Evolution Strategies (ES) where the search distribution is modeled as a Gaussian. The goal is to learn the parameters of this distribution to efficiently explore the space of solutions. The complexity of the method depends on how expressive the Gaussian is:

First-order: Diagonal Gaussian search distribution Here, the Gaussian distribution is parameterized as follows:

$$p_\omega(\theta) = \mathcal{N}(\mu, \sigma^2 I), \text{ with } \omega = \{\mu, \sigma\}$$

This variant is simple and highly scalable, making it suitable for high-dimensional problems. However, it uses an isotropic variance (i.e., the same in all directions), which often leads to slower convergence due to limited directional guidance during exploration.

Second-order: Full-Covariance Gaussian search distribution Here, the Gaussian distribution is parameterized as follows:

$$p_\omega(\theta) = \mathcal{N}(\mu, \Sigma), \text{ with } \omega = \{\mu, \Sigma\}$$

This allows for more flexible and efficient exploration by adapting the shape and orientation of the distribution through the covariance matrix. While it typically converges faster than the diagonal version, its computational cost scales poorly with dimensionality, making it less practical for very high-dimensional problems.

12.1.1 Canonical Evolution Strategy (CES)

A basic first-order algorithm, using a diagonal Gaussian distribution, is the Canonical Evolutionary Strategy (CES). It updates only the mean μ , keeping σ fixed.

Algorithm 20 Canonical Evolutionary Strategy Algorithm

INPUT: $M \in \mathbb{N}$ number of elites

init. $w_i = \frac{\log(M+0.5) - \log(i)}{\sum_{j=1}^M \log(M+0.5) - \log(j)}$

repeat

Sample N parameter vectors $\theta_i \sim \mathcal{N}(\mu_k, \sigma^2 I)$

Evaluate parameter vectors $g_i = g(\theta_i)$

Sort $(\theta_i, \dots, \theta_N)$ according to g_i (best ones come first)

update μ (σ is typically fixed):

$$\mu_{k+1} = \mu_k + \sum_{i=1}^M w_i (\theta_i - \mu_k) = \sum_{i=1}^M w_i \theta_i$$

until Result is good enough

12.1.2 Cross-Entropy Method (CEM)

The Cross-Entropy Method (CEM) is one of the most popular second-order ES variants. It updates both the mean and the covariance matrix of the search distribution using elite samples.

Algorithm 21 Cross Entropy Method

INPUT: $M \in \mathbb{N}$ number of elites

repeat

Sample N parameter vectors $\theta_i \sim \mathcal{N}(\mu_k, \Sigma_k)$

Evaluate parameter vectors $g_i = g(\theta_i)$

Sort $(\theta_i, \dots, \theta_N)$ according to g_i (best ones come first)

update mean :

$$\mu_{\text{elites}} = \frac{1}{M} \sum_{i=1}^M \theta_i \quad \mu_{k+1} = (1 - \alpha) \mu_k + \alpha \mu_{\text{elites}}$$

update covariance:

$$\Sigma_{\text{elites}} = \frac{1}{M} \sum_{i=1}^M (\theta_i - \mu_{\text{elites}})(\theta_i - \mu_{\text{elites}})^T$$
$$\Sigma_{k+1} = (1 - \alpha) \Sigma_k + \alpha \Sigma_{\text{elites}}$$

until Result is good enough

However both these methods have some key limitations:

- Lack of a clear optimization objective: Since they rely on ranking the candidates rather than directly optimizing the expected reward, they do not explicitly maximize the expected return
- Monotonic weighting transformation: The method's ranking approach is monotonic, which can result in convergence to a local optimum, especially in deterministic environments.

- Challenges in stochastic settings: In environments with stochastic rewards, the behavior of ranking-based methods becomes less predictable and harder to analyze.

12.2 Trust-Region Methods for Stochastic Search

One issue with approaches that iteratively update the search distributions is that the updates can either be too moderate or too greedy, often leading to suboptimal behavior. To address this, we can introduce the concept of trust regions, which we encountered in a previous section and constrains the updates in a way that avoids drastic changes, promoting more stable learning. In this case, the objective is defined as:

$$\theta = \arg \max_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] \text{ s.t. } KL(p || p_{\text{old}}) \leq \epsilon$$

This can then be solved using Lagrange multipliers, yielding a closed-form solution for the search distribution. For a more detailed explanation and an alternative method involving trust regions, refer to [HN24].

12.3 Self-Test Questions

1. **What the advantages/disadvantages of an episode-based RL formulation are**
 - + Clear Objective Definition: the agent’s goal is often clearly defined in terms of episodes
 - + Non-Markovian Rewards: can be more intuitive and better aligned with the overall task goals.
 - Delayed Feedback: In some tasks, the consequences of an agent’s actions may not be immediately apparent, leading to delayed rewards.
2. **How episode-based RL is connected to black-box optimization / evolutionary strategies (ES)**
3. **How we can use search distributions for black-box optimization**

A search distribution is a probabilistic model used to sample candidate solutions to the black-box objective function. The idea is to treat the problem as one of searching for the optimal set of parameters (or solutions) by sampling from a distribution, then refining the distribution based on feedback (evaluations of the objective function aka black box).
4. **What is the difference between first-order and second order methods for optimizing search distributions? Pros/Cons?**

The difference is w.r.t. the complexity of the gaussian search distribution.

 - first-order : $\mathcal{N}(\mu, \sigma^2 I)$ (better for high dimension, converges slowly)
 - second order : $\mathcal{N}(\mu, \Sigma)$ (fast convergence, more directed exploration in parameter space, computable inefficient for high dimensions)
5. **How a simple policy gradient algorithm for the episode-based case looks like**
6. **What are the qualitative differences to “deterministic” gradient descent using finite differences**
7. **Why most ES use rankings instead of the fitness values**

Evolutionary Strategies (ES) often use rankings instead of raw fitness values to avoid issues like noise, outliers, or scaling problems in the fitness function. Rankings are less sensitive to extreme fitness values, ensuring that the selection process is more stable and robust.
8. **How does the cross-entropy method work and what is its connection to evolutionary strategies**

See algorithm 21. CEM can be viewed as a special case of evolutionary strategies, particularly in

its use of a population-based search and selection mechanism. Instead of using a simple mutation-based update as in ES, CEM updates the search distribution based on the best-performing candidates.

9. **How to use trust-regions in for episode-based RL?**

See 12.2

10. **Why can we use exact trust regions in this case (as opposed to Lecture 6)?**

Because for Gaussian distributions there is a closed form solution of the KL-Divergence which does not require sampling

12.4 Resources

A great resource on Evolution Strategies is the blog post by Weng (2019) [[Wen19](#)].

13 Bayesian Optimization

This section provides a simplified understanding of Bayesian Optimization (BayesOpt) and is not intended to be a detailed treatment of the topic. A comprehensive explanation would require additional concepts and methods that are beyond the scope of this discussion.

BayesOpt is another method designed for black-box, derivative-free global optimization. This means it is tailored for optimization problems where the objective function is unknown, complex, or cannot be directly analyzed (black-box), and or where derivatives are unavailable or difficult to compute (derivative-free).

Rather than directly optimizing the objective function, BayesOpt relies on a probabilistic surrogate model, typically a Gaussian Process (GP) to approximate the function's behavior. This surrogate model helps guide the search for the optimal solution. BayesOpt is particularly focused on global optimization, meaning it seeks the best solution across the entire search space rather than just local optima.

Through an acquisition function, BayesOpt strikes a balance between exploration (sampling areas with high uncertainty) and exploitation (focusing on areas that are already promising), making it an efficient approach when function evaluations are expensive or time-consuming. But it is limited by the number of dimensions (mostly ≤ 20) since it is very computationally expensive.

13.1 Surrogate Model (Gaussian Process Regression)

In Bayesian optimization, the surrogate model is typically constructed using Gaussian Process (GP) regression, a Bayesian statistical approach. A GP model treats function values as random variables, assuming a prior distribution that is multivariate normal, characterized by a mean vector and a covariance matrix. The mean vector is computed by evaluating a mean function μ_θ at each sampled input point x_i , while the covariance matrix Σ_θ is determined using a kernel function. This kernel encodes our belief that input points closer together in the input space are likely to have similar function values. The kernel is crucial as it dictates the correlation between points in the function's input space.

13.2 Acquisition Function

The acquisition function is central to guiding the optimization process. It decides where to evaluate next by balancing exploration (sampling in areas of high uncertainty) and exploitation (focusing on areas where the model predicts good outcomes). Popular acquisition functions include Upper Confidence Bound (UCB), Knowledge Gradient, Entropy Search, and Predictive Entropy Search, among others.

For instance, the UCB (Upper Confidence Bound see 5.4) acquisition function is a widely used strategy. It combines the predicted mean $\mu(x)$ from the Gaussian Process with the uncertainty $\sigma(x)$ at each point. The UCB function is defined as:

$$\text{UCB}(x) = \mu(x) + \sqrt{\beta} \cdot \sigma(x)$$

The goal is to maximize the UCB function to find the next point that either has high predicted values (exploitation) or high uncertainty (exploration). To identify the next point x_{next} for evaluation, the acquisition function is optimized over the input space:

$$x_{\text{next}} = \arg \max_x \text{UCB}(x)$$

13.3 Optimization Loop

The optimization process proceeds iteratively. In each iteration:

- A new input point is selected based on the acquisition function.
- The objective function is evaluated at that point.
- The surrogate model (e.g., Gaussian Process) is updated with the new data.

The loop continues until a predefined stopping criterion is met, such as reaching a maximum number of iterations, exceeding a time limit, or observing negligible improvement in the objective function.

13.4 Self-Test Questions

1. **Why we do we use a Bayesian treatment for learning the fitness function?**
Bayesian methods allow for efficient optimization of expensive, noisy, or unknown functions by using a probabilistic model (e.g., Gaussian Process) to capture uncertainty and guide the search for optimal points without needing derivative information.
2. **What a GP is and how to obtain the GP posterior**
A Gaussian Process (GP) is a probabilistic model that defines a distribution over functions. The posterior is obtained by updating the prior distribution with observed data using Bayes' rule, resulting in a model that reflects the updated beliefs about the function.
3. **How to use acquisition functions in BO?**
Acquisition functions guide where to sample next. They balance exploration (uncertain regions) and exploitation (promising regions), commonly using Expected Improvement (EI) or alternatives like Knowledge Gradient (KG).
4. **Why BO can give you the global maximum instead of just a local one?**
BO uses a probabilistic model (GP) to model uncertainty and avoid local optima by selecting points that provide the most information, thus increasing the likelihood of finding the global maximum.
5. **Difference between considering the exploration-exploitation tradeoff and pure exploration phase.**
Exploration-exploitation balances uncertainty and predicted value, improving efficiency. Pure exploration focuses solely on areas with high uncertainty, which can be useful early in optimization but may not lead to the best result.
6. **The key ideas of the acquisition functions**
Acquisition functions guide the optimization by quantifying the tradeoff between exploring uncertain areas and exploiting known good areas to maximize the objective.
7. **How to deal with hyper-parameters of the GP**
Hyperparameters (e.g., length scale) are learned through maximum likelihood estimation or Bayesian inference to best fit the GP model to the observed data.
8. **What are the pros and cons of BO in comparison to Stochastic Search?**
 - Pros of BO: Efficient with fewer evaluations, good for expensive functions, global optimization.
 - Cons of BO: Computationally intensive for high dimensions.
 - Stochastic Search pros: Simple and scalable for high dimensions, no model required.
 - Cons of Stochastic Search: Less efficient for expensive functions, risk of local optima.

14 Planning

In the previous sections, we focused on model-free algorithms (except dynamic programming), where we **learned** policies or value functions by interacting with the environment, without any explicit model. Now, we shift to model-based settings, where we assume either that we can learn a model of the environment or are provided with one.

Planning refers to any computational process that uses a model to produce or improve a policy for interacting with the modeled environment. We have already encountered planning in the context of value and policy iteration (Section 4). In the following, we assume a model is available and explore various planning algorithms. (The next chapter will address how a model can potentially be learned.)

14.1 Open-/Closed-Loop Planning

In planning, we distinguish between open-loop and closed-loop planning. In open-loop planning, the environment provides no feedback during execution. We are given an initial state and must immediately output a sequence of actions. For deterministic dynamics this corresponds to solving an optimization problem to find the sequence of actions that maximizes the return :

$$a_1, \dots, a_T = \arg \max_{a_1, \dots, a_T} \sum_{t=1}^T r(s_t, a_t) \text{ s.t. } s_{t+1} = p(s_t, a_t)$$

In an environment with stochastic dynamics the optimization is stated as

$$a_1, \dots, a_T = \arg \max_{a_1, \dots, a_T} \mathbb{E} \left[\sum_{t=1}^T r(s_t, a_t) \middle| a_1, \dots, a_T \right] \quad s_{t+1} \sim p(s_{t+1} | a_t, s_t)$$

with $p(s_1, \dots, s_{T+1} | a_1, \dots, a_T) = p(s_1) \prod_{t=1}^T p(s_{t+1} | a_t, s_t)$

In open-loop planning, trajectories in stochastic environments are often suboptimal. Consider the analogy of taking an exam. Let's say state s_1 represents the moment you receive the exam. The planned action sequence, in this case, corresponds to the answers you are going to give despite not knowing what the questions will be. Even if you're well-prepared, your answers are likely to be wrong, resulting in suboptimal outcomes since you can not look into the future. The fundamental problem is that open-loop plans cannot adapt to unexpected events or changes in the environment. We've already discussed open-loop methods like Random Shooting, Cross-Entropy Method (CEM), CMA-ES in section 12.

In contrast, closed-loop planning allows the environment to provide feedback. In this setting, you're given a state and asked to produce an action. After taking the action, the environment updates and returns the next state, continuing the cycle. This feedback loop enables more adaptive decision-making, as actions can be adjusted based on the state transitions and rewards observed during execution.




Figure 13: Illustration of open- and closed-loop planning, inspired from "CS 285: Lecture 10" [LSM⁺20, RAI20]

14.2 Closed-Loop Planning in Discrete Domains - Monte-Carlo Tree Search

For discrete domains, we have already explored dynamic programming approaches to derive policies given a model. However, in scenarios where the state space is large (e.g., Go, Chess, Atari, etc.), such methods become infeasible. An alternative is to use traditional tree search algorithms, but these also face limitations. For environments with huge branching factors, the search tree grows exponentially, making it computationally expensive.

The approach we will be looking at next focuses on incrementally and asymmetrically building the search tree. Monte Carlo Tree Search (MCTS) addresses this by concentrating on the most promising moves and expanding the search tree through random sampling of the search space. MCTS consists of four main stages:

1. **Selection:** start from root R and select successive child nodes until a leaf node L is reached. The root is the current game state and a leaf is any node from which no simulation (playout) has yet been initiated. The decision which child to take next is determined by the highest UCB-Score 5.4
2. **Expansion:** unless the leaf L ends the game decisively (e.g. win/loss/draw) for either player, create one (or more) child nodes and choose node C from one of them. Child nodes are any valid moves from the game position defined by L.
3. **Simulation:** complete one random playout from node C. This step is sometimes also called playout or rollout. The rollout is typically performed with a default policy π_θ . A playout may be as simple as choosing uniform random moves until the game is decided (for example in chess, the game is won, lost, or drawn).
4. **Backpropagation:** use the result of the playout to update information in the nodes on the path from C to R. For 1-0 (win-loss games), that means updating the win-loss counts of all the ancestor nodes. Note that the win count updates alternate between 0 or 1 (e.g. for black-white players) as you go up the tree.




Figure 14: The different phases of MCTS, image from KIT

However, a major challenge arises when using Monte Carlo Tree Search (MCTS) in environments with continuous state and action spaces. In such cases, MCTS becomes significantly harder to apply effectively, as it's highly unlikely to encounter exactly the same states during planning. This makes it difficult to build and reuse a meaningful search tree. To address this issue, we will explore an alternative method that is better suited for planning in continuous domains in the following chapter. But before we do that, we will first take a closer look at how MCTS has been used for planning in reinforcement learning.

14.2.1 Alpha-Go and its successors

AlphaGo combined deep neural networks with Monte Carlo Tree Search (MCTS) and was trained using a combination of supervised learning from expert human games and reinforcement learning

through self-play. It used separate networks to predict the best move (policy) and the likely winner (value).

AlphaGo Zero is the successor to AlphaGo. AlphaGo Zero learned entirely from scratch, using only self-play without any human data. It also simplified the architecture by using a single neural network to predict both the policy and value, making it not only more efficient but also more powerful. This version surpassed the original AlphaGo despite having no prior knowledge of human strategies.

AlphaZero generalized the approach from AlphaGo Zero. It used the same learning framework as AlphaGo Zero but was applied to different games, including Chess, Shogi, and Go. AlphaZero learned purely through self-play and used MCTS alongside a neural network, just like AlphaGo Zero, but it demonstrated that this method could be adapted to multiple environments with known rules.

MuZero was the next major advancement. While AlphaZero still required a perfect simulator or full knowledge of the environment’s rules, MuZero removed this assumption entirely. It learned not just the policy and value function, but also an internal model of the environment’s dynamics—enabling planning without access to the real rules. Instead of learning to predict exact future states, MuZero learned a compact latent representation, from which it could plan and estimate rewards, values, and future outcomes.

AlphaTensor was designed to discover new algorithms instead of learning to play games like Chess or Go. Specifically, it learns to find efficient matrix multiplication algorithms by framing the task as a single-player reinforcement learning problem.

A nice visualisation of how MCTS works and how it got used in AlphaZero/Go can be found [here](#).

14.3 Resources

Much of the content presented here is based on Sergey Levine’s CS 285: Lecture 10 [LSM⁺20, RAI20]. A helpful visualization of how Monte Carlo Tree Search (MCTS) works, particularly in the context of AlphaZero can be found on [VAR20]. For a technical breakdown of AlphaGo see the article [Hui18], and for an overview of MuZero refer to DeepMind’s blog post [SAH⁺20].

15 Model Predictive Control / Closed-Loop Planning in Continuous Domains – Optimal Control

Model Predictive Control (MPC) is an optimal control strategy designed for managing complex systems, particularly those with constraints. At its core, MPC relies on predicting the system's future behavior and using this prediction to determine the optimal action at the current time step. As new data becomes available, the controller updates its plan, making MPC well-suited for dynamic and uncertain environments. The process unfolds as follows:

1. **Model of the System:** MPC begins with a mathematical model that describes how the system evolves in response to different inputs.
2. **Prediction Horizon:** At each time step, the controller uses the model to predict the system's future behaviour over a fixed time window, known as the prediction horizon. It simulates how the system would respond to different control inputs over this period.
3. **Optimization Problem:** The controller then formulates and solves an optimization problem to find the optimal sequence of control actions over the prediction horizon. This sequence is chosen to minimize a predefined cost function.
4. **Apply Only the First Action:** Instead of applying the entire sequence, MPC executes only the first action from the optimized plan. After this step the system state is then updated, and the process repeats.

By continuously updating its decisions in this way, MPC can react to unexpected changes and uncertainties in the environment. In the following we will explore how to solve such optimization problems (step 3) in different settings. We will focus on finite-horizon problems without constraints. Before we begin, we'll introduce some standard notation from optimal control theory and clarify how it relates to terms commonly used in reinforcement learning (RL):

- x represents the state of the system (s).
- u represents the control input or action (a).
- $c(x, u)$ is the cost function, which is the counterpart to the reward function in RL (note that RL typically maximizes reward, while control minimizes cost)
- $f(x, u)$ defines the system dynamics, describing how the state evolves. This corresponds to the transition function or transition probabilities in RL.

15.1 Linear Quadratic Regulator

We will start by looking at the Linear Quadratic Regulator (LQR). In LQR, "linear" refers to the system dynamics being linear, and "quadratic" refers to the cost function being quadratic. A typical LQR problem can be formulated as:

$$\min_{u_1, \dots, u_T} c(x_1, u_1) + c(f(x_1, u_1), u_2) + \dots + c(f(f(\dots)), u_T)$$

$$f(x_t, u_t) = F_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \mathbf{f} \quad c(x_t, u_t) = \frac{1}{2} \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T C_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T \mathbf{c}_t$$

To find the optimal sequence of actions, we minimize the cost function over time. The details of how to solve this can be found in the recommended resources. However, LQR is limited to systems with deterministic and linear state transitions.

15.2 Linear-quadratic-Gaussian Control

When dealing with stochastic transitions, we extend the LQR framework to the Linear Quadratic Gaussian (LQG) model. The overall problem formulation remains the same as in LQR, with the

key difference being how we define the system dynamics. In LQG, transitions are modeled as linear functions with additive Gaussian noise:

$$p(x_{t+1}|x_t, u_t) = \mathcal{N}(F_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \mathbf{f}, \Sigma_t)$$

In other words, we take the deterministic transition function from LQR and simply add Gaussian noise to account for uncertainty. The good news is that the same algorithm used to solve the LQR problem can also be applied to the LQG setting.

15.3 Approximating Non-Linear Systems (iLQR)

While LQR and LQG are effective for linear systems, real-world systems are often non-linear and require a different approach. To handle non-linear dynamics, we can approximate both the dynamics and the cost function locally using a Taylor expansion around the current trajectory point $(\hat{x}_t, \hat{u}_t)$.

$$\begin{aligned} f(x_t, u_t) &\approx f(\hat{x}_t, \hat{u}_t) + \nabla_{x_t, u_t} f(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix} \\ c(x_t, u_t) &= c(\hat{x}_t, \hat{u}_t) + \nabla_{x_t, u_t} c(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix} + \frac{1}{2} \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix}^T \nabla_{x_t, u_t}^2 c(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix} \end{aligned}$$

we then define

$$\bar{f}(x_t, u_t) = \underbrace{F_t}_{\nabla_{x_t, u_t} f(\hat{x}_t, \hat{u}_t)} \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \mathbf{f} \quad \bar{c}(x_t, u_t) = \frac{1}{2} \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T \underbrace{C_t}_{\nabla_{x_t, u_t}^2 c(\hat{x}_t, \hat{u}_t)} \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T \underbrace{\mathbf{c}_t}_{\nabla_{x_t, u_t} c(\hat{x}_t, \hat{u}_t)}$$

With these approximations, we effectively transform the non-linear control problem into a locally linear-quadratic one, which we can now solve using the standard LQR algorithm.

15.4 Resources

A more detailed explanation of Model Predictive Control and the solution to the equations above can be found in Sergey Levine's CS 285: Lecture 10 [LSM⁺20, RAI20].

16 Model-Based Reinforcement Learning (MBRL)

All the algorithms we have looked at so far were either model-free, meaning that we had no model of the world and just learned value functions by interacting with the environment, or the dynamic programming approach, which assumes that we do have a model and can solve for the optimal solution without having to interact with the world.

In model based reinforcement learning we learn a model from experience and use the learned model to do planning. By model we mean the transition probabilities/ dynamics. One way to learn the dynamics is to use equations of motion and then identify the parameters from the given data but this is not really scalable for complex systems. The more common approach is to use neural networks to learn the dynamics models. The most basic model based reinforcement learning algorithm would look like algorithm 22 where we use the experience $\{S_1, A_1, \dots, S_T\}$ and reduce it to a supervised learning problem.

Algorithm 22 Basic Model-Based RL

```
for  $i = 1, 2, \dots$  do
    collect data under the current policy
    learn dynamics model from past data
    improve policy by using dynamics model
end for
```

Model-based reinforcement learning (MBRL) offers several advantages, such as data efficiency, improved safety through simulation, and faster adaptation capabilities. However, the simple MBRL algorithm described in algorithm 22 also has some drawbacks. The problem arises from the circular dependency between data collection, model learning, and policy improvement. If any of these components perform poorly, everything can start to fail. This also introduces now two sources of approximation error: one from learning the model itself, and the other from using the learned model to derive the value function. There are many potential reasons for failure in any of these areas, but for now, we are focusing on the model learning aspect.

16.1 Distributional shift

One reason for bad performing policies which got trained by a learned model is distributional shift. Suppose we have trained a model on a given dataset of trajectories until it achieves low error on this data. However, since our model is not perfect, it will likely make small deviations from the true actions (e.g., slightly steering left) during testing. These errors accumulate over time, causing the model to encounter states it rarely or never saw during training. As a result, the model performs poorly on these unseen states.

The core problem is that the state distribution on which the model was trained differs from the distribution the model generates during testing. Consequently, the model struggles to extrapolate to these new regions of the state space. One way to address this is by augmenting the training dataset with real dynamics tuples (s, a, s') collected by executing the model's own actions in the actual environment. This approach, which we will also revisit in the Imitation Learning chapter, helps the model adapt to the distribution it will face at test time.

16.2 Model Exploitation

The second challenge is model exploitation, a phenomenon related to overfitting, as depicted in Figure 15. In this figure, the green curve represents the true Q-values, while the red curve shows the Q-values induced by the learned dynamics model. Since the model is used to calculate these Q-values, any inaccuracies in the model can lead to erroneous spikes in the predicted values. The policy, relying on these faulty predictions, may exploit these artifacts—choosing actions that appear optimal based on the model but are actually suboptimal or even harmful in the real environment. This can result in the policy overfitting to the inaccurate dynamics model (choosing this faulty spikes), leading the agent to take unrealistic actions and explore areas of the state space it has never encountered, ultimately causing a distributional shift.




Figure 15: Illustration of overfitting: The green curve represents the true underlying function (ground truth), while the blue dots are noisy samples drawn from it. The red curve shows the model's learned approximation, which fits the training data closely but fails to generalize well. From [Com24].

16.3 Uncertainty

One way to mitigate model exploitation is by incorporating uncertainty into the dynamics model. Rather than predicting a single next state for a given state-action pair, the model instead returns a distribution over possible next states. This allows the agent to evaluate actions based on their expected rewards while taking the uncertainty in the dynamics into account. By considering the full distribution of possible outcomes, the agent becomes less likely to overfit to inaccurate model predictions and is encouraged to choose actions that are more robust to uncertainty in the environment.




Figure 16: Gaussian process regression on noise-free dataset from [sl]

However, there are several challenges with this approach. From earlier chapters, we know that exploration is crucial for discovering optimal policies. If our model is too pessimistic, it may avoid risky actions, leading to a lack of exploration. As a result, the agent might not discover better solutions or improve its policy. In the following sections, we will explore the different types of uncertainty and how we can mitigate their impact to encourage more effective exploration while still accounting for the uncertainty in the learned dynamics.

16.3.1 Epistemic uncertainty

Epistemic uncertainty refers to the uncertainty about the world due to a lack of experience (or data). This type of uncertainty can be reduced by gathering more data to improve our understanding. However, another approach to managing this uncertainty is to explicitly incorporate it into the models themselves.

Bayesian Learning Bayesian learning addresses epistemic uncertainty by treating model parameters as probabilistic rather than fixed. Instead of assigning a single value to each weight, we represent them as probability distributions—commonly Gaussians defined by a mean and variance. This probabilistic framework allows us to quantify how uncertain we are about the learned parameters.

Ensembles Another simpler method to address epistemic uncertainty is by using an ensemble of models. Instead of training a single model, you train k models, each of which learns the task independently. For prediction, we can then average the outputs of all models to obtain a more robust prediction.

To ensure diversity among the models in the ensemble, it is helpful to train each model on independent training sets. This independence allows the models to explore different aspects of the problem and helps reduce the overall uncertainty in the predictions.

16.3.2 Aleatoric uncertainty

Aleatoric uncertainty refers to the inherent uncertainty within a system, meaning that no matter how much experience or data you gather, some aspects will always remain uncertain. A classic example of aleatoric uncertainty is predicting the outcome of a dice roll. The outcome is uncertain, and no matter how many dice rolls you make, you can never predict the exact result with certainty using a deterministic function.

In reinforcement learning, there are three primary sources of aleatoric uncertainty, each corresponding to a component of the Markov Decision Process (MDP):

- **Stochastic Rewards:** When the reward function is stochastic, there is an irreducible uncertainty about the true value of the reward.
- **Stochastic Observations:** Uncertainty in the observations received from the environment.
- **Stochastic Actions:** When actions in the environment are uncertain, resulting in variability in outcomes even if the agent's actions are deterministic.

To address aleatoric uncertainty, practitioners often use probabilistic neural networks, which can model this uncertainty explicitly by outputting a distribution over possible outcomes, rather than just point estimates. This approach allows the agent to account for inherent randomness in the environment and make better decisions under uncertainty.

16.4 Replanning

Since no learned model can perfectly predict the outcome of every action, one way to reduce the impact of model inaccuracies is through replanning. Unlike open-loop planning, which commits to a fixed sequence of actions and is prone to compounding errors, replanning continuously updates the plan based on new observations at each time step. This enables the agent to correct for past mistakes and adapt to changes in the environment. This strategy is commonly known as Model Predictive Control (MPC).

16.5 Probabilistic Ensembles with Trajectory Sampling (PETS)

PETS is a model-based reinforcement learning (MBRL) method introduced by [CCML18]. It represents a powerful approach to reinforcement learning by combining probabilistic modeling with planning. Rather than learning a policy directly, PETS simulates future outcomes using a learned model of the environment and chooses actions based on this simulated future.

16.5.1 Core Components

- **Probabilistic Ensembles:** PETS does not rely on a single neural network to model environment dynamics. Instead, it uses an ensemble of neural networks, each trained to predict the next state given the current state and action. These networks are probabilistic, typically outputting parameters of a Gaussian distribution over the next state.
- **Trajectory Sampling:** Rather than only choosing the next action based on immediate predictions, PETS samples entire sequences of future trajectories. This is achieved by:
 - Sampling multiple models from the ensemble.
 - Simulating rollouts using these models over several timesteps.
 - Sampling from the predicted distributions at each step (the models are probabilistic models so their outputs are distributions which we need to sample from in order to get the state).
 - Evaluating the total expected reward of each trajectory.
- **Model Predictive Control (MPC):** PETS uses the ensemble and trajectory sampling to perform planning at every timestep via MPC. The loop is as follows:
 - Trajectory Sampling from the current state.
 - Choose the sequence with the highest predicted reward.
 - Execute only the *first* action of the best sequence in the real environment.
 - Observe the next state, update the model with new data, and repeat the process.




Figure 17: Illustration of the PETS method from [CCML18].

16.6 Model-Based Policy Optimization (MBPO)

MBPO [JFZL21] combines model-free and model-based reinforcement learning into a unified approach. While a dynamics model is learned from collected data, the policy is not limited to being trained exclusively in a model-free or model-based manner. Instead, MBPO leverages both real environment interactions and short model-generated rollouts to train the policy. As the agent interacts with the environment, new data is gathered to continually refine the learned model and improve policy performance.

16.7 Model-Based-RL under Partial Observability

Previous methods typically assume that environment dynamics can be represented as functions mapping state-action pairs to next states. However, this assumption relies on access to the true underlying states of the environment, which is often not the case in practice.

This limitation becomes clear in image-based tasks, such as those found in Atari games. From a single image, it is generally impossible to directly infer latent properties like the velocity or direction of a moving object. In such cases, the agent receives only partial observations (e.g., raw pixels), and the true state must be inferred.

To address this challenge, we require a more general framework: the Partially Observable Markov Decision Process (POMDP). A POMDP extends the standard MDP formulation by modeling environments where the agent does not have full observability of the underlying state. Formally, a POMDP is defined by:

- state space S (not directly observable)

- action space A
- observation space O (observable)
- transition model $p(s_{t+1}|s_t, a_t)$ (not known)
- reward function $r(s_t, a_t)$
- observation model $p(o_t|s_t)$
- initial state distribution $p(s_0)$

In general, POMDPs represent hidden Markov models, as illustrated in Figure 18. If we knew the states S , the transition probabilities $p(s_{t+1}|s_t, a_t)$, the observation likelihoods $p(o_t|s_t)$ and if everything were Gaussian, we could compute the hidden states from the observations. This would allow us to apply the methods we’ve been using all along to learn policies. However, in reality, none of these assumptions typically hold.




Figure 18: A POMDP illustrated as a hidden Markov chain

A common approach in model-based reinforcement learning under partial observability is to first learn a mapping from high-dimensional observations (e.g., images) to a lower-dimensional latent state space. This latent space is designed to capture the essential features of the environment needed for decision-making, even if it does not correspond directly to the true underlying state. To learn this mapping, Variational Autoencoders (VAEs) are typically used, providing both an encoder (from observation to latent state) and a decoder (from latent state back to observation). Once this representation is learned, the dynamics model is trained directly in the latent space.

There are several possible choices for modeling the encoder. The most accurate approach is to learn the posterior $q_\psi(s_t | o_{1:t}, a_{1:t})$, which incorporates the full history of observations and actions. However, while this posterior provides the most precise estimate of the latent state, it is also the most complex to learn. A simpler, though less accurate, alternative is to use a single-step encoder, which learns the posterior $q_\psi(s_t | o_t)$ based solely on the current observation.

In recent years, several methods have been proposed that excel in model-based reinforcement learning. While a detailed discussion is beyond the scope, here are some notable approaches:

- PlaNet (introduced Recurrent State Space Models (RSSM)) [HLF⁺19]
- Dreamer (Dream to Control) [HLBN20]

16.8 Self-Test Questions

1. What is the difference between model-based and model-free RL?

In model-based RL, we learn a model of the environment and use it for planning. In model-free RL, we don’t focus on learning a model; instead, we learn value functions or policies by directly interacting with the environment.

2. Why is model learning in MBRL not just supervised learning?

In supervised learning, our data comes from a single underlying distribution p and we aim to learn the mapping from features to labels. However, model learning in model-based RL (MBRL) is cyclic. We use the model to improve a policy, which then generates new data that is used to

further optimize the model. This feedback loop results in a shifting distribution of data, so there is no fixed objective, unlike in supervised learning.

3. **What are the different types of uncertainty, how can they be captured and why is that important for MBRL?**

There are two primary types of uncertainty:

- **Epistemic uncertainty:** This refers to the uncertainty about the world due to a lack of data. It can be reduced by gathering more data or incorporating uncertainty directly into the model (e.g., through Bayesian learning).
- **Aleatoric uncertainty:** This refers to uncertainty inherent in the system, meaning no matter how much data we gather, some uncertainty remains irreducible. This type of uncertainty can be captured by using probabilistic models.

Understanding and addressing these types of uncertainty is crucial in model-based reinforcement learning (MBRL). If uncertainty is not properly accounted for, it can lead to model exploitation, where the agent might overfit to the learned model and make poor decisions based on inaccurate predictions, ultimately impacting the performance of the policy.

4. **How can you plan with an imperfect model?**

One approach is replanning after each action is taken, turning an open-loop problem into a closed-loop process. This way, the agent can adjust based on new observations and model corrections.

5. **What is partial observability, why is it important?**

Partial observability occurs when we cannot fully observe the state of the environment. In many real-world scenarios, this happens due to missing sensor data or limited context (e.g., in image-based tasks like Atari, where a single image doesn't provide information about the ball's velocity or direction). Since we can't directly observe the state, this complicates the use of algorithms that assume full observability.

6. **What are the state space assumptions?**

- observations only depend on the current state $p(o_t|s_t)$
- states depend only on the previous state and the action $p(s_{t+1}|s_t, a_t)$

7. **How can we use Variational Inference to learn deep SSMs/sequential VAEs?**

8. **How can we use those for model based RL under partial observability?**

Sequential VAEs and deep state-space models encode observations into a lower-dimensional latent space that represents the states. This latent space can then be used by a dynamics or inference model to predict future states, enabling planning and decision-making despite partial observability.

9. **How can you use a model to learn a parametric policy?**

A parametric policy can be learned by training policy networks and value-function networks on imagined or simulated trajectories generated by the learned model. These simulated trajectories provide a way to refine the policy without needing real-world data.

16.9 Resources

Much of the general content presented here is based on Sergey Levine's CS 285: Lecture 11 [LSM⁺20, RAI20]. A highly informative blog post [MIL23] by one of the authors of Dreamer V1, detailing the method. I also recommend checking out the blog post by Danijar Hafner [Haf], one of the key contributors behind many of the world models papers.

17 Imitation Learning

Defining complex tasks, such as driving a car, is often very challenging. In reinforcement learning, agents typically learn by receiving rewards that signal how good their actions are. However, in many real-world scenarios, especially those involving human skills like driving, designing an explicit reward function is difficult and often impractical. Humans, on the other hand, are naturally good at demonstrating how to perform such tasks.

The idea behind imitation learning is to infer the underlying policy by observing an expert, without access to a reward function. Therefore, the setting for imitation learning is as follows:

- set of trajectories $D = \{\tau_1, \dots, \tau_N\}$
- state action pairs $\tau_n = ((s_1, a_1), \dots, (s_{T_n}, a_{T_n}))$ generated by the expert
- transition/ system dynamics $P(s'|s, a)$
- underlying/observation policy $\pi_{\text{dem}}(a_t|s_t)$
- agent policy $\pi_\theta(a_t|s_t)$

17.1 Behavioural Cloning

One approach is to reduce the problem to a supervised learning task, where we aim to learn the mapping from states to actions. Specifically we want to minimize the Kullback–Leibler between the expert policy π_{dem} and learned policy π_θ . The objective can be written as:

$$\begin{aligned}\pi_\theta^* &= \arg \min_{\theta} \mathbb{E}_{D \sim s} [D_{\text{KL}}(\pi_{\text{dem}}(a|s) || \pi_\theta(a|s))] \\ &= \arg \min_{\theta} \mathbb{E}_{D \sim (s,a)} [-\log \pi_\theta(a|s)] + C\end{aligned}$$

Thus, the problem reduces to maximum likelihood estimation (MLE). However, there are a few notable issues with this approach.

17.1.1 Distributional Shift

The primary problem in respect to behaviour cloning is the distributional shift. Specifically, the network is trained only on the dataset D from the expert, but at test time, it receives its own potentially incorrect predictions as inputs (due to stochastic policy, stochastic system dynamics etc.). This leads to the agent entering regions of the state space it hasn't learned from, due to small compounding errors. As a result, for states it has never observed, the agent doesn't know how to behave, leading to poor policies. The core issue is that the distribution of inputs the model encounters at test time differs from the distribution it was trained on (distributional shift).

A solution to this problem is an algorithm called **D**Agger (Dataset Aggregation). DAgger starts similarly to standard behavior cloning: we train a policy using expert demonstrations. The key difference is that, after we deployed it in the environment to collect new data to train on we query the expert to label the correct actions that should have been taken in each state. These corrected state-action pairs are added to the dataset, and the policy is retrained on the aggregated data. This process is repeated iteratively, gradually improving the policy and reducing distributional shift.

Algorithm 23 DAgger

```

for  $k = 1 \dots n$  do
  train policy from data  $D \rightarrow \pi_\theta$ 
  run policy to collect new data  $\pi_\theta \rightarrow D_\theta$ 
  where necessary ask expert for advice  $D_\theta \rightarrow D_\theta \oplus a_{\text{expert}}$ 
  Aggregate Data  $D \cup D_\theta \rightarrow D$ 
end for
```

However, the obvious problem with this approach is that collecting additional training data is very expensive. Beyond that, there are two other challenges that make it difficult to learn the optimal policy.

Markovian System Another problem preventing us from learning good policies is the assumption that the system is Markovian. In many cases, the states we’ve visited before have a significant impact on our decision-making at the current time step. This means humans do not act in a purely Markovian way. A simple solution to this is to use non-Markovian policies, which we could implement with Recurrent Neural Networks (RNNs).

Multimodal Behaviour Yet another problem is multimodal behaviour. Consider the case of flying a drone, where the dataset includes a state where the drone must fly around a tree, either to the left or to the right. In the dataset, half the time we go left and the other half, we go right. For a discrete policy, after training, it might give a 50% probability of going right or left. However, for a continuous policy, regression will average out the decision between going left or right, resulting in suboptimal behaviour. One solution to this is to output a mixture of Gaussians to account for the multimodal nature of the decision.

Even with these improvements, the fact remains that expert demonstrations provide an upper bound for behavioural cloning. This means that we are only as good as the expert. However, if we could understand the intention behind the expert’s actions, we might be able to develop better policies beyond simply mimicking the expert.

17.2 Inverse Reinforcement Learning

In Inverse Reinforcement Learning (IRL), the goal is to infer the reward function from a given set of expert demonstrations so that the expert policy is optimal under this reward function, and then use this reward function to learn a policy. However, there are two main issues with this approach:

1. the demonstrations could have been produced by many policies
2. there are many reward functions for which the expert’s policy is optimal

In the following we are going to look at some methods which try to implement inverse reinforcement learning.

17.2.1 Feature Matching

In the early days, researchers viewed reward functions as being based not directly on states and actions, but rather on features f of these states and actions:

$$r_\psi(s, a) = \sum_i \psi_i f_i(s, a) = \psi^T f$$

In reinforcement learning (RL), features are numeric values that represent the state (and sometimes the action) of the environment in a form that supports learning. These values act as inputs to the learning algorithm, helping the agent understand its situation and make decisions.

For example, in the CartPole-v1 environment, the state is described by four features: the cart’s position and velocity, the angle of the pole, and the velocity at the pole’s tip. These features capture everything the agent needs to know to learn the task. While early RL relied on handcrafted features, modern approaches especially those using deep learning often learn features directly from raw sensor data. However, not all RL algorithms require explicit features. Tabular methods can work with small, discrete state spaces where each state is just an index. But for large or continuous environments, function approximators like neural networks are used, and these require meaningful features to work effectively.

The core idea is now that if the right set of features is chosen, we can reproduce expert-like behaviour by ensuring that the agent matches the expected values of those features. In other words, by aligning the feature expectations of the agent with those of the expert we can induce similar behaviour. Formally the objective becomes finding a parameter vector ψ such that:

$$\mathbb{E}_{\pi^{r_\psi}}[f(s, a)] = \mathbb{E}_{\pi^*}[f(s, a)]$$

However, there is still the issue of ambiguity, as different ψ vectors can result in the same expected feature values.

17.2.2 Max-Entropy Inverse Reinforcement Learning

Max-Entropy Inverse Reinforcement Learning tries to address the problem of the ambiguity in the reward functions by using the principle of maximum entropy. The principle of maximum entropy states that the probability distribution which best represents the current state of knowledge about a system is the one with the largest entropy. The entropy $H(X)$ is defined as

$$H(X) = - \sum_{x \in X} p(x) \log p(x)$$

The objective is now to maximize the entropy $H(X)$, while ensuring matching the feature count and ensuring that p is a proper distribution over the expert data set D (τ are the trajectories)

$$\arg \max_{\psi} - \sum_{\tau \in D} p_{\psi}(\tau) \log p_{\psi}(\tau) \quad \text{s.t.} \quad \sum_{\tau} p_{\psi}(\tau) f(\tau) = \frac{1}{|D|} \sum_{i \in D} f(\tau_i) \text{ and } \sum_{\tau} p_{\psi}(\tau) = 1$$

The solution for this constrained problem can be derived using Lagrange multipliers, which leads to the solution:

$$p_{\psi}(\tau;) \propto \frac{1}{Z(\psi)} \exp r_{\psi}(\tau) \text{ with } Z(\psi) = \sum_{\tau} \exp r_{\psi}(\tau)$$

Under this principle good behaviour is exponentially more likely. Now, we have a probability distribution over trajectories which depends on the reward model, although we do not know the exact model (we do not know the right ψ). However, we know that the distribution is a normalized exponential distribution. This allows us to perform log-maximum likelihood estimation on the parameters of the reward model.

$$\begin{aligned} \max_{\psi} \log \prod_{\tau \in D} p_{\psi}(\tau) &= \max_{\psi} \log \prod_{\tau \in D} \frac{\exp r_{\psi}(\tau)}{\sum_{\tau^*} \exp r_{\psi}(\tau^*)} \\ &= \max_{\psi} \sum_{\tau \in D} \left(\log \exp r_{\psi}(\tau) - \log \sum_{\tau^*} \exp r_{\psi}(\tau^*) \right) \\ &= \max_{\psi} \sum_{\tau \in D} \log \exp r_{\psi}(\tau) - |D| \log \sum_{\tau^*} \exp r_{\psi}(\tau^*) \\ &= \max_{\psi} \sum_{\tau \in D} r_{\psi}(\tau) - |D| \log \sum_{\tau^*} \exp r_{\psi}(\tau^*) \end{aligned}$$

Now taking the derivative of with respect to ψ yields:

$$\begin{aligned} \nabla_{\psi} \left(\sum_{\tau \in D} r_{\psi}(\tau) - |D| \log \sum_{\tau^*} \exp r_{\psi}(\tau^*) \right) &= \sum_{\tau \in D} \frac{d}{d\psi} r_{\psi}(\tau) - |D| \frac{d}{d\psi} \log \sum_{\tau^*} \exp r_{\psi}(\tau^*) \\ &= \sum_{\tau \in D} \frac{d}{d\psi} r_{\psi}(\tau) - |D| \underbrace{\frac{1}{\sum_{\tau^*} \exp r_{\psi}(\tau^*)} \sum_{\tau^*} \exp r_{\psi}(\tau^*)}_{p_{\psi}(\tau^*)} \frac{d}{d\psi} r_{\psi}(\tau^*) \\ &= \sum_{\tau \in D} \frac{d}{d\psi} r_{\psi}(\tau) - |D| \sum_{\tau^*} p_{\psi}(\tau^*) \frac{d}{d\psi} r_{\psi}(\tau^*) \\ &= \sum_{s \in \tau \in D} \frac{d}{d\psi} r_{\psi}(s) - |D| \sum_s p_{\psi}(s) \frac{d}{d\psi} r_{\psi}(s) \end{aligned}$$

The only remaining component is the state occupancies/frequencies $p_{\psi}(s)$. In the tabular setting, and assuming the dynamics are known, we can compute these frequencies using dynamic programming, provided we have already obtained the policy.

Algorithm 24 Computing state frequencies

```

{Goal is to compute  $p_\psi(s)$ }
 $\{\mu_t(s)$ : prob of visiting  $s$  at timestep  $t\}$ 
 $\mu_1(s) = p(s_1 = s)$  {initial state distribution of the MDP}
for  $t = 1 \dots T$  do
     $\mu_{t+1}(s) = \sum_a \sum_{s'} \mu_t(s') \pi(a|s') P(s|s', a)$ 
end for
 $p_\psi(s) = \sum_t \mu_t(s)$ 

```

The whole max entropy inverse reinforcement learning algorithm is defined as :

Algorithm 25 Max Entropy Inverse Reinforcement Learning

```

initialize  $\psi$  and gather demonstrations  $D$ 
for  $m = 1 \dots M$  do
    solve for optimal policy  $\pi(a|s)$  w.r.t reward  $r_\psi$  (i.e. value iteration)
    solve for state frequencies  $p_\psi(s)$  through algorithm 24
    compute gradient  $\nabla_\psi \mathcal{L} = \sum_{s \in \tau \in D} \frac{d}{d\psi} r_\psi(s) - |D| \sum_s p_\psi(s) \frac{d}{d\psi} r_\psi(s)$ 
     $\psi \leftarrow \psi + \alpha \nabla_\psi \mathcal{L}$ 
end for

```

However, there are some caveats to this approach. The first issue is that we need to compute the optimal policy at every iteration. This is already somewhat cumbersome but in addition to that we saw in the Dynamic Programming chapter that this process breaks down in high- dimensional state and action spaces. Also computing the state frequencies and the gradient becomes infeasible or too slow in such high- dimensional settings. An improved version of this algorithm, which also works in settings where the dynamics are unknown and can be applied to large or continuous state and action spaces, can be found in [FLA16].

17.2.3 Adversarial Methods

The idea behind adversarial methods is to train a network to produce samples that appear to come from the expert demonstration $p^*(x)$. One way to achieve this is through the Generative Adversarial Network (GAN) architecture, where we have a generator $p_\theta(x|z)$ that produces "fake" samples, while simultaneously training a discriminator $D_\psi(x) = p_\psi(\text{real data}|x)$ to distinguish between data generated by the generator and data from the expert. The objective is for the generator to become so good that the discriminator can no longer tell the difference. The training objectives are then defined as:

$$\begin{aligned}
 \psi &= \arg \max_{\psi} \frac{1}{N} \sum_{x \sim p^*} \log D_\psi(x) + \frac{1}{M} \sum_{x \sim p_\theta} \log (1 - D_\psi(x)) \\
 \theta &= \arg \max_{\theta} \mathbb{E}_{x \sim p_\theta} [\log D_\psi(x)]
 \end{aligned}$$

Inverse Reinforcement Learning can also be framed as a two-player game, similar to the setup in GANs. In this analogy, the expert represents the real/true data, while the policy we want to learn plays the role of the generator. The discriminator in this case is the reward function which aims to assign high rewards to trajectories from the expert and low rewards to those generated by the learned policy, essentially trying to distinguish human-like behaviour from the learned one. The policy is then optimized to produce trajectories that make it increasingly difficult for the discriminator (reward function) to tell them apart from the expert's, thereby improving its performance. The objectives for the generator and discriminator remain largely the same as in the standard GAN framework, with a few key differences. The generator (the policy) is updated using policy gradient methods. If we have access to the reward function, this update can take a form similar to the following:

$$\nabla_\theta \mathcal{L} \approx \frac{1}{|\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) r_\psi(\tau)$$

The objective for the discriminator also stays the same, but there are two different ways to construct $D_\psi(x)$ which we will explore in the following.

Adversarial Imitation from Reinforcement Learning (AIRL) It has been show that for a generator with density $p_\theta(\tau)$, the optimal discriminator is the following:

$$D^*(\tau) = \frac{p^*(\tau)}{p_\theta(\tau) + p^*(\tau)}$$

From Section 17.2.2, we know that in the maximum entropy framework, the expert trajectory distribution $p^*(\tau)$ follows an exponential form over rewards. Using this, we can rewrite the discriminator as:

$$D_\psi(\tau) = \frac{\frac{1}{Z} \exp r_\psi(\tau)}{p_\theta(\tau) + \frac{1}{Z} \exp r_\psi(\tau)}$$

This formulation has a valuable property: we can explicitly recover a reward function $r_\psi(\tau)$ from the discriminator. This is useful because it allows us to transfer the learned reward to new settings and re-optimize the policy under different conditions. A detailed derivation and explanation of AIRL can be found in [FCAL16].

Generative Adversarial Imitation Learning (GAIL) In GAIL [HE16], the discriminator is implemented as a simple binary classifier that distinguishes between expert and generated trajectories. While this formulation is easier to optimize than AIRL, it does not allow us to recover an explicit reward function. As a result, GAIL is no longer a true inverse reinforcement learning method—it is purely imitation learning, focused on mimicking expert behavior without modeling the underlying reward structure.

A very general adversarial framework for inverse reinforcement learning might look something like this:

Algorithm 26 Generic AIRL

```

for  $i = 1, 2, \dots$  do
    sample new set of agent trajectories from current policy
    update discriminator using demonstrations and new agent trajectories
    do a RL policy step (like TRPO) with cost based on discriminator
end for
```

17.3 Challenges of human data

Up to this point, we have assumed that the expert behaves optimally. However this assumption may not hold especially when dealing with human experts. Here are some common challenges when learning from human-generated data:

- **Variance in demonstrations:** Learning a single trajectory is not representative enough
- **Outliers in demonstrations:** Skews behaviour away from intended demonstrations
- **Multimodality in demonstrations:** Mode-averaging leads to bad behaviour
- **Context dependency in demonstrations:** Behaviour should also depend on context
- **Context surplus in demonstrations:** Learning process should ignore irrelevant contexts
- **Sub-optimality in demonstrations:** Learning process should cope with suboptimal data
- **Weak annotation of demonstrations:** Learning process should cope with weak annotations
- **Cumbersome data collection:** Learning process should facilitate effective data collection

In the following, we will explore some approaches that attempt to incorporate suboptimal behaviour.

17.3.1 Trajectory-ranked Reward EXtrapolation (T-REX)

In T-REX [BGNN19] we are given a data set $D = \{\tau_1, \dots, \tau_n\}$ where each trajectory is ordered such that if $i < j \rightarrow \tau_i \prec \tau_j$ meaning that trajectory τ_i is considered to be worse than τ_j . The trajectories consist only of sequences of states (no actions) $\tau_i = (s_0^i, s_1^i, \dots, s_t^i)$.

The objective of T-REX is to learn a reward function $r_\theta(s)$ such that the sum of rewards over the states in trajectory τ_i is less than the sum of rewards over the states in trajectory τ_j whenever $\tau_i \prec \tau_j$. Mathematically:

$$\sum_{s \in \tau_i} r_\theta(s) < \sum_{s \in \tau_j} r_\theta(s) \iff \tau_i \prec \tau_j$$

To achieve this, the objective is to minimize the cross-entropy loss $-\mathbb{E}_{x \sim p}[\log q(x)]$ where in our case x are the pair of trajectories (τ_i, τ_j) and $q(x)$ will be modelled via the softmax function resulting in:

$$L(\theta) = - \sum_{\tau_i \prec \tau_j} \log \frac{\exp \sum_{s \in \tau_j} r_\theta(s)}{\exp \sum_{s \in \tau_j} r_\theta(s) + \exp \sum_{s \in \tau_i} r_\theta(s)}$$

Once the reward function is learned, it can be used with algorithms like PPO to learn a policy.

However, a limitation of this approach is the need for many labeled demonstrations, which can be cumbersome. A modification is to allow the agent to provide its own demonstrations and use sparse rewards as a proxy for ranking the trajectories.

17.3.2 Goal-Conditioned Imitation Learning

In earlier approaches such as behaviour cloning, the objective was to learn a policy $\pi(a|s)$ that could achieve a specific goal by mimicking expert demonstrations. However, to make this method robust, a large amount of expert data was required. This could be achieved either by starting with a "good enough" data set or by gradually augmenting the dataset with the agent's own trajectories, corrected by the expert as done in methods like DAgger. However, this process is often expensive and resource intensive.

In Goal-Conditioned Imitation Learning the expert data consists of trajectories for multiple different goals. Instead of learning a policy $\pi(a|s)$, we condition it on both the state and the goal, i.e., $\pi(a|s, g)$, enabling the policy to generalize across different targets. This goal diversity leads the expert to visit a wider range of states, providing richer coverage of the state space. As a result, the policy learns from a broader set of experiences, improving its ability to generalize to new goals. The basic concept looks as follows:

1. Collect "arbitrary" demonstrations
2. Create tuples of the form $(s_t, a_t, s_{\text{goal}})$, where s_{goal} is the final state of the demonstration to which the state-action pair (s_t, a_t) belongs.
3. Create Dataset which consist of these tuples
4. Learn a goal-conditioned policy

Learning Latent Plans from Play Learning Latent Plans from Play [LKX⁺19] is a method based on Goal-Conditioned Imitation Learning that leverages play data to learn generalizable control policies. Play-LMP is a unified model that self-supervises from play data and operates in three main steps:

1. **Sampling:** It samples random temporal windows from a large memory of play data.
2. **Latent Plan Learning:** It learns a latent plan space that captures the diverse behaviours seen in the play data.
3. **Policy Training:** It trains a policy conditioned on the current state, goal state, and a sampled latent plan to reconstruct the actions within the selected window.

The latent plan space is learned using two stochastic encoders:

- **Plan Proposal Network:** Takes only the initial and final state (goal) and outputs a distribution over possible latent plans that could connect the two.
- **Plan Recognition Network:** Processes the full state-action sequence and outputs a latent distribution representing the specific behavior that was executed.

To ensure that the proposed plans are consistent with the behaviours observed during play, the KL divergence between the two latent distributions (recognition and proposal) is minimized. This encourages the Plan Proposal Network to assign high probability to behaviours that actually occurred in the play data. Additionally there is also an **Action Decoder Network** which receives the current state, the final goal of the sequence, and the sampled latent plan vector as input, and it outputs the next action. All of these networks are Conditional Variational Autoencoders (CVAEs).

At inference time, only the Plan Proposal and Action Decoder Networks are used. For a more detailed explanation, you can visit the [website](#)




Figure 19: Architecture used for the Learning from play method

This approach has several key advantages:

- It is cheap: Large amounts of play data can be collected quickly since it doesn't require scene staging, task segmenting, or resetting to an initial state.
- It is general: It encompasses both functional and non-functional behavior, relaxing the need for a predefined task distribution.
- It is rich: Play naturally involves repeated, varied behavior, leading to high coverage of the possible interaction space. Additionally, the data is unlabeled, multimodal, and suboptimal.

17.4 Resources

There are many excellent resources on imitation learning and reinforcement learning. Much of the content presented here is based on Sergey Levine's "CS 285: Lecture 2 – Imitation Learning" and "Lecture 20 – Inverse Reinforcement Learning" [LSM⁺20, RAI20]. Additional resources that helped deepen my understanding of feature matching include "Stanford CS234 – Reinforcement Learning — Offline RL 1 — 2024, Lecture 8" [Onl24], and "Maximum Entropy Inverse RL and Adversarial Imitation Learning" by Katerina Fragkiadaki [Fra19]. Another valuable reference on adversarial approaches is "Deep RL Bootcamp – Lecture 10B: Inverse Reinforcement Learning" [Pri18].

18 Offline Reinforcement Learning

Offline Reinforcement Learning is another paradigm of Reinforcement Learning. All previous methods we looked at are considered to be Online-RL methods. In Online-RL, the agent gathers data by directly interacting with the environment. It then uses this experience immediately or via a replay buffer to update its policy, after which it has to interact with the environment again to collect more data.




Figure 20: Visualisation of online-RL as on-policy (left) and off-policy (right), inspired from [LKTF20]

Offline-RL is inspired by general deep learning. There we are just given a model and a large dataset to learn on. Throughout the training, the model does not interact in any way with anything outside the given data. But they are still able to generalise very well. In Offline-RL, the agent only uses data collected from other agents and/or human demonstration to learn a policy, while not interacting with the environment once during training.




Figure 21: Visualisation of Offline Reinforcement Learning, inspired from [LKTF20]

In the following, we denote by D the static dataset of transitions, defined as $D = \{(s_i, a_i, s'_i, r_i)\}$. The state-action tuples in this dataset are sampled such that the states s are drawn from the stationary distribution $d^{\pi_\beta}(s)$ of a Markov chain (the distribution that remains unchanged over time under the dynamics induced by a policy). The actions a are sampled from the behaviour policy $\pi_\beta(a | s)$, which is generally unknown.

The question that arises is whether we can hope to learn good policy without ever interacting with the environment. While we won't formally prove it, the hope/claim is that Offline RL can uncover the "good stuff" from a dataset that contains a mix of both good and bad behaviours. Importantly, the "good stuff" isn't limited to simply identifying the best trajectories in the dataset, but also includes the potential to learn policies that outperform any single behaviour present in the data. This is based on the hope that Offline RL can generalize: that observing good behaviour in one context can inform good decisions in others. Ideally, the agent would be able to "stitch together" different pieces of good

behaviour from across the dataset to construct an even better overall policy.




Figure 22: Visualization of the stitching process in Offline Reinforcement Learning. In the dataset, we originally observe only partial trajectories—specifically, the segments from $A \rightarrow D$ and $D \rightarrow E$ (as shown on the left). The stitching procedure combines these separate trajectories to construct a new, potentially optimal path from A to the goal E , even though this complete path was never explicitly observed in the data. Inspired from [AK22]

At first glance, Offline RL might seem quite similar to Imitation Learning (IL), since both involve learning from precollected datasets. This raises the question: can stitching also occur in IL methods? The key difference between the two settings lies in the availability of reward signals. Offline RL has access to reward information, which enables the agent to distinguish between good and bad segments of behaviour. This allows it to selectively stitch together high-quality parts of trajectories to construct better policies.

In contrast, Imitation Learning typically does not have access to rewards. It relies solely on mimicking the observed behaviour, making it difficult to identify which parts of a trajectory are optimal or suboptimal. As a result, IL may end up stitching together poor behaviours or failing to prioritize the best ones.

18.1 Challenges in Offline-RL

One problem that Offline-RL suffers from is overestimation. In [KF^{TL}19] they did an experiment where they collected data from an agent that had been trained with some sort of SAC and had achieved reasonably good results. They used the data generated by this agent to fill a replay buffer which then was used for doing Offline-RL on the same task. As you can see in figure 23, the actual return is really bad while the Q values are very high. So the Offline-RL agent is certain that its actions will lead to high reward but in reality the opposite is true. In the following, we will explore from where this problem arises and discuss potential strategies to mitigate it.




Figure 23: Performance of SAC on HalfCheetah task with off-policy expert data w.r.t. number of training samples (n), from [KF^{TL}19]

18.1.1 Intuitive Explanation

The intuitive way to understand the challenge of Offline RL is through counterfactuals. Counterfactuals are if statements where the if part is not realised. For example, you are on your way home and arrive at a fork in the road where you can either go left or right. Both ways lead to your home. So you decide to go left, thinking it would be faster, only to learn that you will get stuck in traffic and arrive home late. The counterfactual is now "I should have went right", or in other words, "If I had went right, it would have been better", even though you do not know the outcome of going right because you never took it. The core problem is that you are making assumptions about the outcomes of actions based solely on the outcomes of the actions you have taken without ever observing what would have happened had you chosen differently.

In Offline RL, such reasoning is necessary. Because for Offline RL to produce policies that outperform the behaviour seen in the dataset it must consider actions beyond those explicitly observed in the data. So when updating/executing the policy/Q-values, we typically compare all possible actions in a given state to identify the best one. This inherently involves counterfactual reasoning. Because the action we believe to be optimal, because it has the highest predicted Q-value, might never been actually observed in the dataset. If we select it anyway, we are relying entirely on predictions, with no evidence to back them up and so doing counterfactual reasoning.

However, counterfactual reasoning in Offline RL brings about a deeper, more subtle problem: distributional shift which we already saw when we looked at Imitation Learning and Model-Based-Reinforcement-Learning. Offline RL inherently involves making predictions about what would happen under a different policy than the one that generated the data. So it seeks to learn a new policy that deviates from the behavior policy ideally to achieve better performance which naturally leads to a shift in the distribution of visited states and actions. This shift manifests in two key ways. First, the learned policy may visit parts of the state space that were rarely or never seen in the dataset. Second, even in familiar states, it may prefer different actions than those taken by the behavior policy (see above explanation). In both cases whether policies, value functions, or models—are forced to make predictions under conditions different from those they were trained on (distributional shift).

Unlike online RL methods, which can try the action and learn from its outcome, offline methods do not have this luxury. As a result, Offline RL must find ways to handle uncertainty around unseen actions (out-of-distribution (OOD)) typically by regularizing, constraining, or penalizing those actions during learning.

18.1.2 Mathematical Explanation

We now turn to the mathematical explanation of the problem. In general machine learning, our objective is typically framed as risk minimization:

$$\theta \leftarrow \arg \min_{\theta} \mathbb{E}_{x \sim p(x), y \sim p(y|x)} [(f_{\theta}(x) - y)^2]$$

After training, we might ask: Given a new input x^* , is the model's prediction $f_{\theta}(x^*)$ likely to be accurate? In expectation, the error is low but only if $x^* \sim p(x)$ so if it comes from the same distribution as the training data. If x^* is drawn from a different distribution, this guarantee no longer holds.

Moreover, if we move away from evaluating under expectation and consider individual inputs, we lose even more certainty. We can not guarantee low error, even if x^* was from the training set. This can be easily exploited if we choose $x^* \leftarrow \arg \max_x f_{\theta}(x)$.

In methods like Q-Learning- or Actor-Critic-methods, we compute the target values y under some policy π_{new} and try to regress them onto the Q-values

$$Q^{k+1} = \arg \min_Q \mathbb{E}_{s,a,s' \sim D \sim \pi_{\beta}} [(Q(s,a) - y(s,a,s'))^2]$$

with $y(s,a,s') = r(s,a) + \mathbb{E}_{a' \sim \pi^k} [Q^k(s',a')] \quad (\text{see footnote } 5)$

Similar to above, we expect the Q values to be accurate if $\pi^k(a|s) = \pi_{\beta}(a|s)$. But the whole point of offline-RL is that $\pi^k \gg \pi_{\beta}$, and to make matters worse, in most methods we choose

$$\pi^{k+1} = \arg \max_{\pi} \mathbb{E}_{s \sim D, a \sim \pi^k(a|s)} [Q^{k+1}(s,a)]$$

so we are explicitly trying to find actions that produce high Q values.

18.2 Approaches to Offline-RL

The most common approach for addressing the distributional shift problem in Offline-RL is to impose constraints on how much the learned policy is allowed to deviate from the behaviour policy, leading to an optimization objective of the form:

$$\pi_{\text{new}} = \arg \max_{\pi} \mathbb{E}_{a \sim \pi(a|s)} [Q(s, a)] \quad \text{s.t.} \quad \text{Diff}(\pi || \pi_{\beta}) \leq \epsilon$$

This approach comes with two key challenges. First, the constraint may not be pessimistic enough. Even if two policies are close in terms of a divergence metric (like KL), they can still differ significantly at specific states. This means the learned policy might take actions in certain states where the Q-values are poorly estimated, leading to unreliable behaviour.

Second, if the constraint is too pessimistic, it can prevent the learned policy from meaningfully improving over the behaviour policy. In such cases, the policy is forced to stay so close to π_{β} that it cannot explore better actions, limiting its potential to exceed the performance of the dataset policy. For simplicity, in the following we always use the KL as the distance measure for the constraint, but it also can be any other measure for probabilities.

To better understand the second challenge, consider Figure 24. We are given the behaviour policy π_{β} (blue) for a single state. The blue dots represent samples from our data set where y values correspond to q values. The q-function fitted to these samples is given by the orange function. If we now train a policy that maximises the q values but also satisfies the constraint, we get π_{KL} (red). You can see that this policy does not only focus on the actions with the highest q-values, but also assigns non-negligible probabilities to actions with low q-values. The reason for this is that the KL constraint forces it to be something like π_{β} . Intuitively, we would want something like π_{support} (green), but that violates the constraint because the variance is so far off.

What comes closer to what we want but is more difficult to implement are support constraints

$$\pi(a|s) > 0 \iff \pi_{\beta}(a|s) \geq \epsilon$$




Figure 24: Visualization of explicit policy constraints using KL divergence. Inspired by CS 285: Lecture 16 [LSM⁺20, RAI20].

⁵ $y = r(s, a) + \mathbb{E}_{a' \sim \pi^k} [Q^k(s', a')]$ may seem unfamiliar at first, as we initially defined the targets as $r(s, a) + \gamma \max_{a'} Q^k(s', a')$. However, we can rewrite this expression as $r(s, a) + \gamma Q^k(s', \arg \max_{a'} Q^k(s', a'))$, which, using the relationship in Equation (3), leads us back to the same formulation.

18.2.1 Explicit Policy Constraints

One common approach to incorporating policy constraints in Offline RL is to solve the constrained optimization problem using the method of Lagrange multipliers. This leads to the following objective:

$$\begin{aligned}\pi_{\text{new}}(a|s) &= \arg \max_{\pi} \mathbb{E}_{a \sim \pi(a|s)} [Q(s, a)] - \lambda \text{KL}(\pi || \pi_{\beta}) \\ &= \arg \max_{\pi} \mathbb{E}_{a \sim \pi(a|s)} [Q(s, a) + \lambda \log \pi_{\beta}(a|s)] + \lambda H(\pi(a|s))\end{aligned}$$

Here, the Lagrangian multiplier λ can either be computed via dual optimization or treated as a hyperparameter. However, since we're not solving for the policy in closed form, we simply optimize the objective directly. An alternative idea is to incorporate the constraint directly into the reward function:

$$r_{\text{diff}}(s, a) = r(s, a) - \text{KL}(\pi || \pi_{\beta})$$

While this explicit approaches using KL constraints are simple to implement, they rely on access to the behaviour policy π_{β} which is often unknown in practice. One idea is to learn $\pi_{\beta}(a|s)$ from the data via a model. Another idea is to cleverly implement the constraint that we only need access to the samples, not the probabilities which we will looking at in the following.

18.2.2 Implicit Policy Constraints

When we once again express the constrained objective using the Lagrangian and this time solve for the optimal policy in closed form, we obtain:

$$\pi^*(a|s) = \frac{1}{Z(s)} \pi_{\beta}(a|s) \exp \frac{A^{\pi_{\text{old}}}(s, a)}{\lambda}$$

Here $Z(s)$ is the normalizing constant to ensure the result is a valid probability distribution, and λ is the Lagrangian multiplier. This optimal policy can now be approximated via maximum likelihood estimation, yielding the following objective:

$$\pi_{\text{new}}(a|s) = \arg \max_{\pi} \mathbb{E}_{(s,a) \sim \pi_{\beta}} \left[\frac{1}{Z(s)} \log(\pi(a|s)) \exp \frac{A^{\pi_{\text{old}}}(s, a)}{\lambda} \right]$$

Notably, this formulation does not require explicit evaluation of π_{β} only to draw samples from it. However, it still suffers from the problem of out-of-distribution actions, as computing advantages using Q-values requires querying actions that lie outside the behaviour policy's distribution. In the following we will look at methods which try to completely avoid outer distribution action in the Q-update and so preventing the overestimation issue.

18.2.3 Implicit Q-Learning

The update rule for the Q-function can be expressed using the value function as follows:

$$Q(s, a) \leftarrow r(s, a) + \gamma \mathbb{E}_{a' \sim \pi_{\text{new}}} [Q(s', a')] \stackrel{\text{eq. (3)}}{=} r(s, a) + \gamma V^{\pi_{\text{new}}}(s')$$

We assume that $V(s')$ is represented by a neural network. When we regress V onto $Q(s, a)$ using mean squared error (MSE) as we've done previously we end up approximating the Q-function corresponding to π_{β} , since our dataset is generated by π_{β} . Thus, we are effectively learning $Q_{\pi_{\beta}}$ rather than $Q_{\pi_{\text{new}}}$. An important observation is that, in large or continuous datasets, individual states are often visited only once, with typically a single action taken per state. However, nearby similar states may be visited with different actions. In practice, the policy generalizes across such states, meaning the targets are not fixed values but rather form a distribution (see Figure 25).




Figure 25: Probability distribution of the value function for a given state

When performing MSE regression, the model tends to predict the mean of the value distribution represented by the red dashed line. However, what we actually want is to focus on the upper quantile, where the value under the optimal policy is supported by the data (illustrated by the cyan dashed line). To achieve this Implicit-Q-Learning [KNL21] uses the expectile regression loss, defined as:

$$L_{\text{expectile}}(x) = \begin{cases} (1 - \mu)x^2, & \text{if } x > 0 \\ \mu x^2 & \text{else} \end{cases}$$

The purpose of this loss function is to penalize negative errors more heavily than positive ones, encouraging predictions of higher values. At first glance, this may appear to contradict the goal of avoiding overestimation. However, overestimation does not occur in this case because all the actions we take are drawn from within our dataset, ensuring that no out-of-distribution (OOD) actions are involved. As a result, we can update our Q- functions without the risk of out-of-distribution actions.

What can be shown is that this implicit Q-learning corresponds to support constraints:

$$V(s) \leftarrow \max_{a \in \Omega(s)} Q(s, a) \quad \text{with} \quad \Omega(s) = \{a : \pi_\beta(a|s) \geq \epsilon\}$$

18.2.4 Conservative Q-Learning

Conservative Q-Learning (CQL) [KZTL20] aims to mitigate overestimation in offline reinforcement learning by explicitly penalizing large Q-values, especially for out-of-distribution (OOD) actions. The core idea is to modify the Q-learning objective as follows:

$$Q^\pi = \arg \min_Q \mathbb{E}_{(s,a,s') \sim D} \left[(Q(s, a) - (r(s, a) + \mathbb{E}_\pi[Q(s', a')]))^2 \right] + \max_\mu \alpha \mathbb{E}_{s \sim D, a \sim \mu(a|s)} [Q(s, a)]$$

Here, the second term penalizes large Q-values under some auxiliary policy μ , encouraging conservatism by lowering Q- values for actions not present in the dataset D .

However, this approach can become overly pessimistic-reducing all Q-values, including those that are not overestimated. To address this, a corrective term is introduced that subtracts the expected Q-value over the dataset distribution. The full CQL loss becomes:

$$Q^\pi = \arg \min_Q \mathbb{E}_{(s,a,s') \sim D} \left[(Q(s, a) - (r(s, a) + \mathbb{E}_\pi[Q(s', a')]))^2 \right] + \underbrace{\max_\mu \alpha \mathbb{E}_{s \sim D, a \sim \mu(a|s)} [Q(s, a)]}_{\text{punish big Q-values}} - \underbrace{\alpha \mathbb{E}_{(s,a) \sim D} [Q(s, a)]}_{\text{unless they from dataset}}$$

18.3 Resources


This section is largely based on Sergey Levine’s CS 285 ”Lecture 15 – Offline Reinforcement Learning“ and ”Lecture 16 – Offline Reinforcement Learning 2“ [LSM⁺20, RAI20]. For a more in-depth comparison between Offline Reinforcement Learning and Imitation Learning, see [KHS22].

19 Transformers in Reinforcement Learning

Transformers, originally developed for sequence modeling in natural language processing, have attracted growing interest in reinforcement learning (RL) due to their ability to capture long-range dependencies and handle complex sequential data. In the following sections, we will first explore the core mechanisms of the transformer architecture, and then discuss how the standard transformer is adapted for use in RL. However, we will not delve deeply into specific RL papers that apply transformers. At a high level, the process of how an input is processed by a transformer can be broken down into a series of steps:

1. **Input Embedding:** Converts input tokens or data into dense vector representations. In NLP, this applies to words/tokens; in RL or vision tasks, states or images are embedded.
2. **Positional Encoding:** Adds position information to embeddings, allowing the transformer to capture sequence order, which it doesn't handle inherently.
3. **Self-Attention Mechanism:** Enables each token to attend to all others using learned query, key, and value vectors. Attention scores—based on query-key similarity—determine how much focus each token gives to others, producing a weighted sum of values.
4. **Feedforward Network:** Applies a fully connected neural network to each position independently for further transformation.
5. **Residuals & Layer Norm:** Residual connections and layer normalization follow attention and feedforward layers to stabilize training and preserve information.
6. **Stacked Layers:** Multiple layers refine the representations, allowing deeper abstraction and more complex understanding.
7. **Output:** Final outputs vary by task—e.g., softmax classification in NLP or decision-making in RL.

Figure 26: Transformer architecture from [VSP⁺23]



In the following, we will take a closer look at the individual steps, primarily focusing on language processing. However, it's important to note that transformers are not limited to language—they can also process other types of inputs, such as images. We focus on language processing here simply because it's a clear and intuitive way to understand the core steps of how transformers work.

19.1 Embedding

The first stage in training a transformer model involves converting raw input into vector representations—a process known as input embedding. A basic approach to representing words is one-hot encoding, where each word is assigned a binary vector the size of the vocabulary, with a single 1 indicating the presence of that word and 0s elsewhere. Although simple, this method has significant drawbacks: the vectors are high-dimensional, sparse, and fail to capture any semantic relationships between words.

To address these limitations, learned word embeddings are used. These are dense, low-dimensional vectors trained to capture semantic meaning, typically using unsupervised learning techniques. Words with similar meanings are mapped to similar vector representations. For example, the vectors for “queen” and “woman” will lie close in the embedding space due to shared contextual usage. Embeddings can also reflect relationships, such as vector arithmetic:

$$\text{vector}(\text{queen}) \approx \text{vector}(\text{king}) + \text{vector}(\text{woman}) - \text{vector}(\text{man})$$




Figure 27: Illustration of potential embeddings, Image from [LDHC22]

Importantly, these embeddings extend beyond representing words in isolation—they evolve through the layers of a transformer to reflect contextual meaning. For instance, the embedding of “mathematician” could shift to refer specifically to Leibniz, depending on the sentence context. While the initial embedding layer does not yet account for surrounding context, it provides a foundational semantic representation that is refined throughout the transformer architecture. (While there are several types of embedding techniques, such as Word2Vec and GloVe, we will not explore them in detail here.)

19.2 Positional Encoding

The position of words in a sentence plays a crucial role in understanding their meaning. Early models, such as Recurrent Neural Networks (RNNs), inherently capture word order through their sequential processing. However, the Transformer architecture does not process tokens sequentially, and therefore lacks an inherent sense of word order. To address this, positional encoding is introduced to inject information about the position of each token in the sequence. An effective positional encoding scheme should satisfy the following criteria:

- **Unambiguous:** Each position in the sequence should have a unique encoding.
- **Deterministic:** It is beneficial if the encoding follows a predictable mathematical pattern, allowing the model to learn or infer position-based relationships.
- **Distance-aware:** The encoding should enable the model to estimate the relative distance between two positions.
- **Generalizable:** The encoding should allow the model to handle longer sequences than those seen during training. Ideally, values should be bounded to avoid numerical instability.

The authors of the *Attention is All You Need* paper [VSP+23] define positional encoding using sinusoidal functions as follows:

$$f(t)_i = \begin{cases} \sin(\omega_k \cdot t), & \text{if } i = 2k \\ \cos(\omega_k \cdot t), & \text{if } i = 2k + 1 \end{cases} = \begin{bmatrix} \sin(\omega_1 \cdot t) \\ \cos(\omega_1 \cdot t) \\ \sin(\omega_2 \cdot t) \\ \vdots \\ \sin(\omega_{d/2} \cdot t) \\ \cos(\omega_{d/2} \cdot t) \end{bmatrix}_{d \times 1}$$

$$\text{with } \omega_k = \frac{1}{10000^{2k/d}}$$

Here, t denotes the position of a token in the sequence (e.g., the 1st, 2nd, or 3rd word), and d is the dimensionality of the embedding space. Each dimension of the positional encoding corresponds to a sinusoid with a different frequency.

This encoding is unique as long as the chosen frequencies span a range that exceeds the maximum sequence length L (where $t \in \{1, \dots, L\}$). Put simply, at least one dimension must have a sinusoid

whose period exceeds L , ensuring that the function values do not repeat and each position maps to a distinct representation. Since sine and cosine are deterministic functions, the positional encoding is also deterministic. Furthermore, the sinusoidal encoding generalizes well to sequences longer than those seen during training, thanks to the periodic nature of sine and cosine functions. Finally, this encoding is distance-aware. Because each dimension of the encoding uses a different frequency, the model can infer relative positions from the differences between encodings.

19.2.1 Binary Representation and Positional Encoding

The intuition behind sinusoidal positional encoding becomes clearer when we consider the binary representation of numbers. At first glance, one might wonder how a combination of sine and cosine waves can meaningfully represent position or order. The answer lies in how binary counting works.




Figure 28: Illustration of the 4-bit binary number system. On the left is the binary representation in tabular form, and on the right is a visualization using periodic step functions (square waves).

So, we can use periodic wave functions (one for each dimension) to encode numbers. This idea closely resembles the sine and cosine functions used in positional encoding. Using discrete binary values directly would be inefficient in the continuous space of neural networks. Instead, we use their smooth, differentiable counterparts: sinusoidal functions. These functions act as continuous analogues of alternating bits. By varying their frequencies we mimic the behaviour of binary bits flipping at different rates. This allows the model to infer both absolute position and relative distance, all within a compact and expressive representation.




Figure 29: Visualization of positional encodings for positions 0 to 49 with an encoding dimension of 120. Each row represents the positional encoding for one position, and each column corresponds to a specific dimension with its associated value.

19.3 Attention Mechanism

In a transformer model, the key objective is for embeddings (vectors) to not only represent individual words but also capture context, allowing for a more nuanced understanding of language. This is achieved through the self-attention mechanism. To understand this, let's first break down a simple example and build on it. Let's consider a sentence like "Kings made laws to keep power over the small towns" and focus on how the attention mechanism works. Each word is initially represented as an embedding vector $\vec{E}_i$.

19.3.1 Query, Key, and Value Vectors

From each embedding $\vec{E}_i$, three vectors are computed using learned weight matrices:

- $\vec{Q}_i = W_Q \vec{E}_i$ (Query)
- $\vec{K}_i = W_K \vec{E}_i$ (Key)
- $\vec{V}_i = W_V \vec{E}_i$ (Value)

These vectors will determine how a token attends to others in the sequence.

19.3.2 Attention Scores

The next step is to measure the relevance of the different tokens in relation to each other. This is done by computing the dot product between the query vector of one word and the key vectors of all other words. The dot product gives a measure of similarity: high positive value indicates that the two words are more relevant to each other in the context (vectors pointing in nearly the same direction) while negative values indicate that the words are not so relevant to each other (vectors nearly pointing in opposite directions).

19.3.3 Softmax

After computing the dot product, we apply a softmax operation to the values. This converts the raw scores into probabilities, allowing the model to focus more on the most relevant words and ignore less relevant ones.

Additionally, in some settings (e.g., in reinforcement learning), we may mask certain values in the attention matrix, ensuring that later words do not influence earlier ones. This is done by setting specific values to negative infinity before applying softmax, which effectively eliminates their contribution.

19.3.4 Updating the Embedding

Now that we know which words are important for each other, the next step is to update the embeddings. The update is based on the value matrix, which is another learned weight matrix. The value matrix is multiplied by each word's embedding, and the resulting vectors are summed to produce the final output.

The value matrix answers the question: If this word is relevant to adjusting the meaning of another word, what specific adjustments should be made? This results in a new set of embeddings that reflect the context. The whole attention procedure can be described via the function

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

To understand why the attention scores are scaled by $\sqrt{d_k}$, consider the dot product between a query vector $\vec{q}_i$ and a key vector $\vec{k}_j$:

$$\sigma_{ij} = \vec{q}_i \cdot \vec{k}_j = \sum_{n=1}^{d_k} q_{in} k_{jn}$$

If we assume that the components of $\vec{q}_i$ and $\vec{k}_j$ are independent, have zero mean, and unit variance (we can make the assumption due to the use of LayerNorm), then each term in the sum $q_{in} k_{jn}$ will

have mean zero and variance 1. Therefore, the variance of the sum (i.e., the dot product) becomes:

$$\text{Var}(\sigma_{ij}) = \sum_{n=1}^{d_k} \text{Var}(q_{in} k_{jn}) = d_k \longrightarrow \text{std}(\alpha_{ij}) = \sqrt{d_k}$$

This means the standard deviation of the attention score grows with the square root of the key dimension. The normalization ensures that the variance of the attention scores remains approximately constant, regardless of dimensionality, leading to more stable gradients and more effective learning.




Figure 30: Illustration of the attention patterns from the first attention head in the first layer of GPT-2, for the input text: "Kings made laws to keep power over the small towns." The left plot shows the raw attention scores (QK^T), while the right plot displays the masked and scaled attention weights after applying the softmax function, i.e., $\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$.

19.3.5 Multi-Head Attention

The process described above outlines a single "head" of attention, where each word attends to other words based on their relevance. However, in practice, we use multi-head attention, which runs multiple attention mechanisms in parallel, each with its own learned query, key, and value matrices.

The benefit of multi-head attention is that it allows the model to focus on different aspects of the input simultaneously. For instance, one head might focus on syntactic relationships, while another might focus on semantic meanings, and yet another could capture long-range dependencies. This parallel processing allows the transformer to learn richer, more nuanced representations.

To make this process more efficient, each attention head typically uses smaller key, value, and query matrices. After the attention outputs are computed using these smaller matrices, the results from all heads are concatenated and passed through a linear layer to combine them. The formula for a Multi-Head attention block with k heads is given by

$$\text{MultiHead}(Q, K, V) = \text{Concat}[\text{head}_1, \dots, \text{head}_k] W_O$$

where W_O is the learnable output projection matrix of dimension $(k \cdot d_h) \times d_{\text{model}}$ with d_h being the dimensionality of the value vectors (the same as the dimensionality of each head's output) and d_{model} the final desired output dimensionality of the multi-head attention (usually equal to the model's embedding size).




Figure 31: Illustration of the attention patterns from all heads in the first layer of GPT-2 for the sentence: *"Kings made laws to keep power over the small towns."*

19.4 Feedforward Neural Network and Add & Norm

The Feedforward Neural Network (FNN) in Transformers is a simple multi-layer perceptron (MLP) composed of fully connected layers with ReLU activations.

The "Add" operation refers to residual (skip) connections, which are critical for stabilizing deep networks. Without them, models can suffer from vanishing or exploding gradients, making training inefficient or unstable. Residual connections allow gradients to flow more directly through the network and help mitigate the degradation problem, where adding more layers leads to worse performance.

Normalization further stabilizes training by ensuring that inputs to each layer are centered and appropriately scaled. While Batch Normalization (BN) normalizes across the batch dimension, it is less suitable for Transformer architectures, which often work with variable-length sequences and small batch sizes. Instead, Transformers use Layer Normalization (LN), which normalizes across feature dimensions for each individual sequence. This approach is more robust in the context of sequence modeling.

19.5 Modifications to the Vanilla Transformer

While the original transformer architecture has demonstrated strong performance across a range of tasks, it also faces certain limitations particularly when modeling long sequences. In this section, we examine follow-up work that addresses these issues.

19.5.1 Transformer-XL

Transformers are capable of learning long-term dependencies, but they are limited by a fixed-length context when applied to language modelling tasks. Transformer-XL [DYY⁺19] proposes an architecture that extends this capability by enabling the learning of dependencies beyond a fixed length, without disrupting temporal coherence. This is achieved by incorporating a memory mechanism into the transformer model.

The approach begins by splitting the input into segments, denoted as X_i . If we were to process these segments using a standard transformer model, there would be no attention between them, as illustrated in Figure 32.

To overcome this, Transformer-XL concatenates the hidden states from the previous segment with the current input at every transformer layer. These past states are treated as fixed memory i.e.

gradients are not propagated through them allowing efficient reuse without increasing memory cost during backpropagation. This approach is illustrated in Figure 33.




Figure 32: Illustration of the vanilla transformer model with a segment length 4. [DYY⁺19]




Figure 33: Illustration of the Transformer-XL model with a segment length 4. [DYY⁺19]

However, this modification introduces a new challenge: absolute positional encodings used in vanilla transformers cannot correctly handle segment-level continuity. For example, tokens X_1^1 (from one segment) and X_1^2 (from the next) would receive identical positional encodings, despite being temporally distinct. This breaks the sequence order and confuses the model.

To solve this, Transformer-XL replaces absolute positional encodings with relative positional encodings. This allows the model to compute attention based on the relative distance between tokens, preserving the temporal relationships across segments. For implementation details, refer to Section 3.3 in the original paper [DYY⁺19].

19.5.2 Stabilizing Transformers for Reinforcement Learning

Despite their success in NLP and vision, vanilla transformers and their memory-augmented variants like Transformer-XL initially performed poorly in reinforcement learning (RL) settings. For instance, [MRCA18] found that standard transformers failed on even basic bandit problems and tabular Markov Decision Processes (MDPs), suggesting that transformers may not be well-suited for RL without further adaptation.

To address this, [PSR⁺19] proposed modifications that significantly improve transformer stability and performance in RL. Their work builds on Transformer-XL and introduces two key changes.

1. TrXL-I: Input-Normalized Transformer-XL The first modification is reordering the application of layer normalization. Instead of applying it before and after each submodule (as in the original transformer), it is applied only to the inputs. This reordering creates an identity map from the input of the transformer at the first layer to the output at the last layer, unlike the canonical transformer, which applies multiple layer normalization operations that non-linearly transform the state encoding. TrXL-I shows a significant improvement in stability and performance over TrXL. One hypothesis for this improvement is that, with the initialization of submodules producing values near zero, the state encoding is passed untransformed to the policy and value heads. This allows the agent to learn a Markovian policies first e.g. $\pi(a|s_1, \dots, s_n) \approx \pi(a|s_n)$ before gradually incorporating longer-term memory, which aligns with how agents typically learn: basic behaviors precede memory-dependent strategies.[PSR⁺19]

2. GTrXL: Gated Transformer-XL To further stabilize learning, the authors replace standard residual connections with gated residual layers, inspired by gating mechanisms in LSTMs. This results in the Gated Transformer-XL (GTrXL) architecture. Gating enables dynamic control over how much information is retained or suppressed at each layer, improving both optimization and generalization. For a detailed explanation of how the gating block is defined, refer to the original paper [PSR⁺19]. The authors explore various gating functions—similar to A helpful supplementary resource on gating mechanisms in neural networks can be found [see](#))




Figure 34: Visualisation of the modified Transformer-XL block from [PSR⁺19]

19.6 Transformers in RL

While this section does not cover all transformer-based approaches in reinforcement learning, several notable works are worth highlighting for further reading. Decision Transformer [CLR⁺21] introduced a novel framing of reinforcement learning as a sequence modeling problem, treating trajectories as input-output sequences. More recently, PoliFormer [ZZE⁺24] demonstrated that scaling on-policy RL with transformer architectures can produce highly capable agents, particularly in navigation tasks.

19.7 Pros and Cons of Transformers in RL

Transformer-based models in RL offer several advantages. They are well-suited for evaluating memory and credit assignment capabilities, and they tend to perform strongly on tasks that require long-term memory. However, they do not inherently improve long-term credit assignment and often suffer from poor sample efficiency. Additionally, many existing RL benchmarks that aim to test memory or credit assignment often conflate the two or rely on tasks with predominantly short-term dependencies, limiting the insights that can be drawn about transformers' strengths in truly long-horizon settings.

19.8 Resources

There are several excellent resources for learning about transformers. A particularly insightful video series is the Neural Networks series by 3Blue1Brown, which includes a well-explained introduction to transformers [3Bl25]. Another highly recommended resource is Lilian Weng's blog post, Attention? Attention! [Wen18a], which provides a clear and intuitive breakdown of the underlying mechanisms.

20 Generative Models

Generative models aim to learn the underlying probability distribution of data, enabling them to generate new samples that closely resemble those in the training set. Unlike discriminative models, which focus on modeling the conditional probability $p(y|x)$ to find decision boundaries between classes, generative models learn the joint distribution $p(x, y)$, typically by modeling $p(x|y)$ along with the class prior $p(y)$. Additionally, once trained, generative models can be employed for classification tasks using Bayes' theorem:

$$p(y|x) = \frac{p(x|y)p(y)}{p(x)} = \frac{p(x|y)p(y)}{\sum_y p(x|y)p(y)}$$

20.1 Autoencoders

Autoencoders (AEs) are a type of neural network used for unsupervised learning, specifically for learning a compressed, efficient representation (encoding) of data. An autoencoder consists of two main parts:

1. **Encoder:** This part takes the input data x and maps it to a lower-dimensional space, often referred to as the latent space or embedding space. The encoder produces an encoded representation $z = f(x)$
2. **Decoder:** The decoder takes this compressed representation z and tries to reconstruct the original input x' from it. The goal is to minimize the difference between the input x and the reconstructed output x' .

The autoencoder is trained to minimize the reconstruction error, meaning it learns to represent the data in a lower-dimensional space while still being able to recover the original data as closely as possible.

While autoencoders are good at learning compact representations of data, they are not inherently designed for generative modeling. The problem is that the encoding learned by the autoencoder might not be structured in a way that facilitates easy sampling and generation. In other words, the latent space may not have meaningful properties (such as smoothness or continuity) that are useful for generating new data.

20.2 Variational Autoencoders (VAEs)

To address these limitations, Variational Autoencoders (VAEs) were introduced. VAEs are a probabilistic extension of autoencoders that are designed specifically for generative modeling. While AEs learn deterministic encodings, VAEs learn probabilistic encodings and allow for sampling from a continuous latent space. They work as follows :

1. **Encoder:** In a VAE, the encoder outputs two vectors: the mean $\mu(x)$ and the standard deviation $\sigma(x)$ of a distribution (often a Gaussian distribution) in the latent space. Instead of producing a single point in the latent space like in regular autoencoders, VAEs learn a distribution over the latent space. The encoder approximates the posterior distribution $p(z|x) \approx q_\psi(z|x)$.
2. **Reparameterization Trick:** To sample from this distribution in a differentiable way (which is required for backpropagation), VAEs use the reparameterization trick 9.4.3. The trick ensures that the gradient can flow through the random sampling process by expressing the latent variable z as:

$$z = \mu(x) + \sigma(x) \cdot \epsilon$$

where ϵ is a random noise sampled from a standard normal distribution (usually $\mathcal{N}(0, 1)$)

3. **Decoder:** After sampling from the distribution, the decoder reconstructs the input x' from the sampled latent variable z as in a standard autoencoder.

In VAEs the prior is chosen to be a standard multivariate normal distribution $p(z) \sim \mathcal{N}(0, I)$. The VAE is trained using maximum likelihood estimation

$$\begin{aligned}
\log p_\theta(x) &= \log \int p_\theta(z, x) dz && \text{(marginalization)} \\
&= \log \int \frac{q_\psi(z|x)}{q_\psi(z|x)} p_\theta(z, x) dz \\
&= \log \mathbb{E}_{z \sim q_\psi(z|x)} \left[\frac{p_\theta(z, x)}{q_\psi(z|x)} \right] \\
&\geq \mathbb{E}_{z \sim q_\psi(z|x)} \left[\log \frac{p_\theta(z, x)}{q_\psi(z|x)} \right] && \text{(Jensen's inequality)}^6
\end{aligned}$$

The resulting expression is known as the Evidence Lower Bound (ELBO), which provides a tractable objective for training the VAE. A more detailed derivation to provide intuitive understanding behind the ELBO and why it is structured the way it is can be found in [Luo22].




Figure 35: Visual illustration of an Autoencoder (top) and a Variational Autoencoder (bottom). Inspired from [Wen18b]

20.3 Diffusion

Diffusion models are a class of generative models that have gained significant attention for their ability to generate high-quality data (e.g., images, audio). These models operate by simulating two key processes: a forward diffusion process and a reverse denoising process.

In the forward process, noise is gradually added to the data until it becomes indistinguishable from a known distribution, typically Gaussian. The model then learns to reverse this process, step by step, denoising the data to recover samples that resemble those from the original dataset.

A key difference between diffusion models and VAEs is that diffusion models operate through an iterative process, unlike VAEs which use a single step for encoding. Additionally, diffusion models do not reduce the dimensionality of the latent space and instead of learning the encoder they employ simple Gaussian noise as an encoder.

⁶Jensens Inequality for concave function: $\mathbb{E}[f(x)] \geq f(\mathbb{E}[x])$

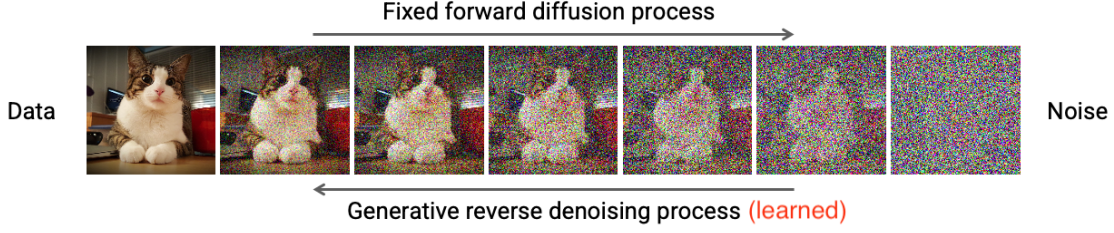


Figure 36: Illustration of the diffusion process

The relationship between the data and latent variables in diffusion models is typically modeled as a Markov chain, where each step depends only on the previous one.

20.3.1 Forward Diffusion Process

The forward diffusion process is a corruption process that takes a data sample (e.g., an image) and gradually adds noise to it in multiple steps, transforming the data into pure noise. This process can be described as follows:

Start with Data Start with an original data point x_0 (e.g., an image).

Add Noise in Steps At each time step t noise is added to the previous state z_{t-1} resulting in a more degraded version z_t . This process is modeled as a Markov chain:

$$q(z_t|z_{t-1}) = \mathcal{N}(z_t; \underbrace{\sqrt{1 - \beta_t}}_{\text{mean}} z_{t-1}, \underbrace{\beta_t I}_{\text{variance}})$$

$$= \sqrt{1 - \beta_t} z_{t-1} + \sqrt{\beta_t} \epsilon \sim \mathcal{N}(0, I)$$

We once again apply the reparameterization trick (see Section 9.4.3). To obtain z_{100} , the transformation would normally be applied sequentially over 100 steps. However, due to the properties of the Gaussian distribution, it is also possible to express this as a single-step transformation from the original input. To do this, we define the following parameters:

$$\alpha_t = 1 - \beta_t \quad \bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$

Here α_t defines the noise schedule, determining how much noise is added at each step. The one-step method can now be derived as follows:

$$\begin{aligned} z_t &= \sqrt{1 - \beta_t} z_{t-1} + \sqrt{\beta_t} \epsilon \\ &= \sqrt{\alpha_t} z_{t-1} + \sqrt{1 - \alpha_t} \epsilon \\ &= \sqrt{\alpha_t} (\sqrt{\alpha_{t-1}} z_{t-2} + \sqrt{1 - \alpha_{t-1}} \epsilon) + \sqrt{1 - \alpha_t} \epsilon \\ &= \sqrt{\alpha_t \alpha_{t-1}} z_{t-2} + \underbrace{\sqrt{\alpha_t - \alpha_t \alpha_{t-1}} \epsilon}_{\mathcal{N}(0, (\alpha_t - \alpha_t \alpha_{t-1}) I)} + \underbrace{\sqrt{1 - \alpha_t} \epsilon}_{\mathcal{N}(0, (1 - \alpha_t) I)} \\ &= \sqrt{\alpha_t \alpha_{t-1}} z_{t-2} + \sqrt{\alpha_t - \alpha_t \alpha_{t-1} + 1 - \alpha_t} \epsilon \quad (\text{sum of ind. gaussians } ^7) \\ &= \sqrt{\alpha_t \alpha_{t-1}} z_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \epsilon \\ &\vdots \quad (\text{repeat same steps from 14}) \\ &= \sqrt{\bar{\alpha}_t} z_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon \\ &\rightarrow q(z_t|z_0) = \mathcal{N}(z_t, \sqrt{\bar{\alpha}_t} z_0, \sqrt{1 - \bar{\alpha}_t} I) \end{aligned} \tag{18}$$

20.3.2 Reverse Diffusion Process

The reverse diffusion process is the key to generating new data points. It attempts to reverse the forward diffusion process and recover the original data from noise, essentially denoising the corrupted data step by step.

Start with Noise The process starts with a noise sample, z_t , and iteratively denoises it by applying learned denoising steps.

Learned Model The reverse process is learned by training a neural network to predict the denoised version of z_t at each step t . Essentially, the model learns to reverse the noising process at each step meaning $p_\theta(z_{t-1}|z_t)$

Generate Data By applying the reverse process step by step (starting from noise and gradually removing noise), the model generates a new data sample that resembles the distribution of the original training data. The reverse process can be modelled as a parametrised Gaussian:

$$p_\theta(z_{t-1}|z_t) = \mathcal{N}(z_{t-1}; \mu_\theta(z_t, t), \Sigma_\theta(z_t, t))$$

20.3.3 Training

Training a diffusion model involves learning how to reverse the diffusion process. Specifically, the model is trained to predict the data at time step $t - 1$ from the noisy data at time step t . This requires learning how to gradually denoise the data and requires a loss function that captures the reconstruction error over the multiple steps. Diffusion models are trained using a likelihood-based approach, aiming to maximize the log-likelihood of the data. This is typically done through the Evidence Lower Bound (ELBO), which incorporates various probabilistic components. A full derivation of this training objective can be found in the original paper by Ho et al. [HJA20]. The resulting bound is:

$$\log p_\theta(x) \geq \log p_\theta(x|z_1) - \sum_{t=2}^T \text{KL}(q(z_{t-1}|z_t, x) || p_\theta(z_{t-1}|z_t)) - \text{KL}(q(z_T|x) || p(z_T))$$

This can be further simplified since the last term does not depend on the parameters θ . Additionally, we know that $p_\theta(z_{t-1}|z_t)$ is the Gaussian we want to learn, and it turns out that $q(z_{t-1}|z_t, x)$ is also a Gaussian. Therefore, we can use the closed-form KL divergence between two Gaussians.

20.4 Score-Based Generative Models

Score-based generative models [SSDK⁺21] are a class of models that learn the score function of a data distribution — the gradient of the log-probability density with respect to the data:

$$\nabla_x \log p(x)$$

As with other generative models, the goal is to learn the data distribution $p_\theta(x)$ and generate new samples from it. A common formulation expresses the data distribution as:

$$p_\theta(x) = \frac{e^{-f_\theta(x)}}{Z_\theta}$$

However, the problem with this direct approach is that it requires knowledge of the normalization constant Z_θ , which involves solving an integral over the entire domain, a task that is often intractable. By instead taking the gradient of the log-probability, we get:

$$\nabla_x \log p_\theta(x) = \nabla_x \log e^{-f_\theta(x)} - \nabla_x \log Z_\theta = -\nabla_x f_\theta(x)$$

This is the core insight of score-based models: instead of modeling $p_\theta(x)$ directly, we model how the probability density changes in space. Once the score function is learned, we can generate samples by following these gradients in a stochastic manner (e.g., via Langevin dynamics).

⁷Sum of two independent Gaussian variables $u \sim \mathcal{N}(\mu_u, \sigma_u^2)$ and $v \sim \mathcal{N}(\mu_v, \sigma_v^2)$ is a gaussian variable $(u + v) \sim \mathcal{N}(\mu_u + \mu_v, \sigma_u^2 + \sigma_v^2)$

20.5 Flow Matching

Flow matching [LCBH⁺23] is a method for training generative models based on the idea of aligning data distributions through normalizing flows. A normalizing flow is a series of invertible transformations applied to a simple distribution (e.g., a Gaussian) to model complex data distributions.

Flow matching introduces a novel way of training generative models by matching the gradients of the data density at different stages of the transformation. This is done by learning the flow that transforms a simple distribution (like Gaussian noise) into a more complex data distribution. Really helpful resource can be found [here](#)

20.6 Connecting Diffusion Models and Reinforcement Learning

An interesting connection exists between diffusion models and reinforcement learning (RL): diffusion models can be used to generate new plausible policies that resemble those from the original dataset. In other words, instead of learning from the dataset directly, we learn a distribution over good trajectories (or policies), and then sample new ones using diffusion processes. This perspective opens up new possibilities in offline RL and imitation learning, where the goal is to generalize and generate diverse, high-performing behaviors from limited data.

20.7 Resources

In this section, we did not delve into the full mathematical derivations behind the various generative modeling approaches, as these tend to be quite lengthy. For readers interested in the detailed math, I recommend reading the original papers. Additionally, there are some excellent resources that explain these concepts more intuitively. A great video series covering diffusion models, score-based diffusion models, and flow matching can be found here [Out]. Lilian Weng’s blog post “What Are Diffusion Models?” [Wen21] is also very well written and highly recommended. For a deeper understanding of score-based models, Yang Song’s blog [Son21] is one of the best resources available. A particularly insightful paper that highlights the connections between the various generative modeling approaches discussed is [Luo22].

21 Distributional Reinforcement Learning

Distributional Reinforcement Learning (DRL) is an extension of traditional reinforcement learning (RL) that focuses on modeling the distribution of return values (rewards) rather than just predicting the expected return.

In traditional RL, the goal is often to maximise the expected cumulative reward (value function) for a given state .

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s \right]$$

However, in distributional reinforcement learning, instead of focusing on a single expected value, we learn the full distribution of the returns. This means that, instead of predicting just the expected return (like a scalar value), we model the entire range of possible future rewards and their probabilities.

Imagine you're playing a video game where, at some points, you can make a risky move that might either yield a high reward or result in a penalty. In traditional RL, you might expect the average reward (e.g., 5 points), but in distributional RL, you can learn the probability distribution of rewards—e.g., 70% chance for +10 points, 20% chance for +5 points, and 10% chance for -5 points. This distributional approach helps the agent better weigh its decisions.

21.1 Resources

The content presented here is intended as a brief overview and will be expanded and improved in the future. For a deeper understanding, a highly recommended lecture is "Deep RL Bootcamp Frontiers Lecture I: Recent Advances, Frontiers, and Future of Deep RL" [\[Pri18\]](#).

References

- [3Bl25] 3Blue1Brown. Neural networks, 2025. https://www.youtube.com/playlist?list=PLZHQB0WTQDNU6R1_67000Dx_ZCJB-3pi.
- [AAB⁺19] Akshay Agrawal, Brandon Amos, Shane Barratt, Stephen Boyd, Steven Diamond, and Zico Kolter. Differentiable convex optimization layers, 2019.
- [Ach18] Josh Achiam. Openai spinning up, 2018. <https://spinningup.openai.com/en/latest/index.html>.
- [AK22] Sergey Levine Aviral Kumar, Ilya Kostrikov. Should i use offline rl or imitation learning?, Apr 25, 2022. <https://bair.berkeley.edu/blog/2022/04/25/rl-or-bc/>, Last accessed on 2025-04-18.
- [BGNN19] Daniel S. Brown, Wonjoon Goo, Prabhat Nagarajan, and Scott Niekum. Extrapolating beyond suboptimal demonstrations via inverse reinforcement learning from observations, 2019.
- [BKM17] David M. Blei, Alp Kucukelbir, and Jon D. McAuliffe. Variational inference: A review for statisticians. *Journal of the American Statistical Association*, 112(518):859–877, April 2017.
- [BPB⁺24] Aditya Bhatt, Daniel Palenicek, Boris Belousov, Max Argus, Artemij Amiranashvili, Thomas Brox, and Jan Peters. Crossq: Batch normalization in deep reinforcement learning for greater sample efficiency and simplicity, 2024.
- [CCML18] Kurtland Chua, Roberto Calandra, Rowan McAllister, and Sergey Levine. Deep reinforcement learning in a handful of trials using probabilistic dynamics models, 2018.
- [CLR⁺21] Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Michael Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling, 2021.
- [Com24] Wikimedia Commons. File:pyplot overfitting.png — wikimedia commons, the free media repository, 2024. [Online; accessed 27-April-2025].
- [Dee21] Google DeepMind. Deepmind x ucl - deep learning lecture series 2021, 2021. <https://www.youtube.com/playlist?list=PLqYmG7hTraZDVH599EI1EWsU0sJbAodm>.
- [DYY⁺19] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. Transformer-xl: Attentive language models beyond a fixed-length context, 2019.
- [FCAL16] Chelsea Finn, Paul Christiano, Pieter Abbeel, and Sergey Levine. A connection between generative adversarial networks, inverse reinforcement learning, and energy-based models, 2016.
- [FLA16] Chelsea Finn, Sergey Levine, and Pieter Abbeel. Guided cost learning: Deep inverse optimal control via policy optimization, 2016.
- [Fra19] Katerina Fragkiadaki. Deep reinforcement learning and controlspring 2019, cmu 10403, 2019. <https://www.andrew.cmu.edu/course/10-403/>.
- [FvHM18] Scott Fujimoto, Herke van Hoof, and David Meger. Addressing function approximation error in actor-critic methods, 2018.
- [Haf] Danijar Hafner. Danijar hafner webiste. <https://danijar.com>.
- [hc] (<https://math.stackexchange.com/users/253273/a-Why-is-the-lagrange-dual-function-concave?> Mathematics Stack Exchange. URL:<https://math.stackexchange.com/q/1374410> (version: 2021-04-15)).

- [HE16] Jonathan Ho and Stefano Ermon. Generative adversarial imitation learning, 2016.
- [HJA20] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models, 2020.
- [HLBN20] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control: Learning behaviors by latent imagination, 2020.
- [HLF⁺19] Danijar Hafner, Timothy Lillicrap, Ian Fischer, Ruben Villegas, David Ha, Honglak Lee, and James Davidson. Learning latent dynamics for planning from pixels, 2019.
- [HN24] Maximilian Hüttenrauch and Gerhard Neumann. Robust black-box optimization for stochastic search and episodic reinforcement learning. *Journal of Machine Learning Research*, 25(153):1–44, 2024.
- [Hui18] Jonathan Hui. Alphago: How it works technically?, May 13, 2018. <https://jonathan-hui.medium.com/alphago-how-it-works-technically-26ddcc085319>.
- [HZAL18] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor, 2018.
- [JFZL21] Michael Janner, Justin Fu, Marvin Zhang, and Sergey Levine. When to trust your model: Model-based policy optimization, 2021.
- [KFTL19] Aviral Kumar, Justin Fu, George Tucker, and Sergey Levine. Stabilizing off-policy q-learning via bootstrapping error reduction, 2019.
- [KHSL22] Aviral Kumar, Joey Hong, Anikait Singh, and Sergey Levine. When should we prefer offline reinforcement learning over behavioral cloning?, 2022.
- [KNL21] Ilya Kostrikov, Ashvin Nair, and Sergey Levine. Offline reinforcement learning with implicit q-learning, 2021.
- [Kno10] David Knowles. Lagrangian duality for dummies, November 13, 2010. https://cs.stanford.edu/people/davidknowles/lagrangian_duality.pdf.
- [KZTL20] Aviral Kumar, Aurick Zhou, George Tucker, and Sergey Levine. Conservative q-learning for offline reinforcement learning, 2020.
- [LA24] Machine Learning and AI Academy. Ai - deep reinforcement learning made easy again! - crossq, February 2, 2024. <https://www.youtube.com/watch?v=UpACd0z6ntw>.
- [LCBH⁺23] Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling, 2023.
- [LDHC22] Mat Leonard, DanB, Ryan Holbrook, and Alexis Cook. Word vectors, 2022. <https://www.kaggle.com/code/matleonard/word-vectors>.
- [Lev18] Sergey Levine. Reinforcement learning and control as probabilistic inference: Tutorial and review, 2018.
- [lh] layman (<https://math.stackexchange.com/users/131740/layman>). Mathematical expectation. $e[e[x]] = e[x]$. Mathematics Stack Exchange. URL:<https://math.stackexchange.com/q/2034867> (version: 2016-11-29).
- [LHP⁺19] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning, 2019.
- [LKTF20] Sergey Levine, Aviral Kumar, George Tucker, and Justin Fu. Offline reinforcement learning: Tutorial, review, and perspectives on open problems, 2020.
- [LKX⁺19] Corey Lynch, Mohi Khansari, Ted Xiao, Vikash Kumar, Jonathan Tompson, Sergey Levine, and Pierre Sermanet. Learning latent plans from play, 2019.

- [LSM⁺20] Sergey Levine, Kyle Stachowicz, Vivek Myers, Joey Hong, and Kevin Black. Cs 285 at uc berkeley, deep reinforcement learning, 2020. <https://rail.eecs.berkeley.edu/deeprlcourse/>.
- [Luo22] Calvin Luo. Understanding diffusion models: A unified perspective, 2022.
- [MIL23] MICHELE MILESI. Dreamer v1 behind our pytorch implementation of the state-of-the-art rl algorithm, 16 JUNE 2023. https://eclecticsheep.ai/2023/06/16/dreamer_v1.html.
- [MKS⁺13] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning, 2013.
- [MRCA18] Nikhil Mishra, Mostafa Rohaninejad, Xi Chen, and Pieter Abbeel. A simple neural attentive meta-learner, 2018.
- [NOJ⁺24a] Michal Nauman, Mateusz Ostaszewski, Krzysztof Jankowski, Piotr Miłoś, and Marek Cygan. Bigger, regularized, optimistic, 2024. <https://sites.google.com/view/bro-agent/introduction>.
- [NOJ⁺24b] Michal Nauman, Mateusz Ostaszewski, Krzysztof Jankowski, Piotr Miłoś, and Marek Cygan. Bigger, regularized, optimistic: scaling for compute and sample-efficient continuous control, 2024.
- [OBV⁺21] Fabian Otto, Philipp Becker, Ngo Anh Vien, Hanna Carolin Ziesche, and Gerhard Neumann. Differentiable trust region layers for deep reinforcement learning, 2021.
- [Onl24] Stanford Online. Stanford cs234 reinforcement learning spring 2024 emma brunskill, 2024. <https://www.youtube.com/playlistlist=PLoROMvodv4rN4wG6Nk6sNpTEbuOSosZdX>.
- [Out] Outlier. Youtube channel: Outlier. <https://www.youtube.com/@outlier/videos>.
- [Pri18] AI Prism. Deep rl bootcamp at berkeley, 2018. <https://www.youtube.com/@aiprism1155/videos>.
- [PSR⁺19] Emilio Parisotto, H. Francis Song, Jack W. Rae, Razvan Pascanu, Caglar Gulcehre, Siddhant M. Jayakumar, Max Jaderberg, Raphael Lopez Kaufman, Aidan Clark, Seb Noury, Matthew M. Botvinick, Nicolas Heess, and Raia Hadsell. Stabilizing transformers for reinforcement learning, 2019.
- [RAI20] RAIL. Deep reinforcement learning: Cs 285 fall 2020 - youtube, 2020. https://www.youtube.com/playlist?list=PL_iWQ0sE6TfURIIhCrIt-wj9ByIVpbfGc.
- [SAH⁺20] Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Muzero: Mastering go, chess, shogi and atari without rules, 23 DECEMBER 2020. <https://deepmind.google/discover/blog/muzero-mastering-go-chess-shogi-and-atari-without-rules/>.
- [SB18] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. A Bradford Book, Cambridge, MA, USA, 2018.
- [sl] scikit learn. Gaussian processes regression: basic introductory example. https://scikit-learn.org/stable/auto_examples/gaussian_process/plot_gpr_noisy_targets.html#sphx-glr-auto-examples-gaussian-process-plot-gpr-noisy-targets-py.
- [SLH⁺14] David Silver, Guy Lever, Nicolas Heess, Thomas Degris, Daan Wierstra, and Martin Riedmiller. Deterministic policy gradient algorithms. In *Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32*, ICML’14, page I-387–I-395. JMLR.org, 2014.

- [SLM⁺17] John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel. Trust region policy optimization, 2017.
- [Son21] Yang Song. Generative modeling by estimating gradients of the data distribution, 2021. <https://yang-song.net/blog/2021/score/>.
- [SSDK⁺21] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations, 2021.
- [SWD⁺17] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [VAR20] JOSH VARTY. Alphazero a step-by-step look at alpha zero and monte carlo tree search, 2020. <http://joshvarty.github.io/AlphaZero/>.
- [vHGS15] Hado van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double q-learning, 2015.
- [Vie19] Tim Vieira. The likelihood-ratio gradient, APR 20, 2019. <https://timvieira.github.io/blog/post/2019/04/20/the-likelihood-ratio-gradient/>.
- [VSP⁺23] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2023.
- [Wen18a] Lilian Weng. Attention? attention!, 2018. <https://lilianweng.github.io/posts/2018-06-24-attention/>.
- [Wen18b] Lilian Weng. From autoencoder to beta-vae, 2018. <https://lilianweng.github.io/posts/2018-08-12-vae/>.
- [Wen18c] Lilian Weng. Policy gradient algorithms, April 8, 2018. <https://lilianweng.github.io/posts/2018-04-08-policy-gradient/>.
- [Wen19] Lilian Weng. Evolution strategies, 2019. <https://lilianweng.github.io/posts/2019-09-05-evolution-strategies/>.
- [Wen21] Lilian Weng. What are diffusion models?, 2021. <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>.
- [Wil19] Cody Marie Wild. It’s only natural: An excessively deep dive into natural gradient optimization, Mar 9, 2019. <https://medium.com/data-science/its-only-natural-an-excessively-deep-dive-into-natural-gradient-optimization-75d464b89dbb>.
- [WSH⁺16] Ziyu Wang, Tom Schaul, Matteo Hessel, Hado van Hasselt, Marc Lanctot, and Nando de Freitas. Dueling network architectures for deep reinforcement learning, 2016.
- [ZZE⁺24] Kuo-Hao Zeng, Zichen Zhang, Kiana Ehsani, Rose Hendrix, Jordi Salvador, Alvaro Herastí, Ross Girshick, Aniruddha Kembhavi, and Luca Weihs. Poliformer: Scaling on-policy rl with transformers results in masterful navigators. *arXiv*, 2024.